

3rd
Week

세번째 봅겠습니다 ?!

▷ 출석 체크도 한 번 해보시면 어떠세요?!

- <https://modulabs.co.kr/>
- 모두연 홈페이지 → 로그인 → 마이페이지 → 참여한 랩·풀잎 → 자세히 보기 → 내 풀잎스쿨 출석 확인하기

▷ Ground Rule

- 가급적 지각/결석 하지 않기
- 가급적 Camera 켜 놓고 수업 참여하기
- 가급적 적극적으로 참여하기
- 3시간이 넘더라도 배고프다고 화내지 않기
- Slack 잊지 않기
- 꼭 끝까지 함께하기

잡담 &
지난 수업 관련 이야기

///

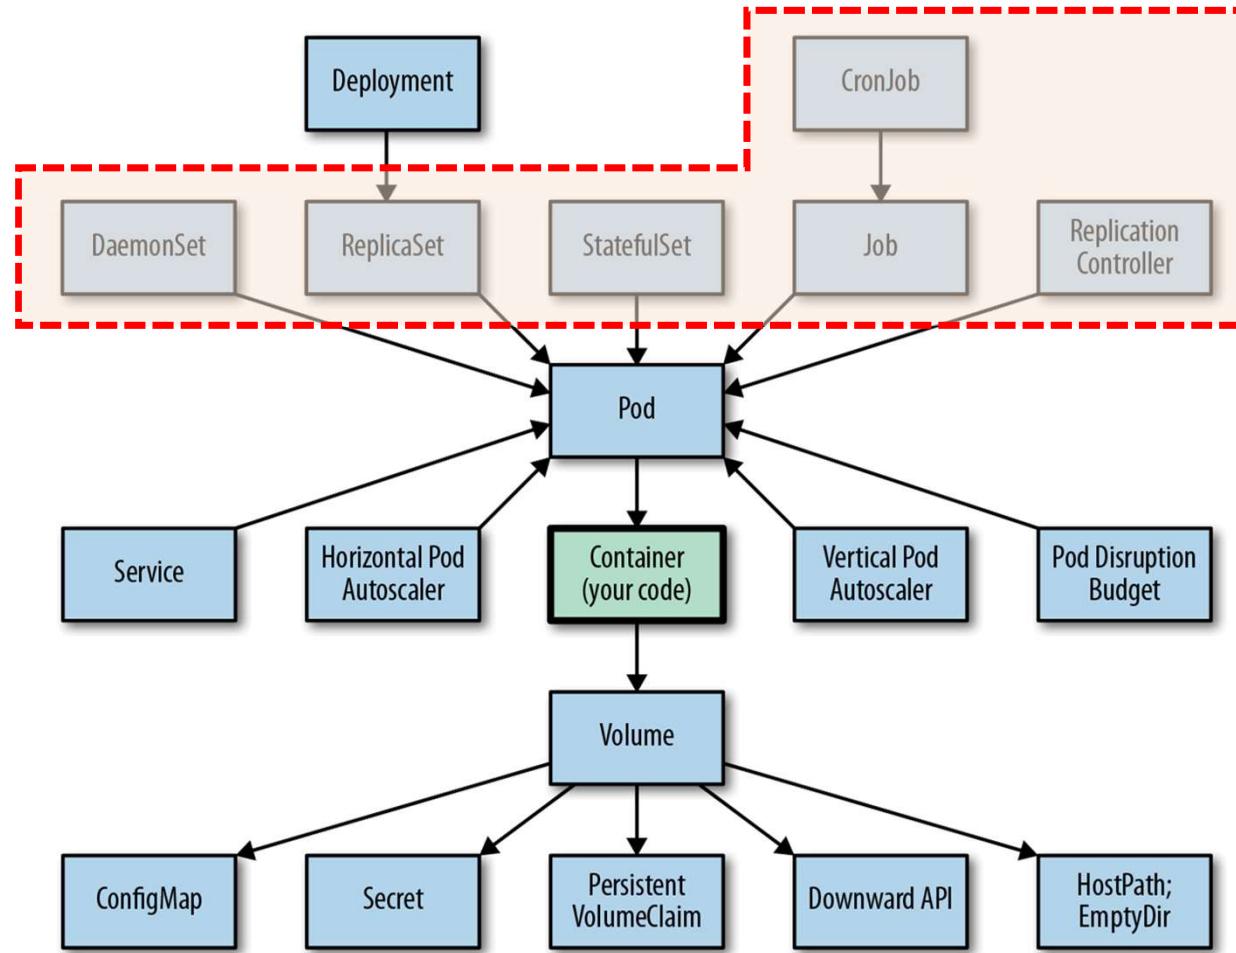
Agenda

- ▷ [10m] make-friendship
- ▷ [30m] Flip - ReplicaSet/DaemonSet/Job/CronJob
- ▷ [30m] Wrap-Up & Hands-On
- ▷ [10m] Break-Time

- ▷ [30m] Flip - Service (ClusterIP/NodePort/ExternalName)
- ▷ [20m] Wrap-Up & Hands-On
- ▷ [10m] Break-Time

- ▷ [20m] Quiz
- ▷ [10m] Q&A

Kubernetes concepts for developers



※ 참고 : <https://www.oreilly.com/library/view/kubernetes-patterns/9781492050278/ch01.html>

///

Flip Learning

(ReplicaSet/DaemonSet/Job/CronJob)

정현찬님

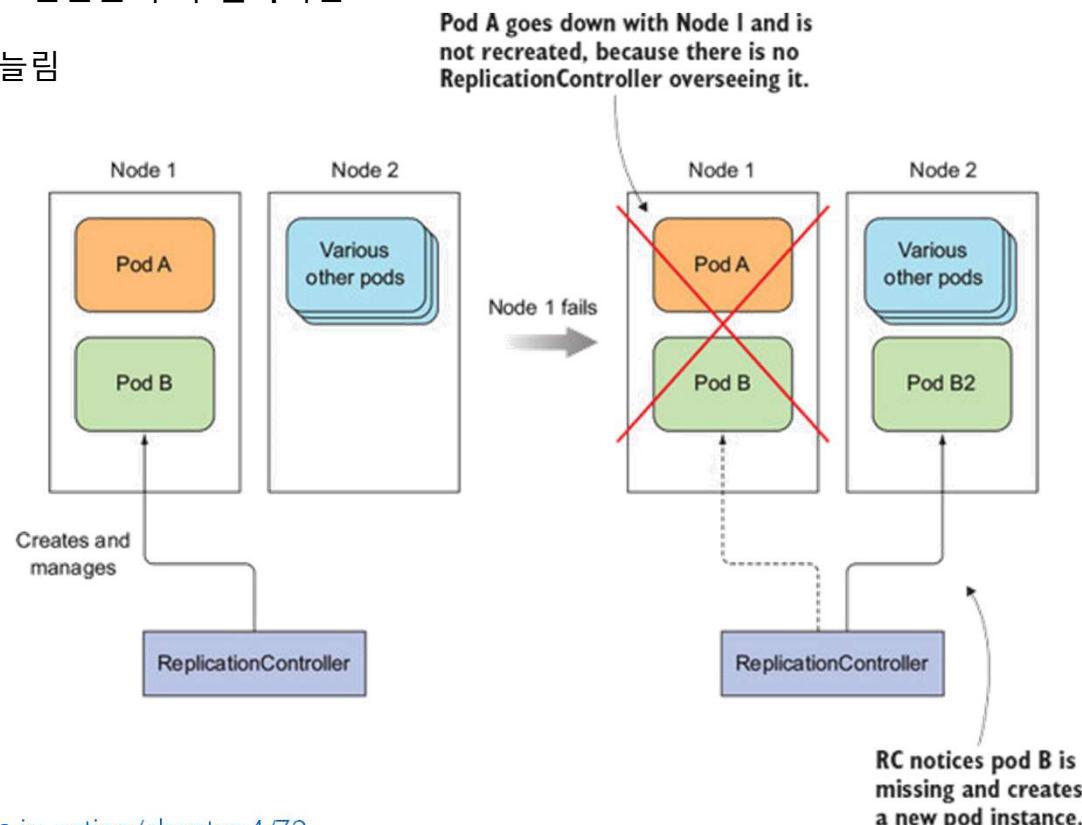
///

Kubernetes

ReplicaSet

ReplicationController - 1/2

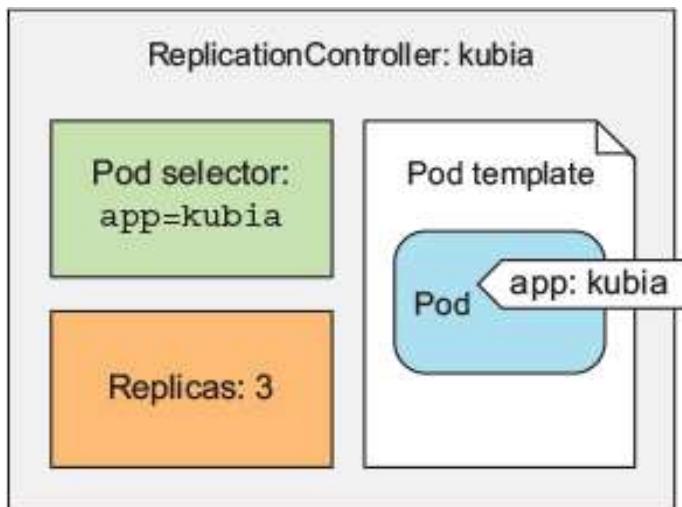
- Pod가 항상 실행되도록 보장하는 리소스
 - . 어떤 이유로든 Pod가 사라지면 사라진 Pod를 감지해 교체 Pod를 생성
- 지속적으로 실행 중인 Pod 목록을 모니터링하고 선언된 수와 일치시킴
 - . 너무 많은 Pod가 실행 중이면 줄이고, 적으면 늘림



※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-4/79>

ReplicationController - 2/2

- ReplicationController 세 가지 필수 요소
 - . **replicas**: Pod가 실행되어야 하는 수
 - . **selector**: ReplicationController 범위에 있는 Pod 결정
 - . **template**: 새로운 Pod replica를 만들 때 사용



```
apiVersion: v1
kind: ReplicationController

metadata:
  name: kubia

spec:
  replicas: 3
  selector:
    app: kubia
  template:
    metadata:
      labels:
        app: kubia
    spec:
      containers:
        - name: kubia
          image: luksa/kubia
        ports:
          - containerPort: 8080
```

※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-4/92>

ReplicaSet

- ReplicationController와 똑같이 동작
 - . 여기에 더해 selector에서 풍부한 표현식을 사용할 수 있음
 - . 특정 키가 있는 label을 갖는 Pod를 매칭
 - . label 조건을 Or/And로 정의
 - 일반적으로 직접 사용하지는 않고,
- Deployment** 리소스를 생성할 때 자동으로 생성됨

```
apiVersion: apps/v1
kind: ReplicaSet

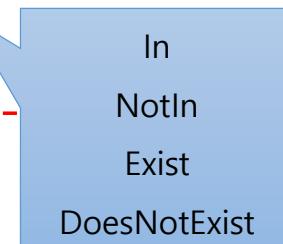
metadata:
  name: kubia

spec:
  replicas: 3

  selector:
    matchExpressions:
      - key: app
        operator: In
        values:
          - kubia

  template:
    metadata:
      labels:
        app: kubia

  spec:
    containers:
      - name: kubia
        image: luksa/kubia
        ports:
          - containerPort: 8080
```



///

K8s : ReplicaSet Hands-On

ReplicaSet Create

YAML 파일 작성해서 ReplicaSet 실행해 보자!

rs-node-web.yaml

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-labels

spec:
  replicas: 3

  selector:
    matchLabels:
      app: node-web

  template:
    metadata:
      labels:
        app: node-web
    spec:
      containers:
      - name: node-web
        image: whatwant/node-web:1.0
      ports:
      - containerPort: 8080
```

```
remote > cd advanced-kubernetes/03-week/ReplicaSet
```

```
remote > kubectl create -f ./rs-node-web.yaml
```

```
replicaset.apps/rs-labels created
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-8677p	1/1	Running	0	7s	10.233.110.61	worker1	<none>	<none>
rs-labels-gdj7n	1/1	Running	0	7s	10.233.103.77	worker2	<none>	<none>
rs-labels-z4frf	1/1	Running	0	7s	10.233.103.76	worker2	<none>	<none>

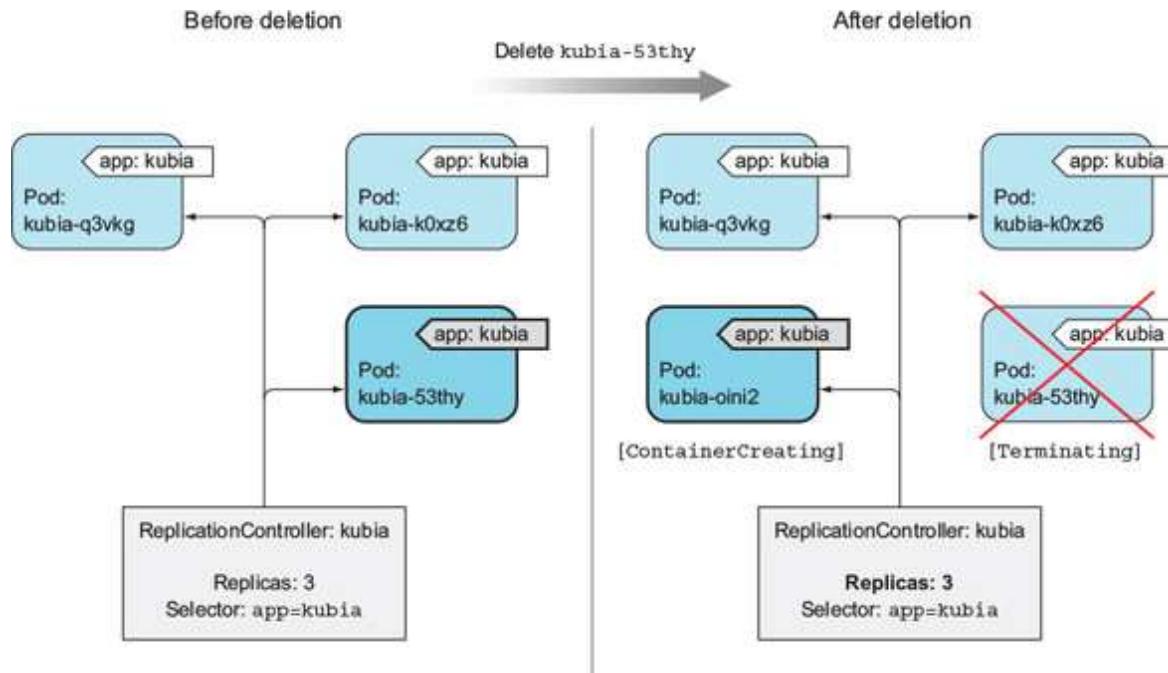
```
remote > kubectl get replicsets -o wide
```

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
rs-labels	3	3	3	5m6s	node-web	whatwant/node-web:1.0	app=node-web

Pod Delete in ReplicaSet - 1/2

Pod가 삭제 되면 어떻게 될까?

삭제가 되는 등의 사유로 'selector'로 맵핑이 되는 Pod의 숫자가 'Replicas'와 맞지 않으면, 'template' 참조하여 Pod를 생성한다.
. 기존 Pod를 복제하는 것이 아니라 현재 시점의 template 정보를 참조하여 신규로 생성 → Update 기법으로 사용 가능



※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-4/141>

Pod Delete in ReplicaSet - 2/2

```
remote > kubectl delete pod rs-labels-8677p
```

```
pod "rs-labels-8677p" deleted
```

```
remote > kubectl get pods -o wide -w
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-8677p	1/1	Running	0	68m	10.233.110.61	worker1	<none>	<none>
rs-labels-gdj7n	1/1	Running	0	68m	10.233.103.77	worker2	<none>	<none>
rs-labels-z4frf	1/1	Running	0	68m	10.233.103.76	worker2	<none>	<none>
rs-labels-8677p	1/1	Terminating	0	68m	10.233.110.61	worker1	<none>	<none>
rs-labels-phbs7	0/1	Pending	0	0s	<none>	<none>	<none>	<none>
rs-labels-phbs7	0/1	Pending	0	0s	<none>	worker1	<none>	<none>
rs-labels-phbs7	0/1	ContainerCreating	0	0s	<none>	worker1	<none>	<none>
rs-labels-phbs7	0/1	ContainerCreating	0	1s	<none>	worker1	<none>	<none>
rs-labels-phbs7	1/1	Running	0	1s	10.233.110.62	worker1	<none>	<none>
rs-labels-8677p	1/1	Terminating	0	68m	10.233.110.61	worker1	<none>	<none>
rs-labels-8677p	0/1	Terminating	0	68m	<none>	worker1	<none>	<none>
rs-labels-8677p	0/1	Terminating	0	68m	<none>	worker1	<none>	<none>
rs-labels-8677p	0/1	Terminating	0	68m	<none>	worker1	<none>	<none>

^C%

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-gdj7n	1/1	Running	0	70m	10.233.103.77	worker2	<none>	<none>
rs-labels-phbs7	1/1	Running	0	2m16s	10.233.110.62	worker1	<none>	<none>
rs-labels-z4frf	1/1	Running	0	70m	10.233.103.76	worker2	<none>	<none>

Replicas Increase/Decrease (+ change Editor) - 1/2

replicas 숫자를 변경하고 싶으면 어떻게 해야 할까?

resource 내역을 수정하면 되는데, 기본 editor로 vi가 설정되어 있다.

개인적인 취향으로 기본 editor를 'nano'로 변경하고 시작해보겠다.

```
remote > export KUBE_EDITOR=nano
```

이제 설정을 변경해보자!

```
remote > kubectl edit replicaset rs-labels
```

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  creationTimestamp: "2021-12-31T12:51:16Z"
  generation: 1
  name: rs-labels
  namespace: default
  resourceVersion: "1472985"
  uid: 6c6ce66f-e894-4177-b2d2-64f3d1f6bd96
spec:
  replicas: 5
  selector:
    matchLabels:
      app: node-web
...
```

Replicas Increase/Decrease (+ change Editor) - 2/2

replicas 숫자를 변경하고 싶으면 어떻게 해야 할까?

```
remote > kubectl get pods -o wide -w
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-gdj7n	1/1	Running	0	160m	10.233.103.77	worker2	<none>	<none>
rs-labels-phbs7	1/1	Running	0	92m	10.233.110.62	worker1	<none>	<none>
rs-labels-z4frf	1/1	Running	0	160m	10.233.103.76	worker2	<none>	<none>
rs-labels-dxjvz	0/1	Pending	0	0s	<none>	<none>	<none>	<none>
rs-labels-zq25x	0/1	Pending	0	0s	<none>	<none>	<none>	<none>
rs-labels-dxjvz	0/1	Pending	0	0s	<none>	worker1	<none>	<none>
rs-labels-zq25x	0/1	Pending	0	0s	<none>	worker2	<none>	<none>
rs-labels-zq25x	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>
rs-labels-dxjvz	0/1	ContainerCreating	0	0s	<none>	worker1	<none>	<none>
rs-labels-zq25x	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>
rs-labels-dxjvz	0/1	ContainerCreating	0	0s	<none>	worker1	<none>	<none>
rs-labels-zq25x	1/1	Running	0	1s	10.233.103.78	worker2	<none>	<none>
rs-labels-dxjvz	1/1	Running	0	1s	10.233.110.63	worker1	<none>	<none>

^C%

replicas 숫자를 변경하고 싶으면 어떻게 해야 할까?

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-dxjvz	1/1	Running	0	2m48s	10.233.110.63	worker1	<none>	<none>
rs-labels-gdj7n	1/1	Running	0	163m	10.233.103.77	worker2	<none>	<none>
rs-labels-phbs7	1/1	Running	0	94m	10.233.110.62	worker1	<none>	<none>
rs-labels-z4frf	1/1	Running	0	163m	10.233.103.76	worker2	<none>	<none>
rs-labels-zq25x	1/1	Running	0	2m48s	10.233.103.78	worker2	<none>	<none>

Quiz

**ReplicaSet을 삭제해도
Pod는 유지가 될까?**

ReplicaSet Delete

```
remote > kubectl get replicaset -o wide
```

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
rs-labels	5	5	5	27s	node-web	whatwant/node-web:1.0	app=node-web

```
remote > kubectl delete replicaset rs-labels
```

```
replicaset.apps "rs-labels" deleted
```

```
remote > kubectl get pods -o wide -w
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-68jpd	1/1	Running	0	24s	10.233.103.81	worker2	<none>	<none>
rs-labels-9gbf7	1/1	Running	0	24s	10.233.103.80	worker2	<none>	<none>
rs-labels-mq65n	1/1	Running	0	39s	10.233.110.65	worker1	<none>	<none>
rs-labels-vl8gw	1/1	Running	0	39s	10.233.110.64	worker1	<none>	<none>
rs-labels-zs9h2	1/1	Running	0	39s	10.233.103.79	worker2	<none>	<none>
rs-labels-zs9h2	1/1	Terminating	0	66s	10.233.103.79	worker2	<none>	<none>
rs-labels-9gbf7	1/1	Terminating	0	51s	10.233.103.80	worker2	<none>	<none>
rs-labels-vl8gw	1/1	Terminating	0	66s	10.233.110.64	worker1	<none>	<none>
rs-labels-mq65n	1/1	Terminating	0	66s	10.233.110.65	worker1	<none>	<none>
rs-labels-68jpd	1/1	Terminating	0	51s	10.233.103.81	worker2	<none>	<none>

```
^C%
```

ReplicaSet Delete (--cascade=orphan)

ReplicaSet은 삭제하지만, 이미 생성된 Pod는 유지하고 싶다면?

ReplicaSet 다시 생성하고 삭제를 진행해보자.

```
remote > kubectl create -f ./rs-node-web.yaml
```

```
replicaset.apps/rs-labels created
```

```
remote > kubectl get replicsets -o wide
```

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
rs-labels	3	3	3	9s	node-web	whatwant/node-web:1.0	app=node-web

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-29r2b	1/1	Running	0	12s	10.233.103.85	worker2	<none>	<none>
rs-labels-747l6	1/1	Running	0	12s	10.233.103.84	worker2	<none>	<none>
rs-labels-xtjbz	1/1	Running	0	12s	10.233.110.67	worker1	<none>	<none>

```
remote > kubectl delete replicaset rs-labels --cascade=orphan
```

```
replicaset.apps "rs-labels" deleted
```

```
remote > kubectl get replicsets -o wide
```

```
No resources found in default namespace.
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-29r2b	1/1	Running	0	94s	10.233.103.85	worker2	<none>	<none>
rs-labels-747l6	1/1	Running	0	94s	10.233.103.84	worker2	<none>	<none>
rs-labels-xtjbz	1/1	Running	0	94s	10.233.110.67	worker1	<none>	<none>

///

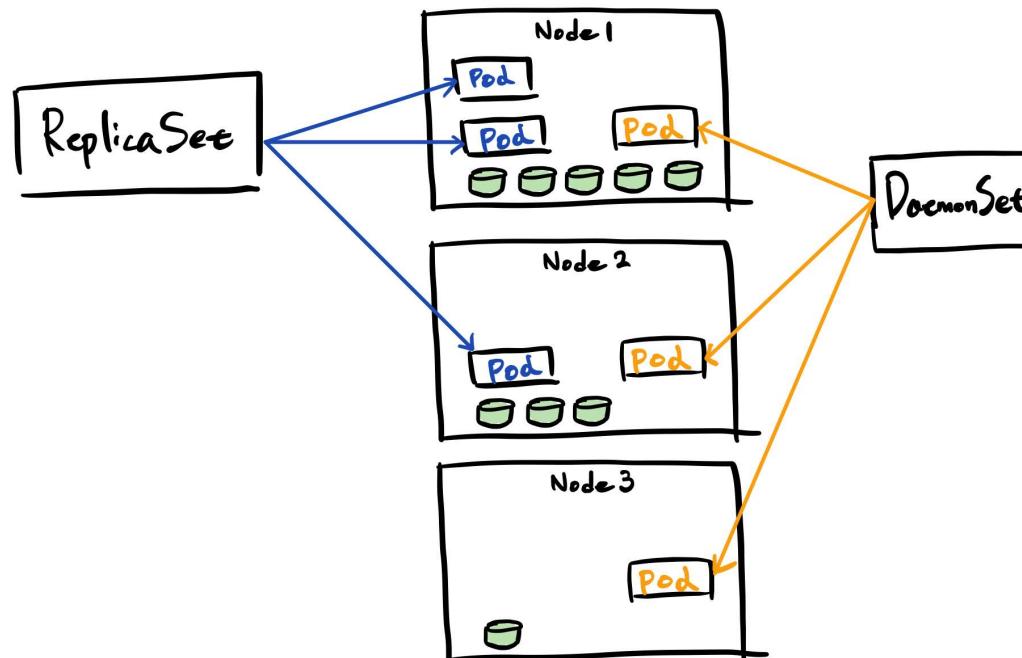
Kubernetes

DaemonSet

DaemonSet is ...

- DaemonSet을 활용하면 cluster의 모든 node에, **node당 하나의 Pod**를 배포할 수 있음
 - . 시스템 수준의 작업을 수행하는 인프라 관련 Pod (로깅, 모니터링), Kube-proxy도 DaemonSet의 일종
- 특정 node에만 Pod를 배포하려면 nodeSelector 속성 지정

o DaemonSet



※ 참고 : <https://zunoxi.github.io/devops/2020/11/07/devops-k8s-daemonset/>

///

K8s : DaemonSet Hands-On

Daemon Create

YAML 파일 작성해서 DaemonSet 생성도 하고 실제로 해보자 !

ds-node-web.yaml

```
apiVersion: apps/v1
kind: DaemonSet

metadata:
  name: ds-labels

spec:
  selector:
    matchLabels:
      app: node-web

  template:
    metadata:
      labels:
        app: node-web

  spec:
    containers:
      - name: node-web
        image: whatwant/node-web:1.0
        ports:
          - containerPort: 8080
```

```
remote > cd advanced-kubernetes/03-week/DaemonSet
```

```
remote > kubectl create -f ./ds-node-web.yaml
```

```
daemonset.apps/ds-labels created
```

```
> kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
ds-labels-4gtfx	1/1	Running	0	8s	10.233.103.86	worker2	<none>	<none>
ds-labels-kljvl	1/1	Running	0	8s	10.233.110.68	worker1	<none>	<none>

```
remote > kubectl get daemonsets -o wide
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE	CONTAINERS	IMAGES	SELECTOR
ds-labels	2	2	2	2	2	<none>	11s	node-web	whatwant/node-web:1.0	app=node-web

Daemon Delete

지우는 방식은 계속 유사하다 ...

```
remote > kubectl get daemonsets -o wide
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE	CONTAINERS	IMAGES	SELECTOR
ds-labels	2	2	2	2	2	<none>	11s	node-web	whatwant/node-web:1.0	app=node-web

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
ds-labels-f8dld	1/1	Running	0	46s	10.233.110.69	worker1	<none>	<none>
ds-labels-mspkq	1/1	Running	0	46s	10.233.103.87	worker2	<none>	<none>

```
remote > kubectl delete daemonset ds-labels
```

```
daemonset.apps "ds-labels" deleted
```

```
remote > kubectl get daemonsets -o wide
```

```
No resources found in default namespace.
```

```
remote > kubectl get pods -o wide
```

```
No resources found in default namespace.
```

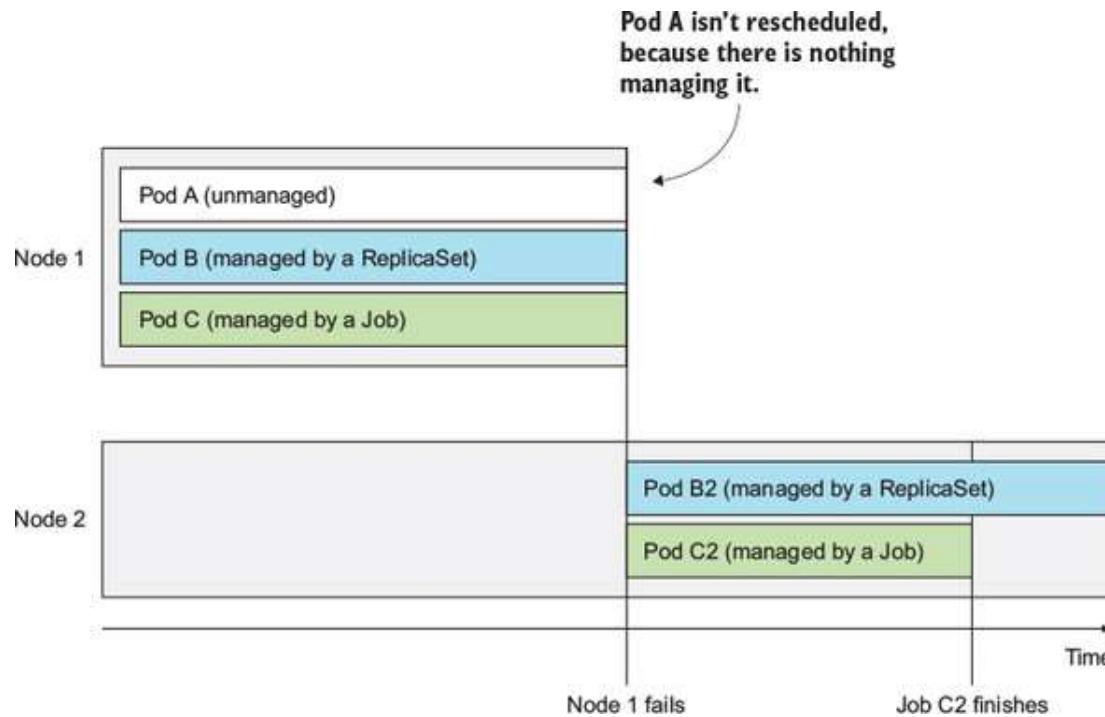
///

Kubernetes

Job

Job is ...

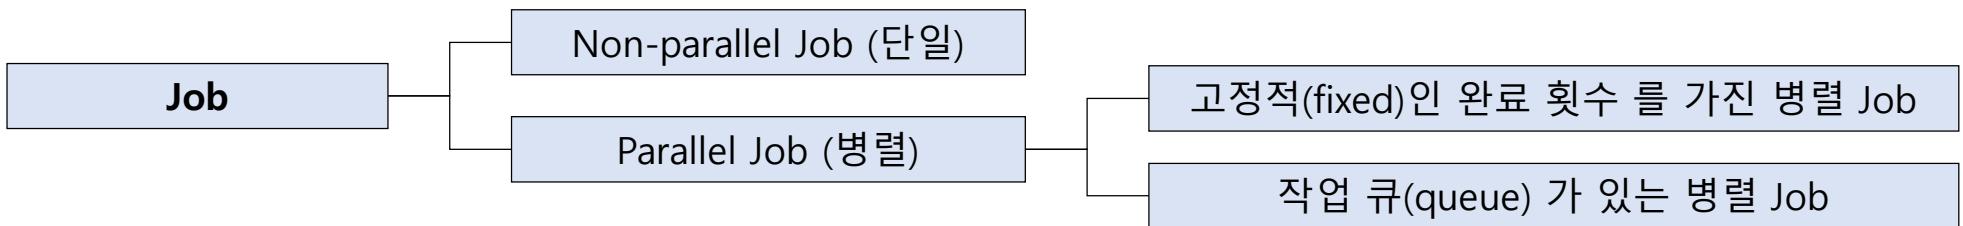
- 완료 가능한 단일 태스크를 수행하는 Pod를 실행
 - . 프로세스 종료 이후에도 다시 실행되지 않음
 - . node 장애 발생時 다시 실행할지 여부를 결정할 수 있음



※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-4/321>

Job is ... (detailed)

- 고정적(fixed)인 완료 횟수 를 가진 병렬 Job
 - .spec.completions 에 0이 아닌 양수 값을 지정
 - . Job은 전체 작업을 나타내며 1에서 .spec.completions 까지의 범위의 각 값에 대해 한 개씩 성공한 Pod가 있으면 완료
- 작업 큐(queue) 가 있는 병렬 Job
 - .spec.completions 를 지정하지 않고, .spec.parallelism 에 양수 값 설정
 - . Job의 모든 Pod가 성공적으로 종료되면, 새로운 Pod는 생성되지 않음
 - . 하나 이상의 Pod가 성공적으로 종료되고, 모든 Pod가 종료되면 Job은 성공적으로 완료



※ 참고 : <https://kubernetes.io/ko/docs/concepts/workloads/controllers/job/>

///

K8s : Job Hands-On

Job YAML

Container에서 command 입력하는 방법과, Job YAML 문법에 대해서 먼저 살펴보자

Job-pi.yaml

```
apiVersion: batch/v1
kind: Job

metadata:
  name: pi

spec:
  template:
    spec:
      containers:
        - name: pi
          image: perl
          command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
          restartPolicy: Never
      backoffLimit: 4
```

→ 파일 값을 계산하는 명령어이며, 평균 10초 정도 소요

▷ 파드 안에서 동작하는 컨테이너를 위한 커맨드와 인자 사용 방법

설명	도커 필드 이름	쿠버네티스 필드 이름
컨테이너에서 실행되는 커맨드	Entrypoint	command
커맨드에 전달되는 인자들	Cmd	arg

▷ restartPolicy (Pod Lifecycle)

- Always : 항상 재시작
- OnFailure : 비정상 종료 발생 時 재시작
- Never : 재시작 하지 않음

▷ BackoffLimit (백오프 정책)

- 실패할 경우 다시 실행시킬 재시도 횟수 지정
- default 값 = 6
- 실패 후 10초, 20초, 40초 간격으로 재시도

※ 참고 : <https://kubernetes.io/ko/docs/concepts/workloads/controllers/job/>

※ 참고 : <https://kubernetes.io/ko/docs/concepts/workloads/pods/pod-lifecycle/>

※ 참고 : <https://kubernetes.io/ko/docs/tasks/inject-data-application/define-command-argument-container/>

Job Create

Job은 앞에서 살펴본 다른 리소스와 차이가 있다.

```
remote > cd advanced-kubernetes/03-week/Job
```

```
remote > kubectl create -f ./job-pi.yaml
```

```
job.batch/pi created
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pi-x94l9	0/1	Completed	0	33s	10.233.103.88	worker2	<none>	<none>

```
remote > kubectl get jobs -o wide
```

NAME	COMPLETIONS	DURATION	AGE	CONTAINERS	IMAGES	SELECTOR
pi	1/1	24s	47s	pi	perl	controller-uid=c4cbbbfc-894b-4a16-97f3-178ed0eab10a

일단 STATUS가 'Completed'이고, 그렇기에 READY 값이 '0/1'로 나오게 된다. 즉, 프로세스가 종료된 것이다.

또한, Jobs의 정보의 형태도 앞에서의 다른 리소스와는 다른 모습을 보인다.

Job Delete

지우는 방식은 계속 유사하다 ...

```
remote > kubectl get jobs -o wide
```

NAME	COMPLETIONS	DURATION	AGE	CONTAINERS	IMAGES	SELECTOR
pi	1/1	24s	7m8s	pi	perl	controller-uid=c4cbbbf8-894b-4a16-97f3-178ed0eab10a

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pi-x94l9	0/1	Completed	0	7m10s	10.233.103.88	worker2	<none>	<none>

```
remote > kubectl delete job pi
```

```
job.batch "pi" deleted
```

```
remote > kubectl get jobs -o wide
```

```
No resources found in default namespace.
```

```
remote > kubectl get pods -o wide
```

```
No resources found in default namespace.
```

///

Problem

특정 작업이 3번 성공
되어야 할 때

throw-dice - 1/2

success / failure 결과가 나오는 container를 하나 골라봤다.

Job-throw-dice.yaml

```
apiVersion: batch/v1
kind: Job

metadata:
  name: throw-dice-job

spec:
  completions: 3
  parallelism: 3
  backoffLimit: 25

  template:
    spec:
      containers:
        - name: math-add
          image: kodekloud/throw-dice

  restartPolicy: Never
```

- ▷ **completions** : 몇 번의 Completed가 나올 때까지 (성공 발생 횟수)
- ▷ **parallelism** : 한 번에 몇 개까지 실행할 것인가
- ▷ **backoffLimit** : 최대 실행 횟수 (에러 발생 횟수)

'6'이 나올 때면 success, 나머지는 failure

※ 참고 : <https://github.com/kodekloudhub/throw-dice>

throw-dice - 2/2

앞에서 설명한 옵션들이 어떻게 작용하는지 잘 살펴보자

remote > cd advanced-kubernetes/03-week/Job

```
remote > kubectl create -f ./job-throw-dice.yaml
```

job.batch/throw-dice-job created

```
remote > kubectl get jobs -o wide -w
```

NAME	COMPLETIONS	DURATION	AGE	CONTAINERS	IMAGES	SELECTOR
throw-dice-job	0/3		0s	math-add	kodekloud/throw-dice	controller-uid=8bca3fe3-d325-4d33-8f21-be14799fd364
throw-dice-job	0/3	0s	0s	math-add	kodekloud/throw-dice	controller-uid=8bca3fe3-d325-4d33-8f21-be14799fd364
throw-dice-job	0/3	5s	5s	math-add	kodekloud/throw-dice	controller-uid=8bca3fe3-d325-4d33-8f21-be14799fd364
throw-dice-job	1/3	5s	5s	math-add	kodekloud/throw-dice	controller-uid=8bca3fe3-d325-4d33-8f21-be14799fd364
throw-dice-job	1/3	7s	7s	math-add	kodekloud/throw-dice	controller-uid=8bca3fe3-d325-4d33-8f21-be14799fd364
throw-dice-job	1/3	15s	15s	math-add	kodekloud/throw-dice	controller-uid=8bca3fe3-d325-4d33-8f21-be14799fd364
...						

```
remote > kubectl get pods -o wide -w
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
throw-dice-job-4ffd2	0/1	Pending	0	0s	<none>	<none>	<none>	<none>
throw-dice-job-4ffd2	0/1	Pending	0	0s	<none>	worker2	<none>	<none>
throw-dice-job-5k7gj	0/1	Pending	0	0s	<none>	<none>	<none>	<none>
throw-dice-job-vmbst	0/1	Pending	0	0s	<none>	<none>	<none>	<none>
throw-dice-job-5k7gj	0/1	Pending	0	0s	<none>	worker2	<none>	<none>
throw-dice-job-vmbst	0/1	Pending	0	0s	<none>	worker1	<none>	<none>
throw-dice-job-vmbst	0/1	ContainerCreating	0	0s	<none>	worker1	<none>	<none>
throw-dice-job-4ffd2	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>
throw-dice-job-5k7gj	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>
throw-dice-job-vmbst	0/1	ContainerCreating	0	1s	<none>	worker1	<none>	<none>

///

Kubernetes

CronJob

CronJob is ...

- 특정 시간 또는 지정된 간격으로 Job을 반복 실행

CRONJOB

- An extension of the Job Controller, it provides a method of executing jobs on a cron-like schedule.
- CronJobs within Kubernetes use **UTC ONLY**.

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: cronjob-example
spec:
  schedule: "*/1 * * * *"
  successfulJobsHistoryLimit: 3
  failedJobsHistoryLimit: 1
  jobTemplate:
    spec:
      completions: 4
      parallelism: 2
      template:
        <pod template>
```

The number of successful jobs to retain

The number of failed jobs to retain

The cron schedule for the job



※ 참고 : <https://www.slideshare.net/RonnyTrommer/devjam-2019-introduction-to-kubernetes>

///

K8s : CronJob Hands-On

CronJob YAML

Job을 주기적으로 실행하기 위한 리소스이다. 한번 해보자!

cj-pi.yaml

```
apiVersion: batch/v1beta1
kind: CronJob

metadata:
  name: batch-job-every-fifteen-minutes

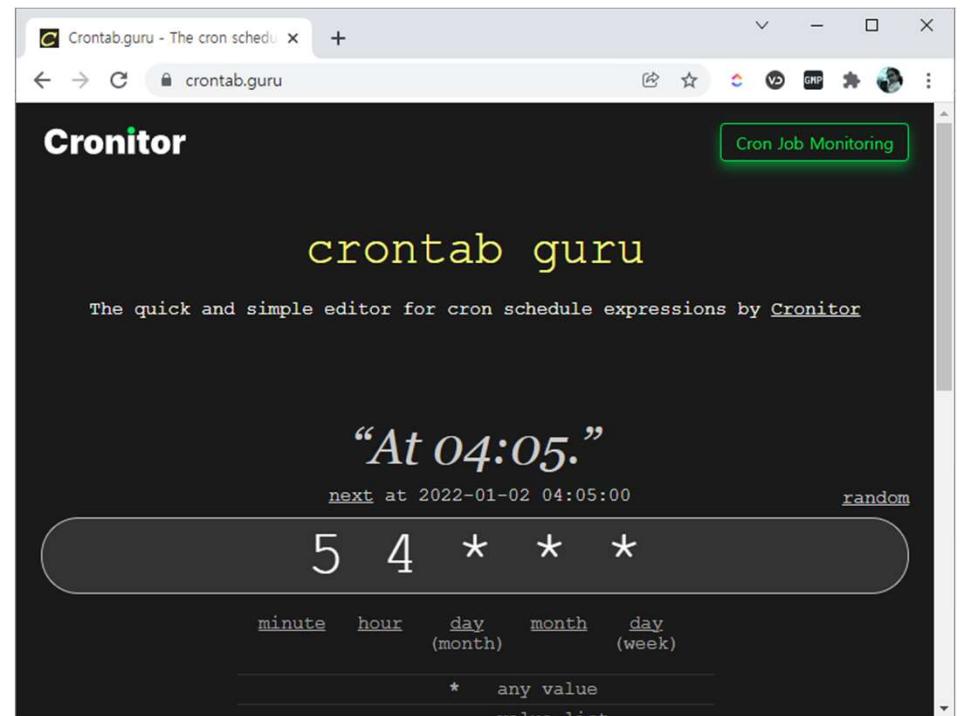
spec:
  schedule: "0,15,30,45 * * * *"
  jobTemplate:
    spec:
      backoffLimit: 3

    template:
      metadata:
        labels:
          app: periodic-batch-job

      spec:
        containers:
        - name: pi
          image: perl
          command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]

  restartPolicy: Never
```

- ▷ schedule 설정은 아래 사이트에서 편하게 확인해볼 수 있다
 - <https://crontab.guru/>



※ 참고 : <https://kubernetes.io/ko/docs/tasks/automated-tasks-with-cron-jobs/>

CronJob Create

CronJob은 앞에서 살펴본 다른 리소스와 결과 확인이 조금 색다르다.

```
remote > cd advanced-kubernetes/03-week/CronJob
```

```
remote > kubectl create -f ./cj-pi.yaml
```

```
cronjob.batch/batch-job-every-fifteen-minutes created
```

```
remote > kubectl get cronjobs -o wide
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE	CONTAINERS	IMAGES	SELECTOR
batch-job-every-fifteen-minutes	0,15,30,45 * * * *	False	0	<none>	9s	pi	perl	<none>

```
remote > kubectl get pods -o wide
```

```
No resources found in default namespace.
```

CronJob Watch

CronJob은 시간이 되면 ACTIVE가 되었다가, 끝나면 다시 InACTIVE가 된다.

```
remote > kubectl get cronjobs -o wide -w
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE	CONTAINERS	IMAGES	SELECTOR
batch-job-every-fifteen-minutes	0,15,30,45 * * * *	False	0	<none>	102s	pi	perl	<none>
batch-job-every-fifteen-minutes	0,15,30,45 * * * *	False	1	1s	5m5s	pi	perl	<none>
batch-job-every-fifteen-minutes	0,15,30,45 * * * *	False	0	11s	5m15s	pi	perl	<none>

CronJob은 시간이 안되었을 때엔, Pod가 아예 보이지 않는다가, 때가 되면 생성되고, 정해진 일을 수행한다.

```
remote > kubectl get pods -o wide -w
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
batch-job-every-fifteen-minutes-1641028500-7jhvv	0/1	Pending	0	0s	<none>	<none>	<none>	<none>
batch-job-every-fifteen-minutes-1641028500-7jhvv	0/1	Pending	0	0s	<none>	worker2	<none>	<none>
batch-job-every-fifteen-minutes-1641028500-7jhvv	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>
batch-job-every-fifteen-minutes-1641028500-7jhvv	0/1	ContainerCreating	0	1s	<none>	worker2	<none>	<none>
batch-job-every-fifteen-minutes-1641028500-7jhvv	1/1	Running	0	5s	10.233.103.90	worker2	<none>	<none>
batch-job-every-fifteen-minutes-1641028500-7jhvv	0/1	Completed	0	8s	10.233.103.90	worker2	<none>	<none>
batch-job-every-fifteen-minutes-1641028500-7jhvv	0/1	Completed	0	8s	10.233.103.90	worker2	<none>	<none>

일을 마친 Pod는 'Completed'인 상태로 누적되어 보인다.

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
batch-job-every-fifteen-minutes-1641028500-7jhvv	0/1	Completed	0	15m	10.233.103.90	worker2	<none>	<none>
batch-job-every-fifteen-minutes-1641029400-qp58w	0/1	Completed	0	22s	10.233.103.91	worker2	<none>	<none>

CronJob Delete

지우는 방식은 계속 유사하다 ...

```
remote > kubectl get cronjobs -o wide
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE	CONTAINERS	IMAGES	SELECTOR
batch-job-every-fifteen-minutes	0,15,30,45 * * * *	False	0	14m	79m	pi	perl	<none>

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
batch-job-every-fifteen-minutes-1641030300-pp2jj	0/1	Completed	0	44m	10.233.103.92	worker2	<none>	<none>
batch-job-every-fifteen-minutes-1641031200-ljgw5	0/1	Completed	0	29m	10.233.103.93	worker2	<none>	<none>
batch-job-every-fifteen-minutes-1641032100-glr6d	0/1	Completed	0	14m	10.233.103.94	worker2	<none>	<none>

```
remote > kubectl delete cronjob batch-job-every-fifteen-minutes
```

```
cronjob.batch "batch-job-every-fifteen-minutes" deleted
```

```
remote > kubectl get cronjobs -o wide
```

```
No resources found in default namespace.
```

```
remote > kubectl get pods -o wide
```

```
No resources found in default namespace.
```

CronJob History - 1/4

cronjob-3-history.yaml

```
apiVersion: batch/v1beta1
kind: CronJob

metadata:
  name: cronjob-3-history

spec:
  schedule: "*/1 * * * *"

  successfulJobsHistoryLimit: 3
  failedJobsHistoryLimit: 1

  jobTemplate:

    spec:
      backoffLimit: 3
      template:
        metadata:
          labels:
            app: periodic-batch-job

        spec:
          containers:
            - name: pi
              image: perl
              command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
          restartPolicy: Never
```

cronjob-0-history.yaml

```
apiVersion: batch/v1beta1
kind: CronJob

metadata:
  name: cronjob-0-history

spec:
  schedule: "*/1 * * * *"

  successfulJobsHistoryLimit: 0
  failedJobsHistoryLimit: 0

  jobTemplate:

    spec:
      backoffLimit: 3
      template:
        metadata:
          labels:
            app: periodic-batch-job

        spec:
          containers:
            - name: pi
              image: perl
              command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
          restartPolicy: Never
```

CronJob History - 2/4

2개의 CronJob을 생성해보자 !

remote > cd advanced-kubernetes/03-week/CronJob

```
remote > kubectl create -f ./cj-pi-1min-3history.yaml
cronjob.batch/cronjob-3-history created
```

```
remote > kubectl create -f ./cj-pi-1min-0history.yaml
cronjob.batch/cronjob-0-history created
```

1분마다 '0-history' & '3-history' cronJob이 실행되었다가 종료된다

```
remote > kubectl get cronjobs -o wide -w
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE	CONTAINERS	IMAGES	SELECTOR
cronjob-3-history	*/1 * * * *	False	0	<none>	0s	pi	perl	<none>
cronjob-0-history	*/1 * * * *	False	0	<none>	0s	pi	perl	<none>
cronjob-0-history	*/1 * * * *	False	1	0s	1s	pi	perl	<none>
cronjob-3-history	*/1 * * * *	False	1	0s	9s	pi	perl	<none>
cronjob-3-history	*/1 * * * *	False	0	10s	19s	pi	perl	<none>
cronjob-0-history	*/1 * * * *	False	0	20s	21s	pi	perl	<none>
cronjob-0-history	*/1 * * * *	False	1	0s	61s	pi	perl	<none>

CronJob History - 3/4

Pod들이 생성되고 완료되고, 삭제(Terminating) 되는 모습을 살펴볼 수 있다.

remote > **kubectl get pods -o wide -w**

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
cronjob-0-history-1641034140-28xxg	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	
cronjob-0-history-1641034140-28xxg	0/1	Pending	0	0s	<none>	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	Pending	0	0s	<none>	worker2	<none>	<none>	
cronjob-0-history-1641034140-28xxg	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>	
cronjob-0-history-1641034140-28xxg	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	1/1	Running	0	5s	10.233.103.96	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	Completed	0	8s	10.233.103.96	worker2	<none>	<none>	
cronjob-0-history-1641034140-28xxg	1/1	Running	0	8s	10.233.103.95	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	Completed	0	8s	10.233.103.96	worker2	<none>	<none>	
cronjob-0-history-1641034140-28xxg	0/1	Completed	0	12s	10.233.103.95	worker2	<none>	<none>	
cronjob-0-history-1641034140-28xxg	0/1	Completed	0	12s	10.233.103.95	worker2	<none>	<none>	
cronjob-0-history-1641034140-28xxg	0/1	Terminating	0	20s	10.233.103.95	worker2	<none>	<none>	
cronjob-0-history-1641034140-28xxg	0/1	Terminating	0	20s	10.233.103.95	worker2	<none>	<none>	
cronjob-0-history-1641034200-zh56f	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	
cronjob-0-history-1641034200-zh56f	0/1	Pending	0	0s	<none>	worker2	<none>	<none>	
cronjob-0-history-1641034200-zh56f	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>	
cronjob-3-history-1641034200-m6df8	0/1	Pending	0	0s	<none>	<none>	<none>	<none>	
cronjob-3-history-1641034200-m6df8	0/1	Pending	0	0s	<none>	worker2	<none>	<none>	
cronjob-3-history-1641034200-m6df8	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>	
cronjob-0-history-1641034200-zh56f	0/1	ContainerCreating	0	1s	<none>	worker2	<none>	<none>	
cronjob-3-history-1641034200-m6df8	0/1	ContainerCreating	0	1s	<none>	worker2	<none>	<none>	
cronjob-0-history-1641034200-zh56f	1/1	Running	0	5s	10.233.103.97	worker2	<none>	<none>	

^C%

CronJob History - 4/4

pods 현황을 계속 찍어보면, history를 몇 개씩 유지하고 있는지 살펴볼 수 있다.

remote > `kubectl get pods -o wide`

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
cronjob-0-history-1641034260-z8s2h	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	Completed	0	2m	10.233.103.96	worker2	<none>	<none>	
cronjob-3-history-1641034200-m6df8	0/1	Completed	0	60s	10.233.103.98	worker2	<none>	<none>	
cronjob-3-history-1641034260-l5hht	0/1	ContainerCreating	0	0s	<none>	worker2	<none>	<none>	
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
cronjob-0-history-1641034260-z8s2h	1/1	Running	0	8s	10.233.103.100	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	Completed	0	2m8s	10.233.103.96	worker2	<none>	<none>	
cronjob-3-history-1641034200-m6df8	0/1	Completed	0	68s	10.233.103.98	worker2	<none>	<none>	
cronjob-3-history-1641034260-l5hht	0/1	Completed	0	8s	10.233.103.99	worker2	<none>	<none>	
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
cronjob-0-history-1641034260-z8s2h	0/1	Completed	0	18s	10.233.103.100	worker2	<none>	<none>	
cronjob-3-history-1641034140-nmdfp	0/1	Completed	0	2m18s	10.233.103.96	worker2	<none>	<none>	
cronjob-3-history-1641034200-m6df8	0/1	Completed	0	78s	10.233.103.98	worker2	<none>	<none>	
cronjob-3-history-1641034260-l5hht	0/1	Completed	0	18s	10.233.103.99	worker2	<none>	<none>	
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
cronjob-3-history-1641034140-nmdfp	0/1	Completed	0	2m54s	10.233.103.96	worker2	<none>	<none>	
cronjob-3-history-1641034200-m6df8	0/1	Completed	0	114s	10.233.103.98	worker2	<none>	<none>	
cronjob-3-history-1641034260-l5hht	0/1	Completed	0	54s	10.233.103.99	worker2	<none>	<none>	
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
cronjob-3-history-1641034200-m6df8	0/1	Completed	0	2m34s	10.233.103.98	worker2	<none>	<none>	
cronjob-3-history-1641034260-l5hht	0/1	Completed	0	94s	10.233.103.99	worker2	<none>	<none>	
cronjob-3-history-1641034320-v4rkn	0/1	Completed	0	34s	10.233.103.102	worker2	<none>	<none>	

///

Break

///

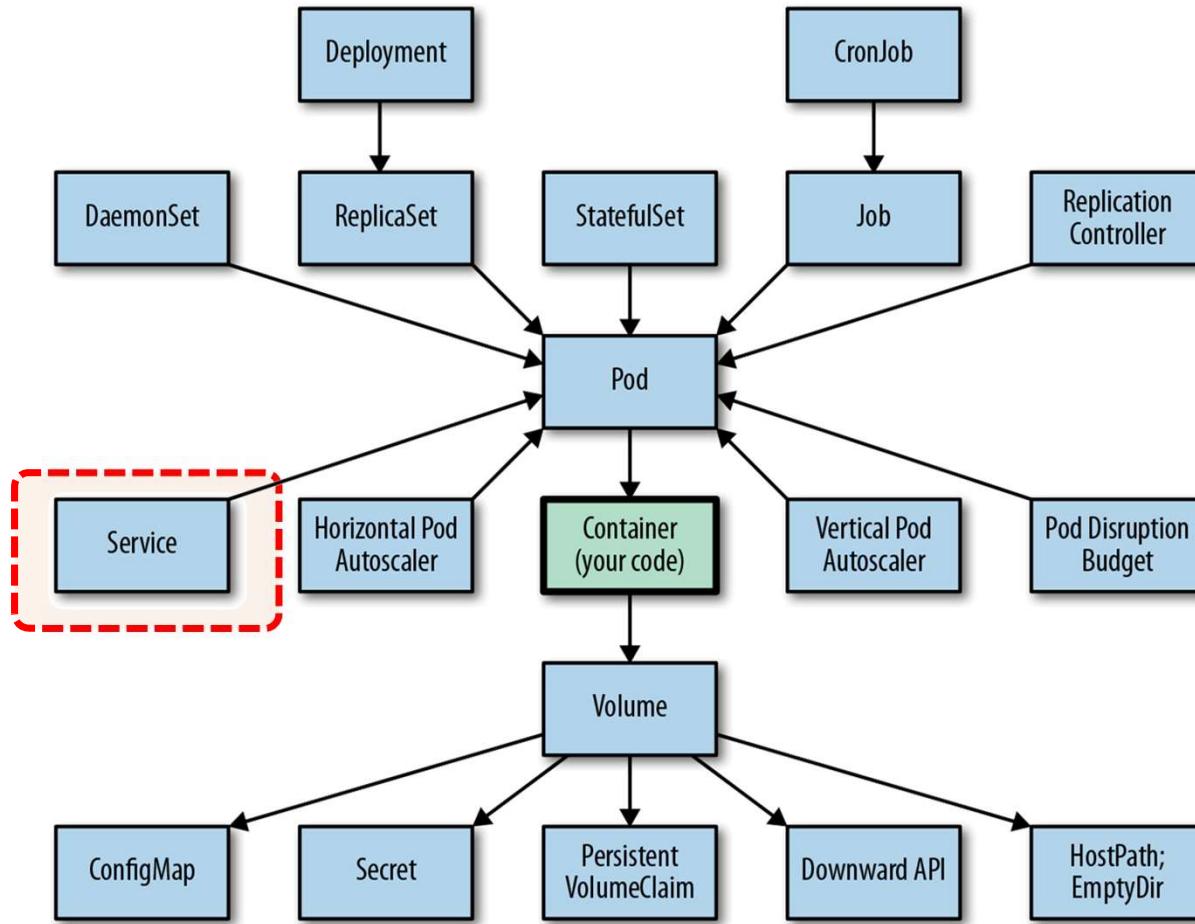
Flip Learning

(Service - ClusterIP/NodePort/ExternalName)

이혜정님

///

Service



※ 참고 : <https://www.oreilly.com/library/view/kubernetes-patterns/9781492050278/ch01.html>

Kubernetes

Service

너무나 정리가 잘되어 있는 자료

▷ Kubernetes 네트워크 이해하기 (1) : 컨테이너 네트워크부터 CNI까지

- <https://hyojun.me/~k8s-network-1>

▷ Kubernetes 네트워크 이해하기 (2) : 서비스 개념과 동작 원리

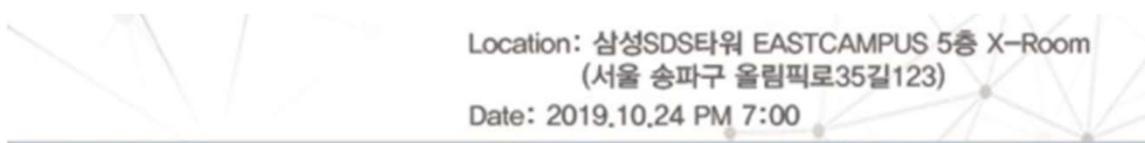
- <https://hyojun.me/~k8s-network-2>

※ 참고 : <https://www.facebook.com/groups/k8skr/permalink/3011677939113859/>

IT 인프라 엔지니어 그룹 맛업

1st IT Infra Engineer Group Meet Up

- <https://festa.io/events/592>

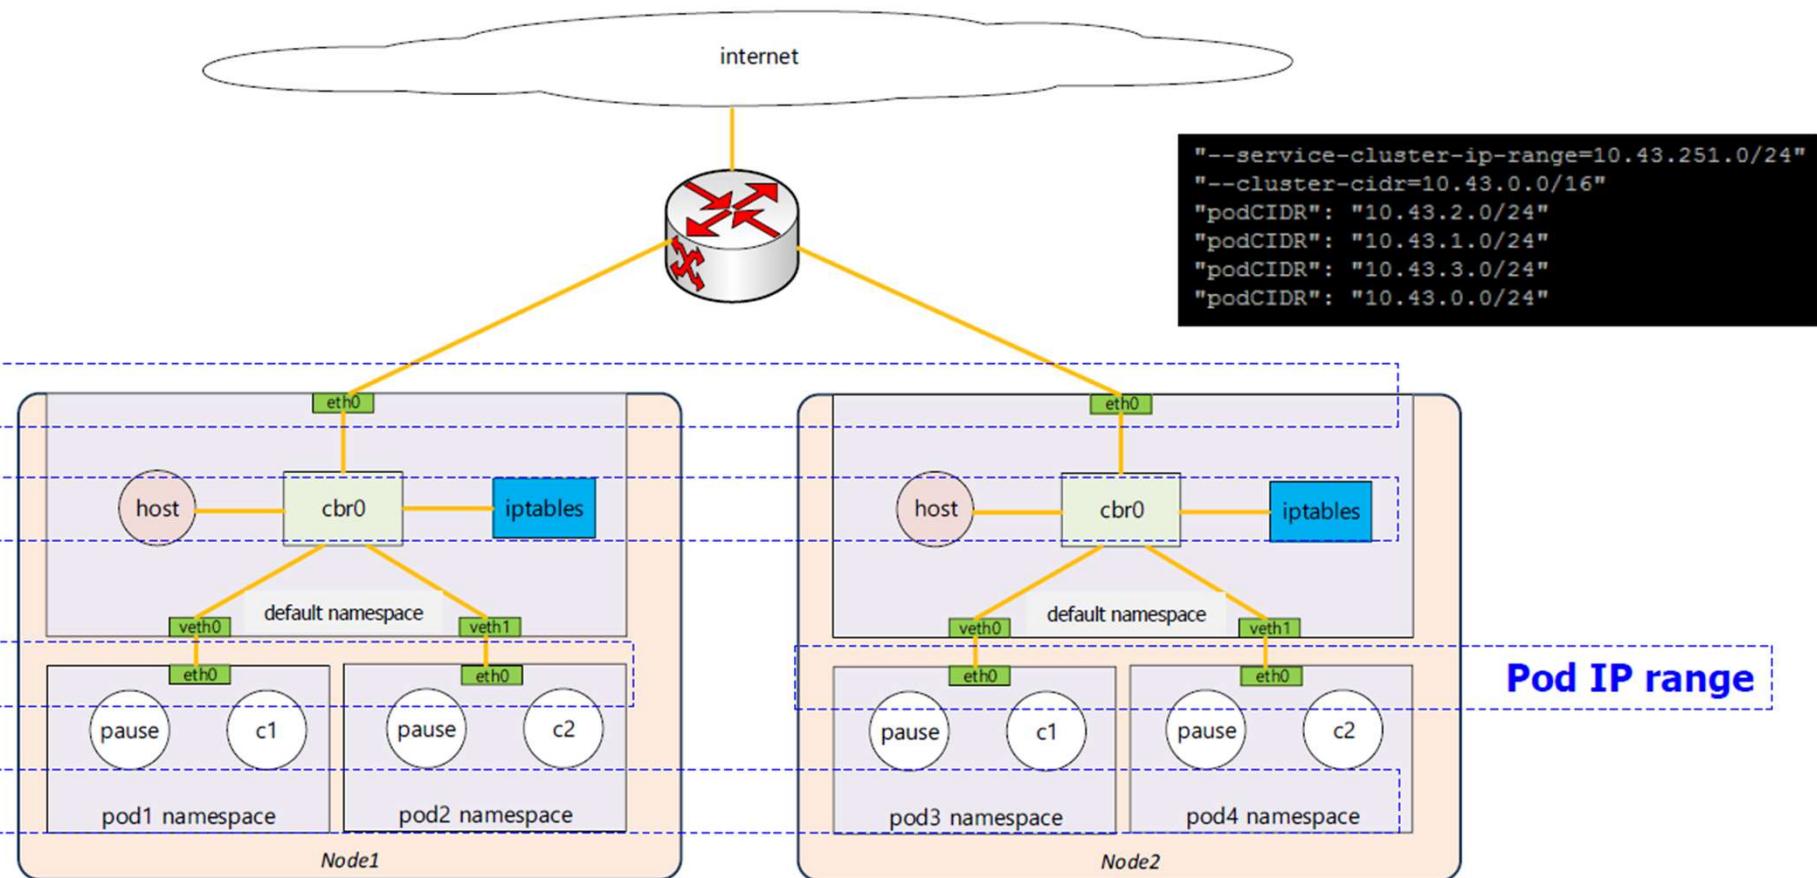


When	Agenda	Speaker
19:00~19:45	자동화를 왜 해야하나요? " Why do I have to do the automation?" Session1: (자동화가 구성하는게 더 시간이 걸리지 않나요? 불편하지 않나요? 등에 대하여 유의미한 경험을 공유하고자 합니다.)	조훈(우아하게 앤서블 저자)
20:00~20:45	Session2: "오픈소스를 활용한 netflow/sflow 수집/시각화"	김용대 (홀스홀리커)
21:00~21:45	Session3: "쿠버네티스 네트워킹(Paket Flow in K8S) (쿠버네티스의 패킷흐름을 네트워크 엔지니어의 관점에서 살펴보기)"	임창식 (오리엔티)

[발표자료] : <https://www.slideshare.net/InfraEngineer/meetup2nd-v12>

K8S Network IP Ranges

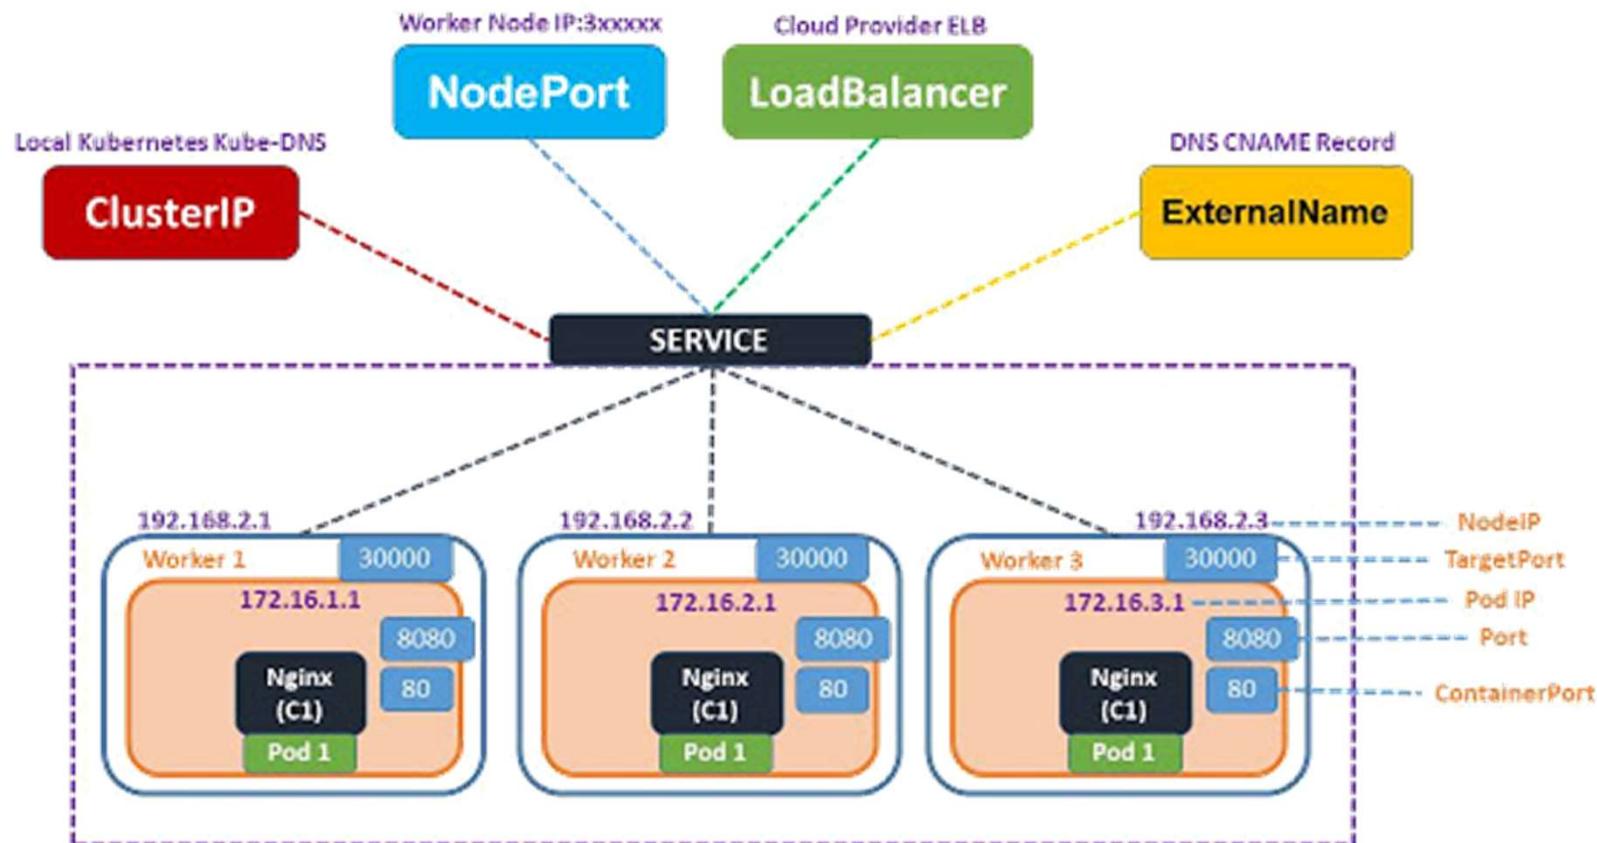
Kubernetes 안에서는 여러 계층의 Network Layer가 존재한다.



※ 참고 : <https://www.slideshare.net/InfraEngineer/ss-186475759>

Service is ...

- 동일한 서비스를 제공하는 Pod 그룹에 지속적인 단일 접점을 만들려고 할 때 생성하는 리소스



※ 참고 : <https://www.learnitguide.net/2020/05/kubernetes-services-explained-examples.html>

Keywords

▷ ClusterIP

- 디폴트 설정으로, 서비스에 클러스터 IP (내부 IP)를 할당한다.
- 쿠버네티스 클러스터 내에서는 이 서비스에 접근이 가능하지만, 클러스터 외부에서는 외부 IP 를 할당 받지 못했기 때문에 접근이 불가능하다.

▷ NodePort

- 클러스터 IP로만 접근이 가능한 것이 아니라, 모든 노드의 IP와 포트를 통해서도 접근이 가능하게 된다.
- 예를 들어 hello-node-svc 라는 서비스를 NodePort 타입으로 선언을 하고 nodePort를 30036으로 설정하면,
- 설정에 따라 클러스터 IP의 80포트로도 접근이 가능하지만, 모든 노드의 30036 포트로도 서비스를 접근할 수 있다.

▷ Load Balancer

- 보통 클라우드 벤더에서 제공하는 설정 방식으로, 외부 IP 를 가지고 있는 LoadBalancer를 할당한다.
- 외부 IP를 가지고 있기 때문에, 클러스터 외부에서 접근이 가능하다.

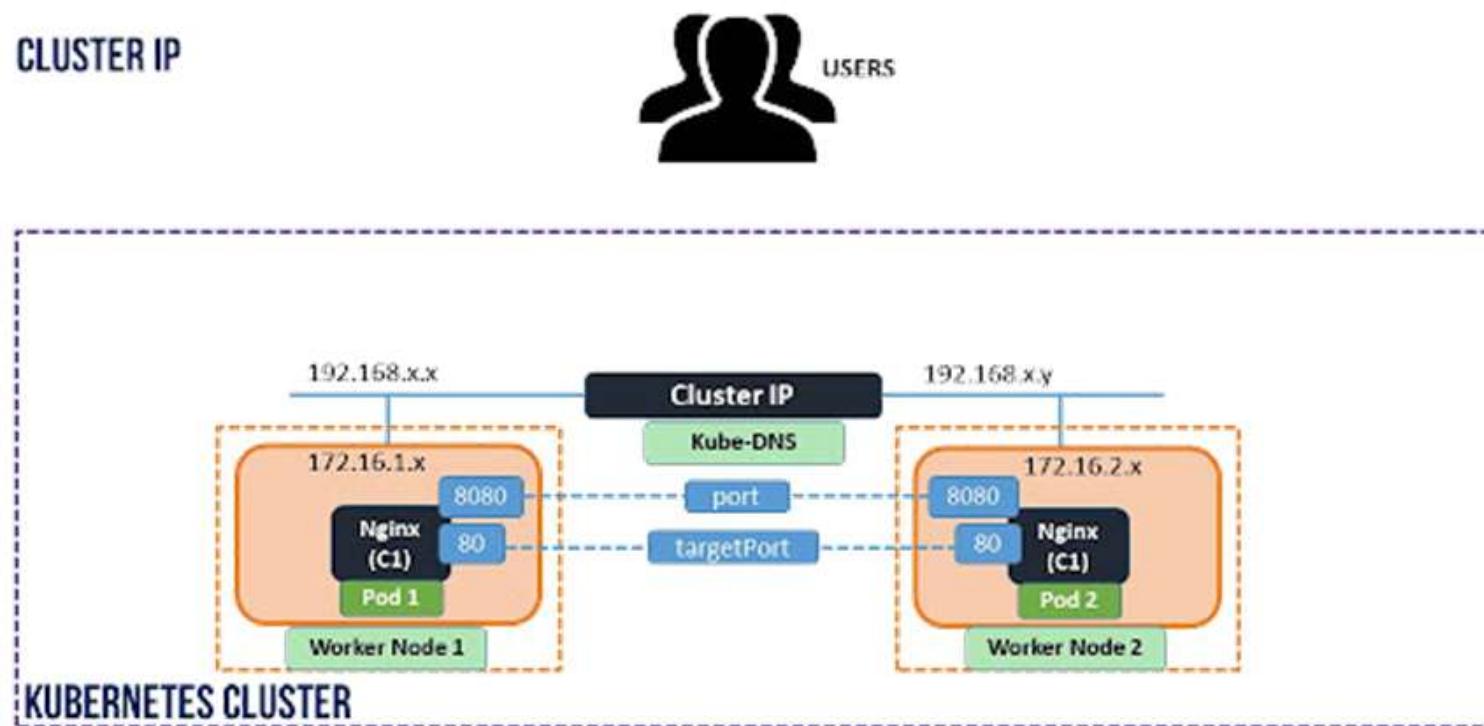
▷ ExternalName

- 외부 서비스를 쿠버네티스 내부에서 호출하고자 할 때 사용할 수 있다.
- Pod들은 클러스터 IP를 가지고 있기 때문에 클러스터 IP 대역 밖의 서비스를 호출하고자 하면, NAT 설정 等 복잡한 설정이 필요하다. 이 때 사용하면 된다.

※ 참고 : <https://bcho.tistory.com/1262>

ClusterIP

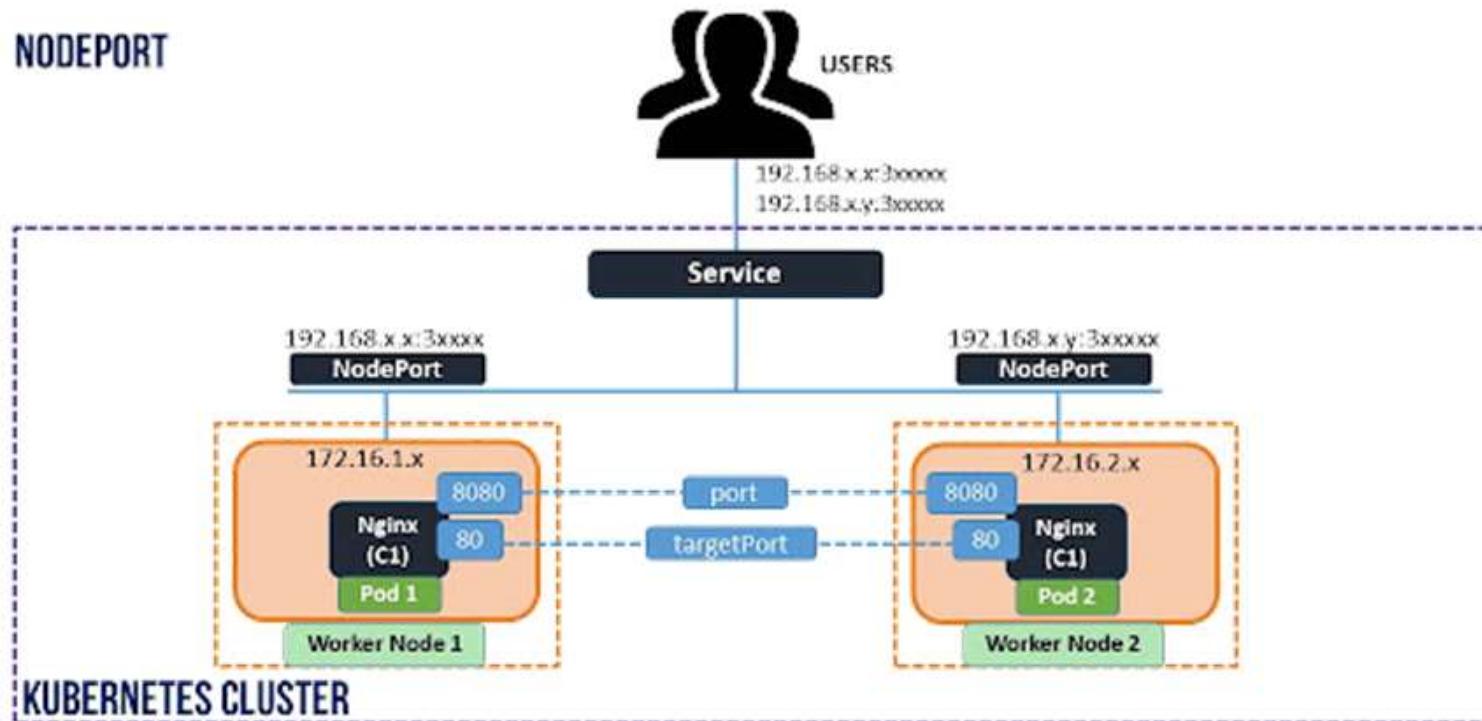
- 디폴트 설정으로, 서비스에 클러스터 IP (내부 IP)를 할당한다.
- 클러스터 내에서는 이 서비스에 접근이 가능하지만, 클러스터 외부에서는 외부 IP 를 할당 받지 못했기 때문에 접근이 불가능하다.



※ 참고 : <https://www.learnitguide.net/2020/05/kubernetes-services-explained-examples.html>

NodePort

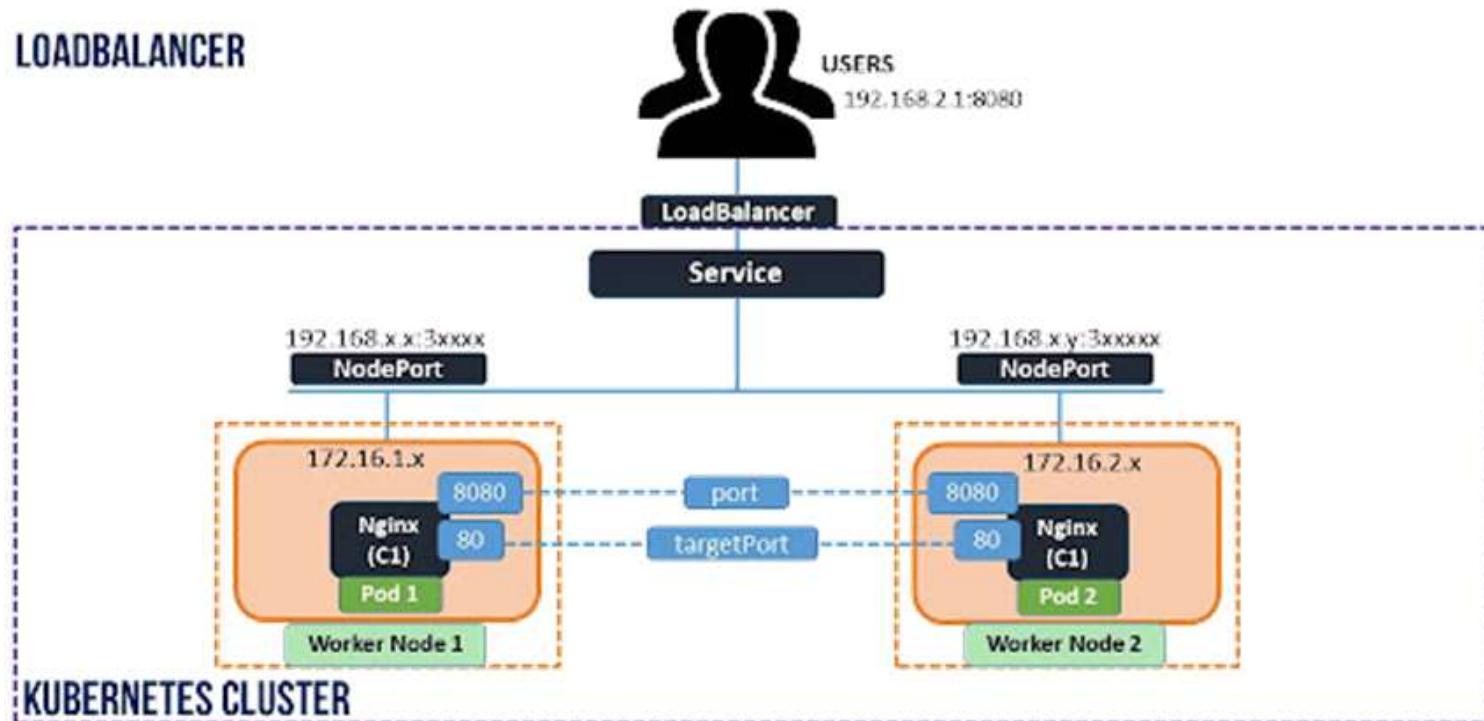
- 클러스터 IP로만 접근이 가능한 것이 아니라, 모든 노드의 IP와 포트를 통해서도 접근이 가능하게 된다.
- 예를 들어 hello-node-svc라는 서비스를 NodePort 타입으로 선언을 하고 nodePort를 30036으로 설정하면,
- 설정에 따라 클러스터 IP의 80포트로도 접근이 가능하지만, 모든 노드의 30036 포트로도 서비스를 접근할 수 있다.



※ 참고 : <https://www.learnitguide.net/2020/05/kubernetes-services-explained-examples.html>

LoadBalancer

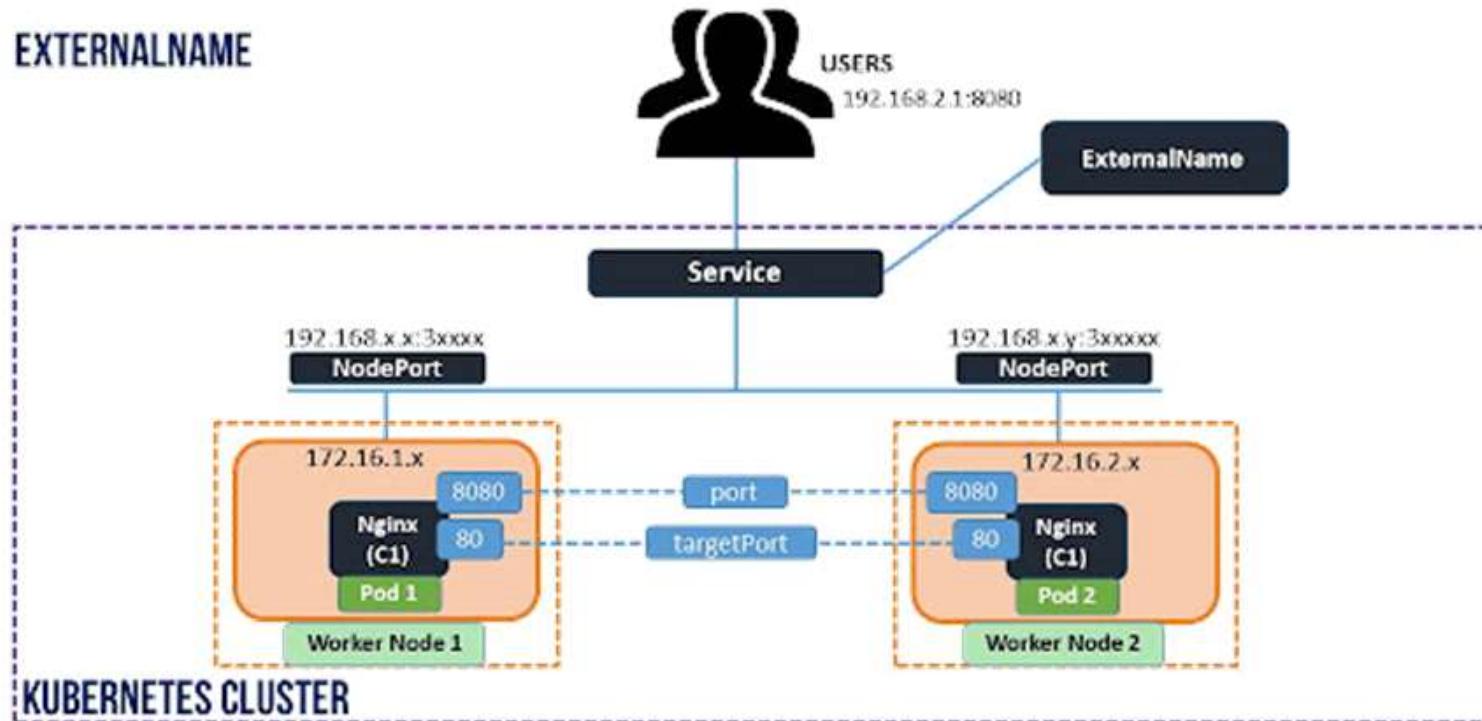
- 보통 클라우드 벤더에서 제공하는 설정 방식으로, 외부 IP 를 가지고 있는 LoadBalancer를 할당한다.
- 외부 IP를 가지고 있기 때문에, 클러스터 외부에서 접근이 가능하다.



※ 참고 : <https://www.learnitguide.net/2020/05/kubernetes-services-explained-examples.html>

ExternalName

- 외부 서비스를 쿠버네티스 내부에서 호출하고자 할 때 사용할 수 있다.
- Pod들은 클러스터 IP를 가지고 있기 때문에 클러스터 IP 대역 밖의 서비스를 호출하고자 할 때 사용하면 된다.



※ 참고 : <https://www.learnitguide.net/2020/05/kubernetes-services-explained-examples.html>

///

K8s : Service (ClusterIP / NodePort / ExternalName) Hands-On

ReplicaSet Create

앞에서 공부한 ReplicaSet을 이용해서 Kubernetes Network에 대해서 알아보겠다.

rs-node-web.yaml

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-labels

spec:
  replicas: 3

  selector:
    matchLabels:
      app: node-web

  template:
    metadata:
      labels:
        app: node-web
    spec:
      containers:
      - name: node-web
        image: whatwant/node-web:1.0
      ports:
      - containerPort: 8080
```

```
remote > cd advanced-kubernetes/03-week/Service
```

```
remote > kubectl create -f ./rs-node-web.yaml
```

```
replicaset.apps/rs-labels created
```

```
remote > kubectl get pods -o wide
```

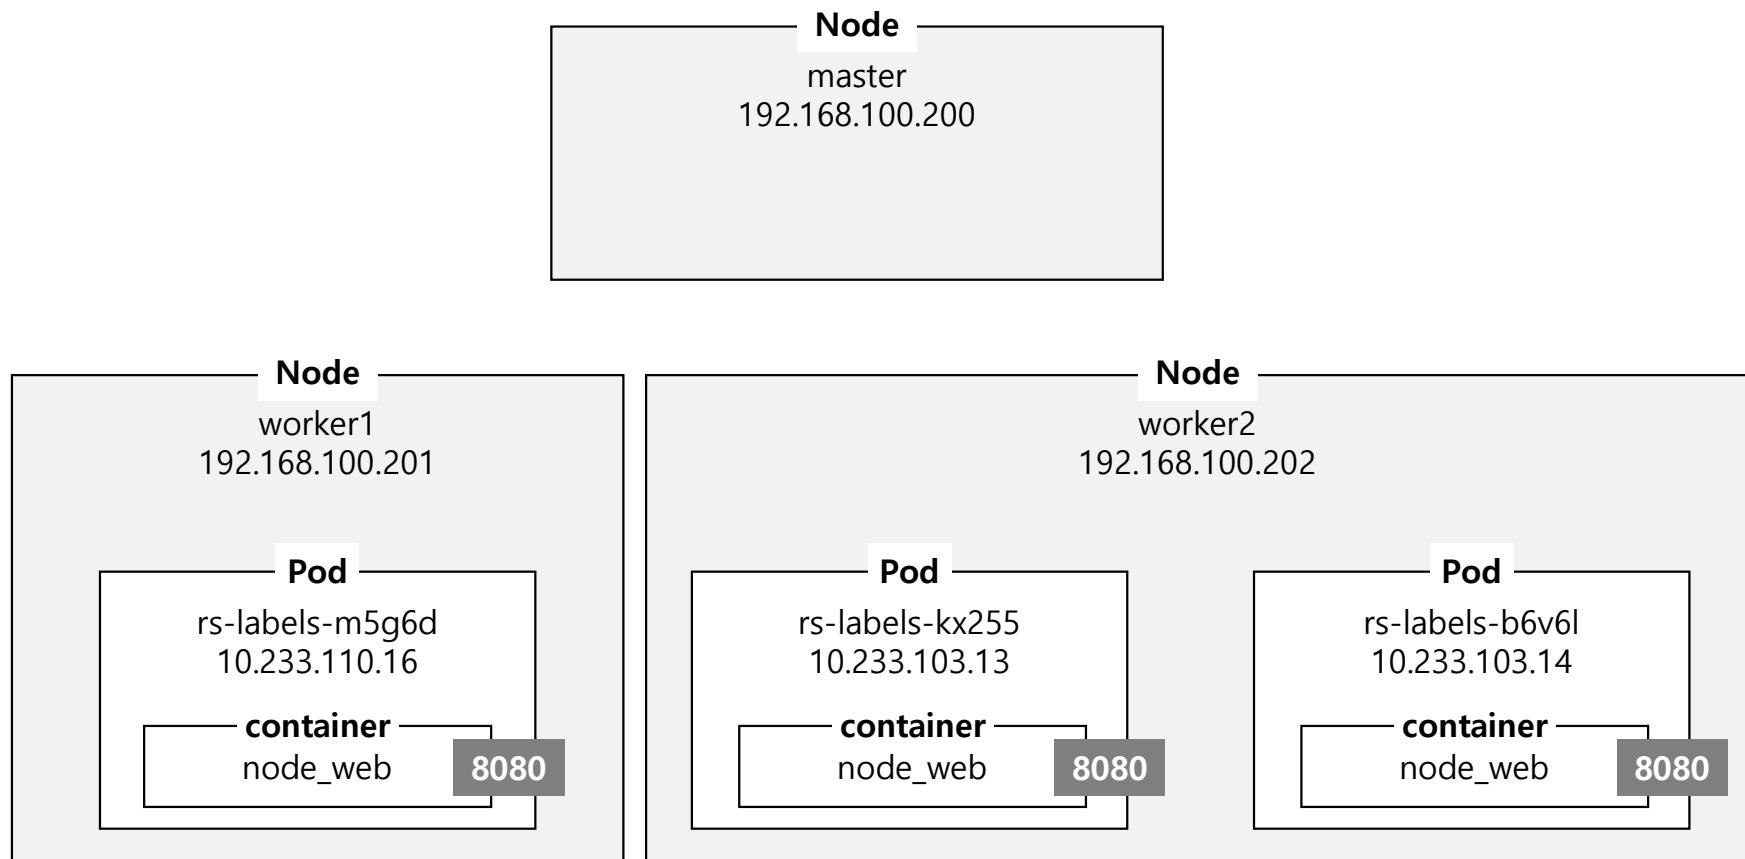
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-b6v6l	1/1	Running	0	85s	10.233.103.14	worker2	<none>	<none>
rs-labels-kx255	1/1	Running	0	85s	10.233.103.13	worker2	<none>	<none>
rs-labels-m5g6d	1/1	Running	0	85s	10.233.110.16	worker1	<none>	<none>

```
remote > kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME
master	Ready	control-plane,master	113d	v1.20.7	192.168.100.200	<none>	Ubuntu 20.04.3 LTS	5.4.0-84-generic	docker://19.3.15
worker1	Ready	<none>	113d	v1.20.7	192.168.100.201	<none>	Ubuntu 20.04.3 LTS	5.4.0-84-generic	docker://19.3.15
worker2	Ready	<none>	113d	v1.20.7	192.168.100.202	<none>	Ubuntu 20.04.3 LTS	5.4.0-84-generic	docker://19.3.15

Node와 Pod의 IP 정보를 살펴보자.

Now Status



Connect Pods

master node에서 Pod를 접근할 수 있을까?

```
remote > ssh vagrant@192.168.100.200
```

```
master > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-b6v6l	1/1	Running	0	85s	10.233.103.14	worker2	<none>	<none>
rs-labels-kx255	1/1	Running	0	85s	10.233.103.13	worker2	<none>	<none>
rs-labels-m5g6d	1/1	Running	0	85s	10.233.110.16	worker1	<none>	<none>

```
master > curl http://10.233.110.16:8080
```

```
You've hit rs-labels-m5g6d
```

K8s Node가 아닌 외부에서는 Pod를 접근할 수 있을까?

```
master > exit
```

```
remote > curl http://10.233.110.16:8080
```

```
^C
```

당연히 외부에서는 접근할 수가 없다.

Service Create

우리가 원하는 것은 하나의 접점으로 Pod들 중 하나에 접근할 수 있도록 하는 것이다.

svc-node-web.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node

spec:
  ports:
    - port: 80
      targetPort: 8080

  selector:
    app: node-web
```

rs-node-web.yaml

```
apiVersion: apps/v1
kind: ReplicaSet
...
template:
  metadata:
    labels:
      app: node-w
```

remote > cd advanced-kubernetes/03-week/Service

```
remote > kubectl create -f ./svc-node-web.yaml
```

service/svc-node created

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	113d	<none>
svc-node	ClusterIP	10.233.24.252	<none>	80/TCP	12s	app=node-web

remote > ssh vagrant@192.168.100.200

master > curl http://10.233.24.252

You've hit rs-labels-kx255

master > curl http://10.233.24.252

You've hit rs-labels-b6v6l

Restart Pods

ReplicaSet과 Service에 대해서 가장 잘 설명할 방법은, Pods를 죽여버리는 것이다

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-b6v6l	1/1	Running	0	85m	10.233.103.14	worker2	<none>	<none>
rs-labels-kx255	1/1	Running	0	85m	10.233.103.13	worker2	<none>	<none>
rs-labels-m5g6d	1/1	Running	0	85m	10.233.110.16	worker1	<none>	<none>

```
master > kubectl delete pods rs-labels-b6v6l
```

```
master > kubectl delete pods rs-labels-kx255
```

```
master > kubectl delete pods rs-labels-m5g6d
```

```
master > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-98pvb	1/1	Running	0	5m38s	10.233.103.15	worker2	<none>	<none>
rs-labels-r52ch	1/1	Running	0	6m18s	10.233.110.18	worker1	<none>	<none>
rs-labels-tws2l	1/1	Running	0	7m8s	10.233.110.17	worker1	<none>	<none>

```
master > ssh vagrant@192.168.100.200
```

```
remote > curl http://10.233.24.252
```

You've hit rs-labels-tws2l

```
remote > curl http://10.233.24.252
```

You've hit rs-labels-98pvb

Service Type - ClusterIP

우리가 앞에서 작성한 Service에 type을 지정하지 않았었다. 그러면 기본값은?

svc-node-web.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: node-web
```

remote > kubectl get services -o wide

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	114d	<none>
svc-node	ClusterIP	10.233.24.252	<none>	80/TCP	5h29m	app=node-web

remote > kubectl describe service svc-node

```
Name:           svc-node
Namespace:      default
Labels:         <none>
Annotations:   <none>
Selector:       app=node-web
Type:          ClusterIP
IP Families:   <none>
IP:            10.233.24.252
IPs:           10.233.24.252
Port:          <unset>  80/TCP
TargetPort:    8080/TCP
Endpoints:     10.233.103.15:8080,10.233.110.17:8080,10.233.110.18:8080
Session Affinity: None
Events:        <none>
```

그렇다. 특별히 지정하지 않으면, 기본값으로 ClusterIP type으로 설정이 된다.

그리고, Endpoints가 어떻게 지정되고 있는지도 살펴보면 좋다~!

Service Update - NodePort

앞에서 선언한 Service를 업데이트 하는 방식으로 진행해보자.

svc-node-web-nodeport.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node

spec:
  type: NodePort
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30123

  selector:
    app: node-web
```

```
remote > cd advanced-kubernetes/04-week/Service
remote > kubectl apply -f ./svc-node-web-nodeport.yaml
```

```
Warning: resource services/svc-node is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.
service/svc-node configured
```

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	114d	<none>
svc-node	NodePort	10.233.24.252	<none>	80:30123/TCP	5h45m	app=node-web

```
remote > curl http://192.168.100.200:30123
```

```
You've hit rs-labels-r52ch
```

```
remote > curl http://192.168.100.201:30123
```

```
You've hit rs-labels-98pzb
```

```
remote > curl http://192.168.100.202:30123
```

```
You've hit rs-labels-98pzb
```

NodePort를 이용해
Node 밖에서도
접근할 수 있음을 볼 수 있다!

Service Environments

Pod에서 Service 설정 값을 확인해볼 수 있다.

※ Pod가 생성될 시점에서의 Service 정보가 반영된다.
즉, Pod가 생성된 이후의 Service 정보는 모른다.

remote > kubectl get services

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	114d
svc-node	NodePort	10.233.24.252	<none>	80:30123/TCP	6h27m

remote > kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
rs-labels-98pzb	1/1	Running	0	6h1m
rs-labels-r52ch	1/1	Running	0	6h2m
rs-labels-tws2l	1/1	Running	0	6h3m

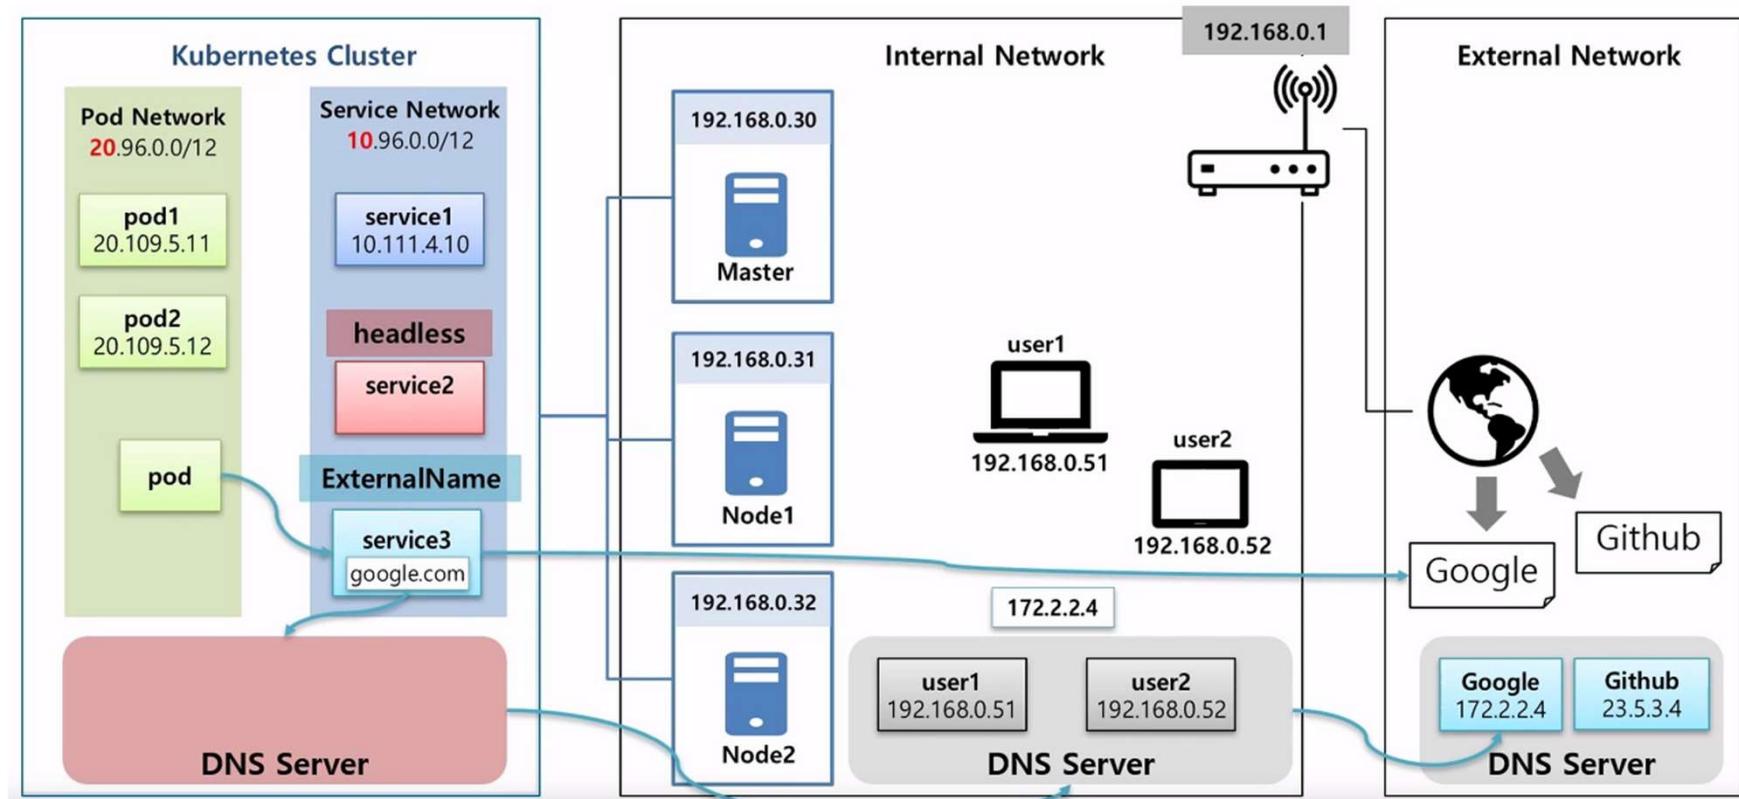
remote > kubectl exec -it rs-labels-tws2l -- env

```
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=rs-labels-tws2l
TERM=xterm
SVC_NODE_SERVICE_HOST=10.233.24.252
SVC_NODE_PORT=tcp://10.233.24.252:80
SVC_NODE_PORT_80_TCP=tcp://10.233.24.252:80
SVC_NODE_PORT_80_TCP_PROTO=tcp
SVC_NODE_PORT_80_TCP_ADDR=10.233.24.252
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_PORT_443_TCP=tcp://10.233.0.1:443
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_SERVICE_PORT=443
...
...
```

```
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=rs-labels-tws2l
TERM=xterm
SVC_NODE_SERVICE_HOST=10.233.24.252
SVC_NODE_PORT=tcp://10.233.24.252:80
SVC_NODE_PORT_80_TCP=tcp://10.233.24.252:80
SVC_NODE_PORT_80_TCP_PROTO=tcp
SVC_NODE_PORT_80_TCP_ADDR=10.233.24.252
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_PORT_443_TCP=tcp://10.233.0.1:443
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_SERVICE_PORT=443
SVC_NODE_SERVICE_PORT=80
SVC_NODE_PORT_80_TCP_PORT=80
KUBERNETES_SERVICE_HOST=10.233.0.1
KUBERNETES_PORT=tcp://10.233.0.1:443
KUBERNETES_PORT_443_TCP_ADDR=10.233.0.1
NODE_VERSION=15.14.0
YARN_VERSION=1.22.5
HOME=/root
```

ExternalName - Recap 1/3

- ExternalName service3에 google 도메인 이름을 넣으면 DNS를 타고 외부 Google 아이피를 가져올 수 있음
- 주후 google을 github으로 service3의 ExternalName만 변경하면 google이 아닌 github에서 가져오게 할 수 있음



※ 참고 : <https://junior-developer.tistory.com/74>

ExternalName - Recap 2/3

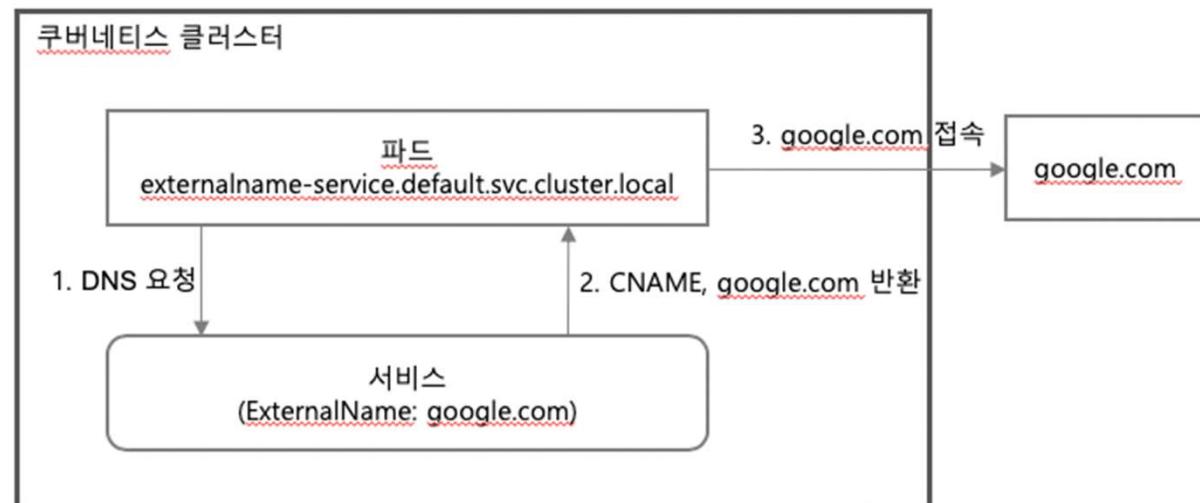
- 외부 서비스의 별칭으로 사용되는 서비스
- . Pod는 서비스의 FQDN을 사용하는 대신 external-service.default.svc.cluster.local로 외부 서비스에 연결
- . Service를 사용하는 Pod에서 실제 서비스 이름과 위치가 숨겨짐
- . 나중에 Service spec 수정하여 다른 Service를 가리킬 수 있음

```
external-name.yaml
```

```
apiVersion: v1
kind: Service

metadata:
  name: external-service

spec:
  type: ExternalName
  externalName: google.com
```



※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-5/144>

ExternalName - Recap 3/3

- 일반적인 selector에 대한 Service가 아닌, DNS 이름에 대한 Service에 맵!
- ExternalName은 IPv4 주소 문자열 및 숫자로 구성된 DNS 이름도 허용
 - . But IPv4 주소와 유사한 ExternalName은 CoreDNS 또는 Ingress-nginx에 의해 확인되지 않음
 - . ExternalName은 정식(canonical) DNS 이름을 지정하기 때문 (Proxy, Forwarding이 아닌 DNS라는 점이 중요!!!)
 - . IP 주소를 하드 코딩하려면, 헤드리스(headless) 서비스 사용 권장
- HTTP 및 HTTPS를 포함한, 몇몇 일반적인 프로토콜에 ExternalName을 사용하는 것은 문제점 존재
 - . 클러스터 내부의 클라이언트가 사용하는 호스트 이름(hostname)이 ExternalName이 참조하는 이름과 다름
 - . 호스트 이름을 사용하는 프로토콜의 경우, 이러한 차이로 인해 오류가 발생하거나 예기치 않은 응답이 발생할 수 있음
 - . HTTP 요청에는 오리진(origin) 서버가 인식하지 못하는 'Host:' 헤더가 존재
 - . TLS 서버는 클라이언트가 연결된 호스트 이름과 일치하는 인증서를 제공할 수 없음

※ 참고 : <https://kubernetes.io/ko/docs/concepts/services-networking/service/#externalname>

///

kubernetes - THE DOCUMENTARY



※ 참고 : <https://www.youtube.com/watch?v=BE77h7dmoQU&t=5s>

실습 환경 : <https://labs.play-with-k8s.com/>

- 여러 개의 instance 생성 가능, but 4시간 무료 사용
- disk 관련된 제약 等 불편함은 존재

The image displays two side-by-side browser windows from the website <https://labs.play-with-k8s.com/>.

The left window, titled "Play with Kubernetes", features a large blue octagonal icon with a white steering wheel in the center. Below the icon, the text "Play with Kubernetes" is prominently displayed in a large, bold, black font. Underneath that, a smaller line of text reads "A simple, interactive and fun playground to learn Kubernetes". A green "Start" button is located at the bottom left.

The right window, titled "Docker Playground", shows a timer at "03:59:45" in large blue digits. Below the timer is a red button labeled "CLOSE SESSION". To the right of the timer, there is a section with the text "Add instances to your playground." and "Sessions and all their instances are deleted after 03:59:45 hours." At the bottom left of this window, there is a link "+ ADD NEW INSTANCE".

///

<https://kahoot.it/>

[Score]

이민준 (4)

이혜정 (3)

김상호 (2)

정현찬 (1)

김정은 (1)

자습 (복습)

- ▷ ReplicaSet / DaemonSet / Job / CronJob
- ▷ Service - ClusterIP / NodePort / ExternalName

///