

4th
Week

네번째 뵙겠습니다 ?!

▷ 출석 체크도 한 번 해보시면 어떠세요?!

- <https://modulabs.co.kr/>
- 모두연 홈페이지 → 로그인 → 마이페이지 → 참여한 랩·풀잎 → 자세히 보기 → 내 풀잎스쿨 출석 확인하기

▷ 지난 주 갑작스럽게 휴강(?)하게 되어 죄송합니다.

- 4 week : 02월 12일
- 5 week : 02월 19일
- 6 week : 02월 26일
- 7 week : 03월 05일
- 8 week : 03월 12일

잡담 &
지난 수업 관련 이야기

///

Agenda

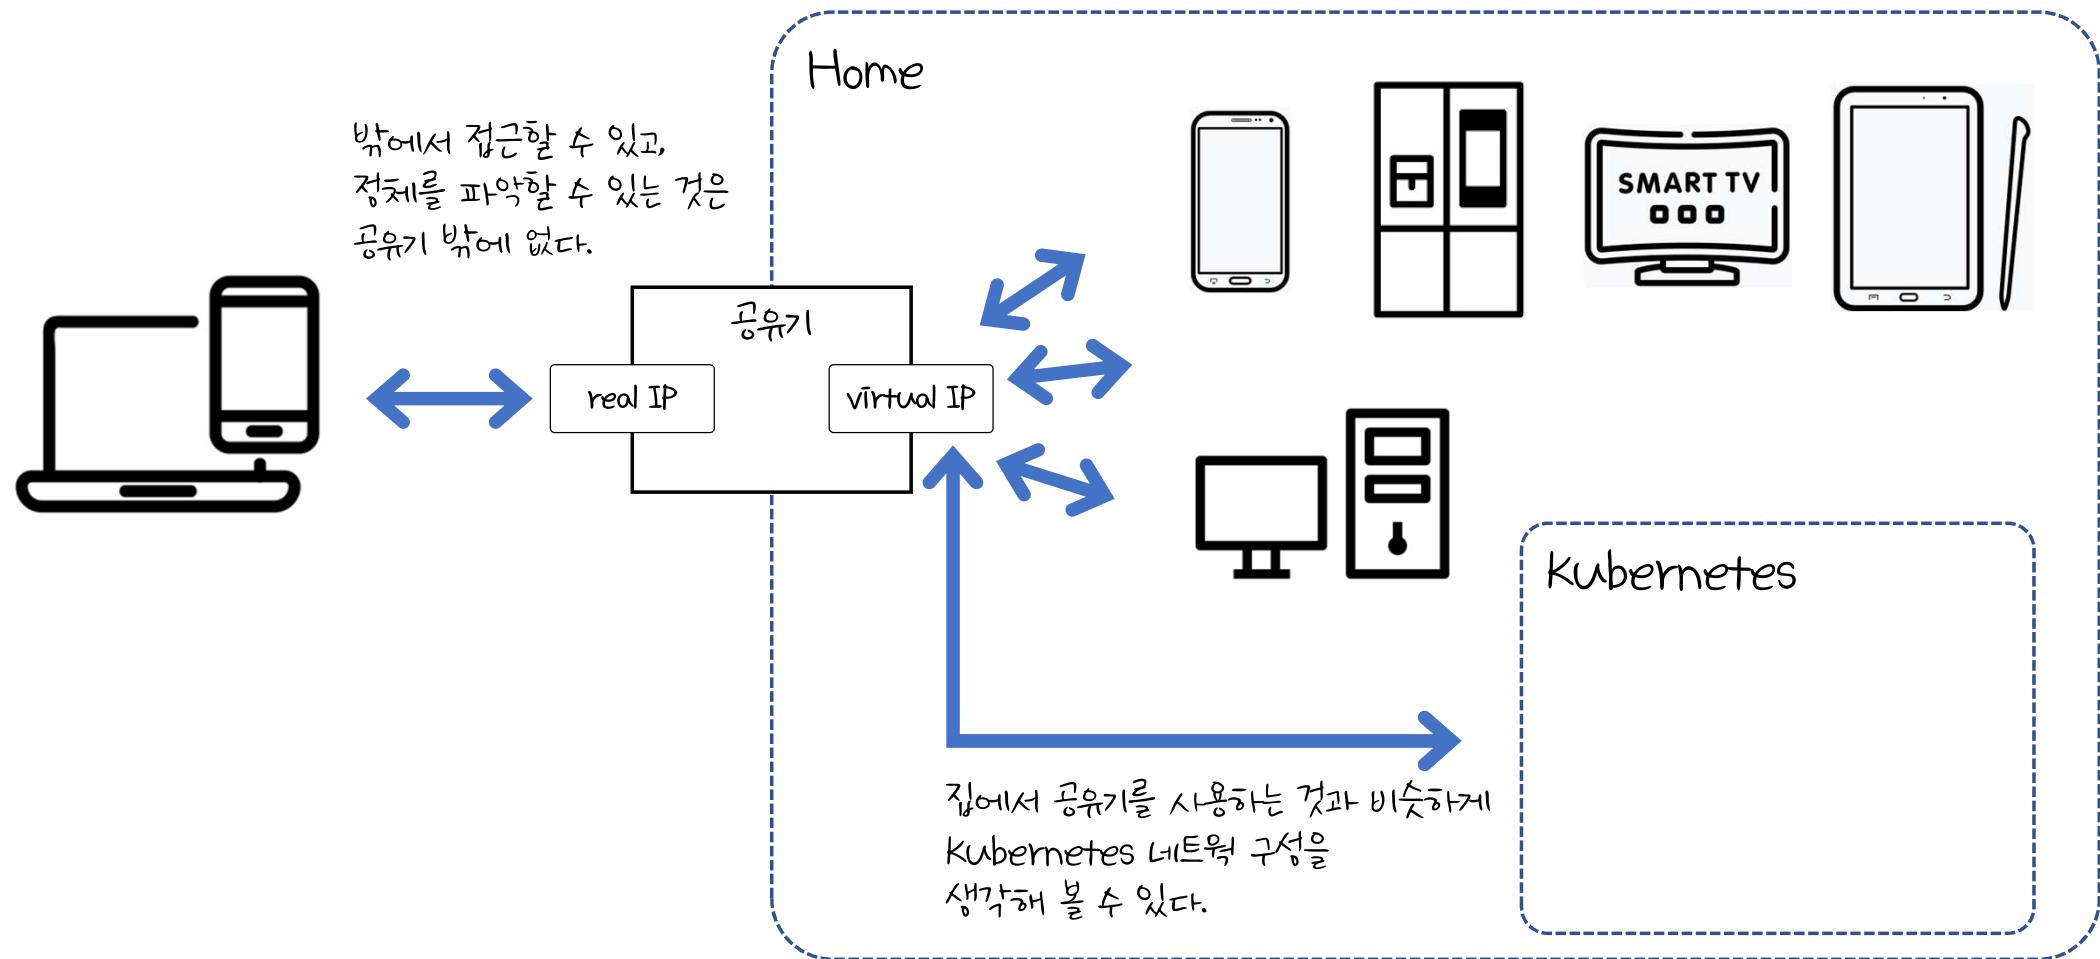
- ▷ [10m] make-friendship
 - ▷ [20m] DNS / FQDN / Endpoints / Headless
 - ▷ [20m] Flip - Ingress
 - ▷ [20m] Wrap-Up & Hands-On
 - ▷ [10m] Break-Time
-
- ▷ [30m] MetalLB
 - ▷ [40m] Advanced Network
 - ▷ [20m] Quiz

///

잠깐만~!

Network overview

Real World (Public Internet)



///

Kubernetes

Network - DNS / FQDN

DNS (Domain Name System)

- 특정 컴퓨터(또는 네트워크로 연결된 임의의 장치)의 주소를 찾기 위해, 사람이 이해하기 쉬운 도메인 이름을 숫자로 된 식별 번호(IP 주소)로 변환



※ 참고 : <https://www.youtube.com/watch?v=e2xLV7pCOLI>

※ 참고 : https://ko.wikipedia.org/wiki/도메인_네임_시스템

FQDN (Fully Qualified Domain Name)

- 명확한 도메인 표기법
 - . 소프트웨어 설치 중 도메인명을 요구하면, YAHOO.COM. 을 입력할지, WWW.YAHOO.COM. 을 입력할지 모호
→ Namespace 계층상에서 최종 호스트명을 포함하는 도메인명을 의미하는 FQDN을 요청하면 명확
- 원칙적으로 도메인의 표기는 네임스페이스상의 경로를 명확히 하기 위해 끝에 도트('루트 도메인)를 포함
 - . 하지만, 보통 도트를 생략하고 사용

www	호스트명
google.com.	도메인명
www.google.com.	FQDN

※ 참고 : <http://doc.kldp.org/KoreanDoc/html/PoweredByDNS-KLDP/fqdn.html>

CoreDNS

- CoreDNS는 쿠버네티스 클러스터의 DNS 역할을 수행할 수 있는, 유연하고 확장 가능한 DNS 서버
- CoreDNS 프로젝트도 CNCF가 관리
- 사용자는 기존 deployment인 kube-dns를 교체 가능

```
remote > kubectl get pods --namespace kube-system -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
calico-kube-controllers-5788f6558-g5fl2	1/1	Running	3 (46m ago)	23h	192.168.100.201	worker1	<none>	<none>
calico-node-67cdg	1/1	Running	2 (46m ago)	23h	192.168.100.202	worker2	<none>	<none>
calico-node-722h6	1/1	Running	2 (47m ago)	23h	192.168.100.201	worker1	<none>	<none>
calico-node-jtspn	1/1	Running	2 (47m ago)	23h	192.168.100.200	master	<none>	<none>
coredns-8474476ff8-8mxwm	1/1	Running	2 (46m ago)	23h	10.233.103.3	worker2	<none>	<none>
coredns-8474476ff8-z96xn	1/1	Running	2 (47m ago)	23h	10.233.70.7	master	<none>	<none>
dns-autoscaler-5ffdc7f89d-hhx9z	1/1	Running	2 (47m ago)	23h	10.233.70.9	master	<none>	<none>
kube-apiserver-master	1/1	Running	3 (47m ago)	23h	192.168.100.200	master	<none>	<none>
kube-controller-manager-master	1/1	Running	3 (47m ago)	23h	192.168.100.200	master	<none>	<none>
kube-proxy-gfts9	1/1	Running	2 (47m ago)	23h	192.168.100.201	worker1	<none>	<none>
kube-proxy-q8w2g	1/1	Running	2 (46m ago)	23h	192.168.100.202	worker2	<none>	<none>
kube-proxy-r6l2t	1/1	Running	2 (47m ago)	23h	192.168.100.200	master	<none>	<none>
kube-scheduler-master	1/1	Running	3 (47m ago)	23h	192.168.100.200	master	<none>	<none>
kubernetes-dashboard-548847967d-ntgzz	1/1	Running	3 (47m ago)	23h	10.233.110.6	worker1	<none>	<none>
kubernetes-metrics-scraper-6d49f96c97-56q6s	1/1	Running	2 (47m ago)	23h	10.233.110.5	worker1	<none>	<none>
metrics-server-6978dd689f-dgst9	1/1	Running	2 (47m ago)	23h	10.233.70.8	master	<none>	<none>
...								

※ 참고 : <https://kubernetes.io/ko/docs/tasks/administer-cluster/coredns/>

Hands-On - FQDN 1/3

FQDN 설정을 위해 ReplicaSet 생성해서 Pod들을 준비해 놓자

rs-node-web.yaml

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-labels

spec:
  replicas: 3

  selector:
    matchLabels:
      app: node-web

  template:
    metadata:
      labels:
        app: node-web
    spec:
      containers:
        - name: node-web
          image: whatwant/node-web:1.0
          ports:
            - containerPort: 8080
```

```
remote > cd advanced-kubernetes/03-week/Service
```

```
remote > kubectl create -f ./rs-node-web.yaml
```

```
replicaset.apps/rs-labels created
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-8vzrh	1/1	Running	0	58s	10.233.103.5	worker2	<none>	<none>
rs-labels-tbshs	1/1	Running	0	57s	10.233.103.4	worker2	<none>	<none>
rs-labels-vhb7v	1/1	Running	0	57s	10.233.110.7	worker1	<none>	<none>

Hands-On - FQDN 2/3

ClusterIP Endpoint으로 Service를 생성하자.

svc-node-web.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node

spec:
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: node-web
```

```
remote > kubectl create -f ./svc-node-web.yaml
```

```
service/svc-node created
```

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	23h	<none>
svc-node	ClusterIP	10.233.53.207	<none>	80/TCP	31s	app=node-web

당연히 Remote 환경에서는 접근을 할 수 없지만, Pod에서는 Service를 통해 접근할 수 있다.

```
remote > curl -s http://10.233.53.207
```

```
^C
```

```
remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://10.233.53.207
```

```
You've hit rs-labels-tbshs
```

Hands-On - FQDN 3/3

- Kubernetes는 각 container의 /etc/resolv.conf 파일을 직접 관리
- FQDN (Fully Qualified Domain Name) : service name . namespace name . svc.cluster.local

```
remote > kubectl exec -it rs-labels-8vzrh -- cat /etc/resolv.conf
search default.svc.cluster.local svc.cluster.local cluster.local skbroadband
nameserver 169.254.25.10
options ndots:5

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node
You've hit rs-labels-vhb7v

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node.default
You've hit rs-labels-vhb7v

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node.default.svc
You've hit rs-labels-8vzrh

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node.default.svc.cluster
^Ccommand terminated with exit code 130

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node.default.svc.cluster.local
^Ccommand terminated with exit code 130
```

Trouble-Shooting

- 앞에서 진행한 결과를 보면 원하는 내용과 다르다. 원인은 잘못된 DNS 정보 삽입. 이유는 Network 설정 이슈!
- 아래 내용은 각 상황에 따라 다를 수 있음을 염두에 두어야 한다.

```
remote > ssh vagrant@192.168.100.200
```

```
master > sudo nano /etc/netplan/00-installer-config.yaml
```

```
master > sudo nano /etc/netplan/50-vagrant.yaml
```

```
master > sudo reboot
```

Network 설정 파일 2개를 옆의 내용 참고해서 수정한 뒤, 재부팅 하자

master 뿐만 아니라, worker1 / worker2 도 마찬가지로 수정해주어야 한다.

공유기 환경에서 DHCP 설정으로 인해 공유기의 DNS 설정이 반영되어 발생하는 문제이다.

00-installer-config.yaml

```
# This is the network config written by 'subiquity'  
network:  
  ethernets:  
    enp0s3:  
      dhcp4: no  
      version: 2
```

50-vagrant.yaml

```
---  
network:  
  version: 2  
  renderer: networkd  
  ethernets:  
    enp0s8:  
      addresses:  
      - 192.168.100.200/24  
      dhcp4: no  
      dhcp6: no  
      gateway4: 192.168.100.1  
      nameservers:  
        addresses: [8.8.8.8,8.8.4.4]
```

Hands-On - FQDN 3/3 - Fix!

- Kubernetes는 각 container의 /etc/resolv.conf 파일을 직접 관리
- FQDN (Fully Qualified Domain Name) : service name . namespace name . svc.cluster.local

```
remote > kubectl exec -it rs-labels-8vzrh -- cat /etc/resolv.conf
search default.svc.cluster.local svc.cluster.local cluster.local
nameserver 169.254.25.10
options ndots:5

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node
You've hit rs-labels-vhb7v

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node.default
You've hit rs-labels-vhb7v

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node.default.svc
You've hit rs-labels-8vzrh

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node.default.svc.cluster
command terminated with exit code 6

remote > kubectl exec -it rs-labels-8vzrh -- curl -s http://svc-node.default.svc.cluster.local
You've hit rs-labels-tbshs
```

///

Kubernetes

Network - Endpoints/Headless

Service

- Service로 노출되는 Pod의 IP 주소와 Port 목록
- Client → kube-proxy → Endpoints 중 하나의 IP/Port 선택 → 들어온 연결을 대상 Pod의 수신 대기 서버로 전달

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-8vzrh	1/1	Running	3 (3h1m ago)	4h2m	10.233.103.13	worker2	<none>	<none>
rs-labels-tbshs	1/1	Running	3 (3h1m ago)	4h2m	10.233.103.12	worker2	<none>	<none>
rs-labels-vhb7v	1/1	Running	3 (3h1m ago)	4h2m	10.233.110.16	worker1	<none>	<none>

```
remote > kubectl describe service svc-node
```

```
Name:           svc-node
Namespace:      default
Labels:          <none>
Annotations:    <none>
Selector:        app=node-web
Type:           ClusterIP
IP Family Policy: SingleStack
IP Families:    IPv4
IP:             10.233.53.207
IPs:            10.233.53.207
Port:           <unset>  80/TCP
TargetPort:     8080/TCP
Endpoints:      10.233.103.12:8080,10.233.103.13:8080,10.233.110.16:8080
Session Affinity: None
Events:          <none>
```

svc-node-web.yaml

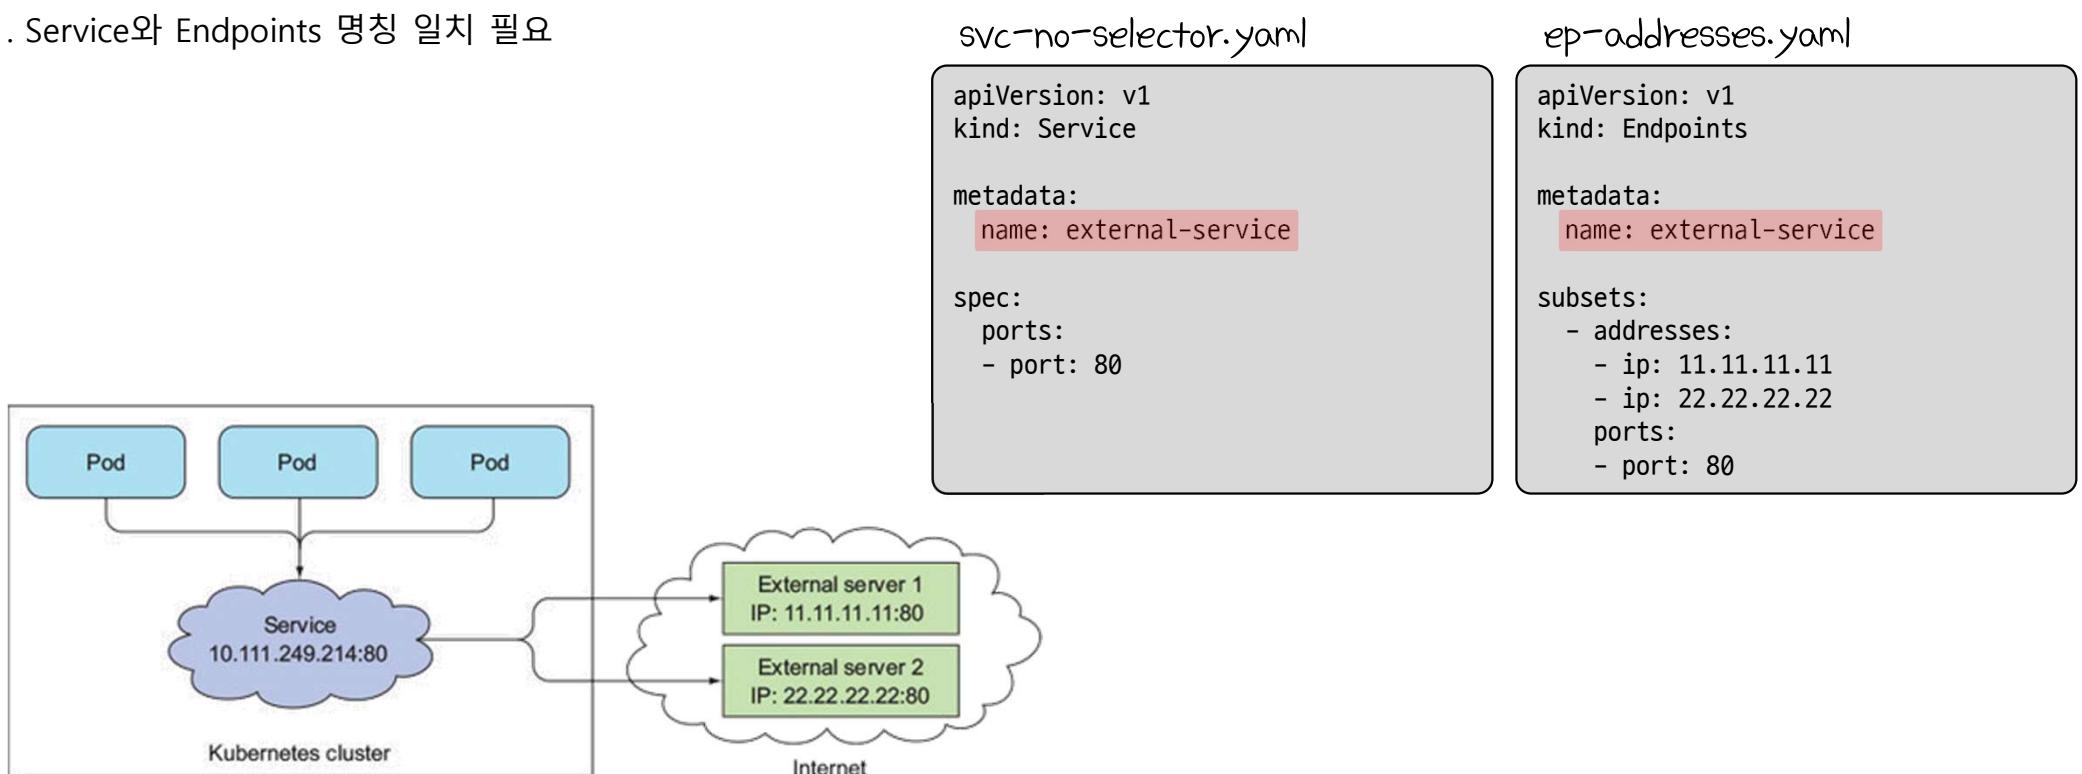
```
apiVersion: v1
kind: Service
metadata:
  name: svc-node

spec:
  ports:
    - port: 80
      targetPort: 8080

  selector:
    app: node-web
```

Endpoints

- Endpoint 리소스도 만들 수 있을까?
- . Selectors를 지정하지 않고, 직접 생성한 Endpoints 활용하기
- . Service와 Endpoints 명칭 일치 필요



※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-5/144>

Headless Service

- Pod 하나에 연결하는 것이 아니라, 생성된 모든 Pod에 연결하고 싶은 경우. 또는, 같은 Service에 묶인 다른 Pod와의 통신이 필요한 경우 사용
- `clusterIP: None` 설정을 하고, FQDN을 통해 접근

svc-node-web-headless.yaml

```
apiVersion: v1
kind: Service

metadata:
  name: svc-node-web-headless

spec:
  clusterIP: None

  selector:
    app: node-web

  ports:
  - port: 80
    targetPort: 8080
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes/04-week/headless

remote > kubectl create -f ./svc-node-web-headless.yaml
service/svc-node-web-headless created

remote > kubectl get services -o wide
NAME           TYPE        CLUSTER-IP      EXTERNAL-IP   PORT(S)    AGE     SELECTOR
kubernetes     ClusterIP   10.233.0.1    <none>       443/TCP   28h    <none>
svc-node       ClusterIP   10.233.53.207 <none>       80/TCP    4h47m   app=node-web
svc-node-web-headless ClusterIP   None          <none>       80/TCP    26s    app=node-web
```

※ 참고 : <https://kubernetes.io/docs/tasks/administer-cluster/dns-debugging-resolution/>

Headless Service Check

Headless Service는 DNS를 통해서 활용하는 것인기에 이를 확인하기 위해 'nslookup' 명령어를 이용해보고자 했다.
'nslookup' 명령어가 설치되어 있는 Pod를 하나 생성하고 이를 사용했다.

```
remote > kubectl apply -f https://k8s.io/examples/admin/dns/dnsutils.yaml
```

```
pod/dnsutils created
```

```
remote > kubectl exec -it dnsutils -- nslookup svc-node-web-headless
```

```
Server: 169.254.25.10
Address: 169.254.25.10#53

Name: svc-node-web-headless.default.svc.cluster.local
Address: 10.233.103.13
Name: svc-node-web-headless.default.svc.cluster.local
Address: 10.233.103.12
Name: svc-node-web-headless.default.svc.cluster.local
Address: 10.233.110.16
```

3개가 다 확인된다

```
remote > kubectl exec -it dnsutils -- nslookup svc-node
```

```
Server: 169.254.25.10
Address: 169.254.25.10#53

Name: svc-node.default.svc.cluster.local
Address: 10.233.53.207
```

1개가 확인되는 것이 일반적이다.

dnsutils.yaml

```
apiVersion: v1
kind: Pod

metadata:
  name: dnsutils
  namespace: default

spec:
  containers:
    - name: dnsutils
      image: k8s.gcr.io/e2e-test-images/jessie-dnsutils:1.3

    command:
      - sleep
      - "3600"

  imagePullPolicy: IfNotPresent

  restartPolicy: Always
```

///

Flip Learning

(Network - Ingress)

박남준님

///

Kubernetes

Network - Ingress

Ingress Controller

- Ingress는 cluster 외부에서 내부 Service로 HTTP와 HTTPS 경로 노출
- 클라이언트가 요청한 호스트와 경로에 따라 요청을 전달할 서비스가 결정
 - . HTTP(S)기반의 L7 LoadBalancing 기능을 제공하는 컴포넌트
- Ingress 구현체 : 각 구현체마다 설정 방법 차이 존재
 - . 구글 클라우드 : <https://github.com/kubernetes/ingress-gce>
 - . 오픈소스 (Ingress NGINX Controller) : <https://github.com/kubernetes/ingress-nginx>
 - . 오픈소스 (NGINX Ingress Controller) : <https://github.com/nginxinc/kubernetes-ingress>
 - . 상용 (F5 BIG IP Controller) : <http://clouddocs.f5.com/products/connectors/k8s-bigip-ctlr>



※ 참고 : <https://kubernetes.io/ko/docs/concepts/services-networking/ingress/>

Ingress Controller Install

Kubernetes에서 배포하고 있는 패키지로 설치해보는데, 아래 링크에서 bare-metal 설치 방법을 찾아보면 된다

- <https://github.com/kubernetes/ingress-nginx/blob/main/docs/deploy/Index.md#bare-metal-clusters>

```
remote > kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.1.1/deploy/static/provider/baremetal/deploy.yaml
```

```
namespace/ingress-nginx created
serviceaccount/ingress-nginx created
configmap/ingress-nginx-controller created
clusterrole.rbac.authorization.k8s.io/ingress-nginx created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx created
role.rbac.authorization.k8s.io/ingress-nginx created
rolebinding.rbac.authorization.k8s.io/ingress-nginx created
service/ingress-nginx-controller-admission created
service/ingress-nginx-controller created
deployment.apps/ingress-nginx-controller created
ingressclass.networking.k8s.io/nginx created
validatingwebhookconfiguration.admissionregistration.k8s.io/ingress-nginx-admission created
serviceaccount/ingress-nginx-admission created
clusterrole.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
role.rbac.authorization.k8s.io/ingress-nginx-admission created
rolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
job.batch/ingress-nginx-admission-create created
job.batch/ingress-nginx-admission-patch created
```

무언가 몽땅 설치하는 것을 볼 수 있다.

Ingress Controller Check

```
remote > kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	2d7h
ingress-nginx	Active	5m47s
...		

```
remote > kubectl get pods -o wide --namespace kube-system
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
...								
nginx-proxy-worker1	1/1	Running	7 (44m ago)	2d7h	192.168.100.201	worker1	<none>	<none>
nginx-proxy-worker2	1/1	Running	7 (44m ago)	2d7h	192.168.100.202	worker2	<none>	<none>
...								

```
remote > kubectl get pods -o wide --namespace ingress-nginx
```

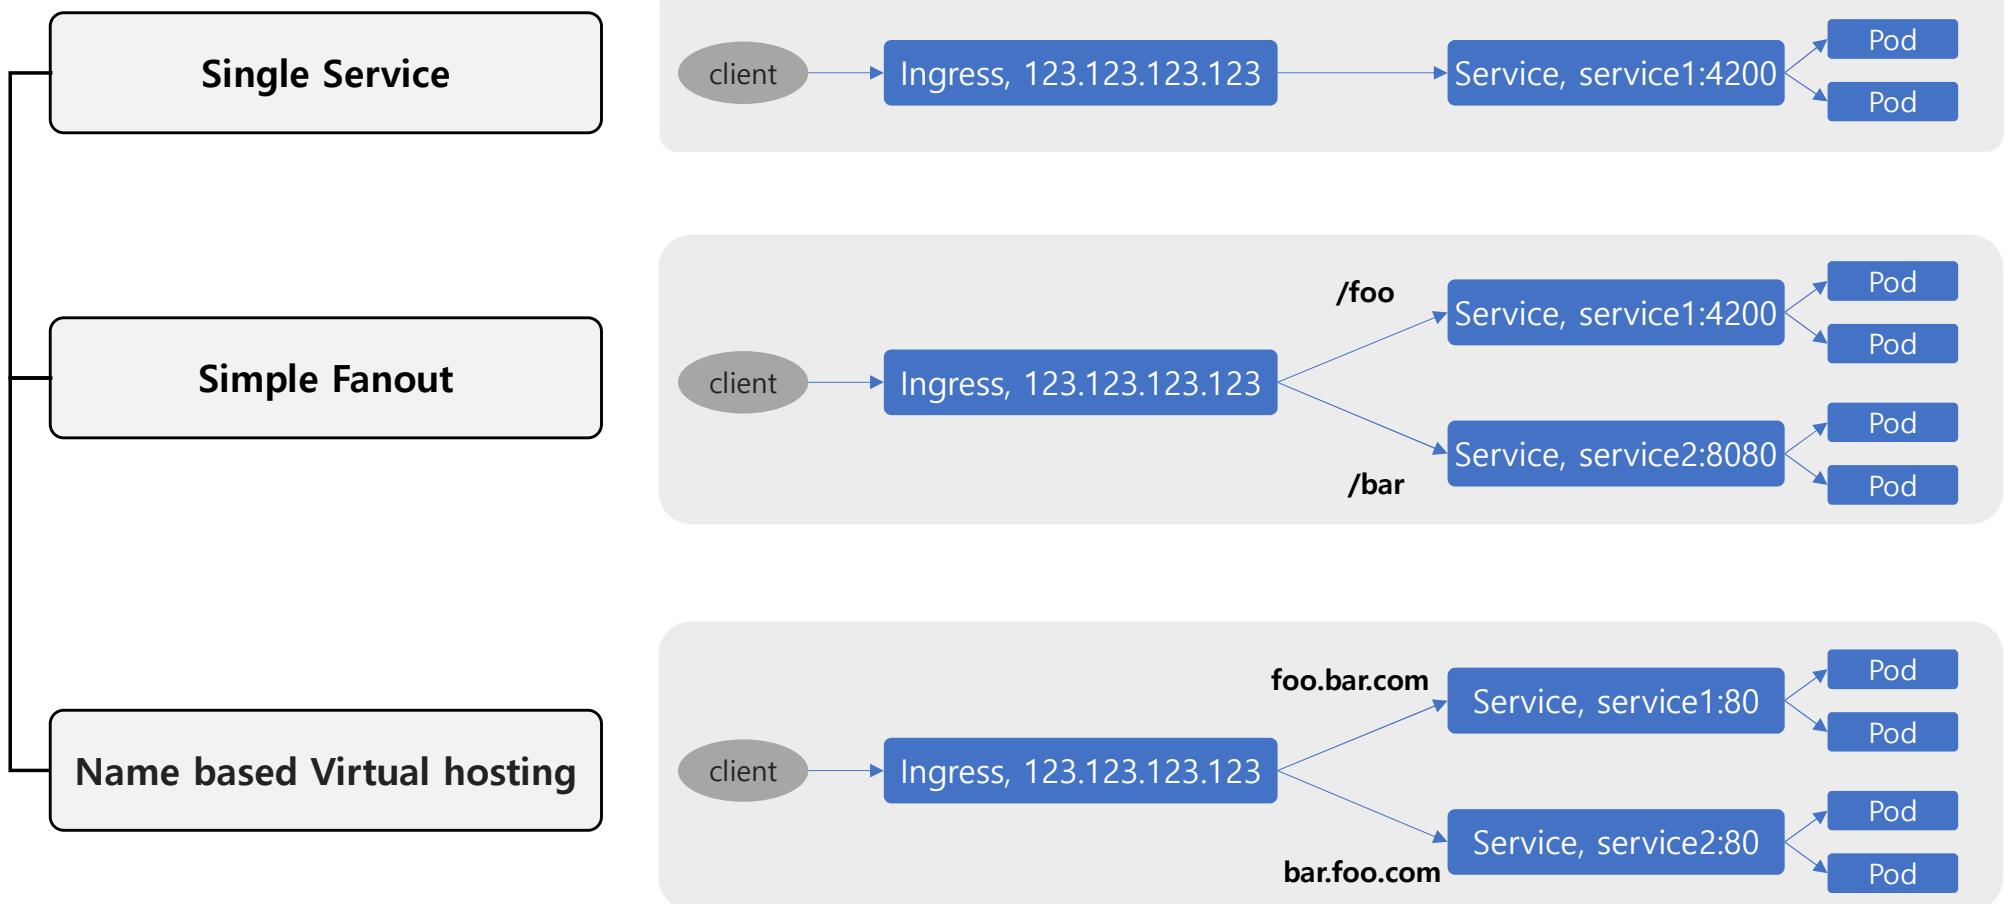
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
ingress-nginx-admission-create--1-gr5zf	0/1	Completed	0	10m	10.233.110.20	worker1	<none>	<none>
ingress-nginx-admission-patch--1-87plx	0/1	Completed	0	10m	10.233.103.16	worker2	<none>	<none>
ingress-nginx-controller-778574f59b-7rnzn	1/1	Running	0	10m	10.233.103.17	worker2	<none>	<none>

```
remote > kubectl get serviceaccounts -o wide --namespace ingress-nginx
```

NAME	SECRETS	AGE
default	1	15m
ingress-nginx	1	15m
ingress-nginx-admission	1	15m

///

Types of Ingress



※ 참고 : <https://kubernetes.io/docs/concepts/services-networking/ingress/>

///

Hands-On: Container Image 준비

'Simple Fanout' 태입의 Ingress 실습을 위해서 Pod를 하나 더 필요해서, docker Image를 하나 더 만들었다.

app.js

```
const http = require('http');
const os = require('os');

console.log("node-web server starting...");

var handler = function(request, response) {
  console.log("Received request from " + request.connection.remoteAddress);
  response.writeHead(200);
  response.end("You've hit " + os.hostname() + "(Ver2.0)\n");
};

var www = http.createServer(handler);
www.listen(8080);
```

Dockerfile

```
FROM node:latest
ADD app.js /app.js
ENTRYPOINT ["node", "app.js"]
```

기존 'node-web:1.0'과 구분이 될 정도만 살짝 변경한 후
'node-web:2.0' 버전으로 push 했다.

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes/04-week/container
```

```
remote > docker build -t node-web:2.0 .
```

```
remote > docker tag node-web:2.0 <user-id>/node-web:2.0
```

```
remote > docker push <user-id>/node-web:2.0
```

Hands-On: ReplicaSet YAML

'ReplicaSet'을 이용해서 2종의 어플리케이션 그룹을 만들어 놓으려고 한다.

rs-web-v1.yaml

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-web-v1

spec:
  replicas: 3
  selector:
    matchLabels:
      app: node-web-v1

  template:
    metadata:
      labels:
        app: node-web-v1

  spec:
    containers:
      - name: node-web
        image: whatwant/node-web:1.0
        ports:
          - containerPort: 8080
        imagePullPolicy: Always
```

rs-web-v2.yaml

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-web-v2

spec:
  replicas: 3
  selector:
    matchLabels:
      app: node-web-v2

  template:
    metadata:
      labels:
        app: node-web-v2

  spec:
    containers:
      - name: node-web
        image: whatwant/node-web:2.0
        ports:
          - containerPort: 8080
        imagePullPolicy: Always
```

Hands-On: ReplicaSet 생성

3개 | 묶음으로 해서 2개의 묶음으로 미리 만들어 놓자

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes/04-week/replica

remote > kubectl create -f ./rs-web-v1.yaml
replicaset.apps/rs-web-v1 created

remote > kubectl create -f ./rs-web-v2.yaml
replicaset.apps/rs-web-v2 created

remote > kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE      NOMINATED NODE   READINESS GATES
rs-web-v1-7zv4l 1/1     Running   0          3m27s  10.233.110.25  worker1   <none>        <none>
rs-web-v1-9lcxc 1/1     Running   0          3m27s  10.233.103.25  worker2   <none>        <none>
rs-web-v1-d59sr 1/1     Running   0          3m27s  10.233.103.26  worker2   <none>        <none>
rs-web-v2-cdv4s 1/1     Running   0          3m8s   10.233.103.27  worker2   <none>        <none>
rs-web-v2-gkhs6 1/1     Running   0          3m8s   10.233.110.26  worker1   <none>        <none>
rs-web-v2-zhccb 1/1     Running   0          3m8s   10.233.103.28  worker2   <none>        <none>
```

Hands-On: Service 생성

'ClusterIP' 유형으로 해도 되지만, 여기에서는 'NodePort' 유형으로 'Service'를 만들어보자.

svc-web-nodeport-v1.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node-web-v1

spec:
  type: NodePort
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30001
  selector:
    app: node-web-v1
```

svc-web-nodeport-v2.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node-web-v2

spec:
  type: NodePort
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30002
  selector:
    app: node-web-v2
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes/04-week/service
```

```
remote > kubectl create -f ./svc-web-nodeport-v1.yaml
remote > kubectl create -f ./svc-web-nodeport-v2.yaml
```

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	7d20h	<none>
svc-node-web-v1	NodePort	10.233.28.43	<none>	80:30001/TCP	8s	app=node-web-v1
svc-node-web-v2	NodePort	10.233.26.164	<none>	80:30002/TCP	5s	app=node-web-v2

///

Hands-On: Status Check

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-web-v1-7zv4l	1/1	Running	0	4h40m	10.233.110.25	worker1	<none>	<none>
rs-web-v1-9lcxc	1/1	Running	0	4h40m	10.233.103.25	worker2	<none>	<none>
rs-web-v1-d59sr	1/1	Running	0	4h40m	10.233.103.26	worker2	<none>	<none>
rs-web-v2-gzbqb	1/1	Running	0	44s	10.233.110.27	worker1	<none>	<none>
rs-web-v2-p85wx	1/1	Running	0	44s	10.233.103.29	worker2	<none>	<none>
rs-web-v2-sd2ds	1/1	Running	0	44s	10.233.103.30	worker2	<none>	<none>

```
remote > kubectl get replicasesets -o wide
```

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
rs-web-v1	3	3	3	4h40m	node-web	whatwant/node-web:1.0	app=node-web-v1
rs-web-v2	3	3	3	50s	node-web	whatwant/node-web:2.0	app=node-web-v2

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	7d20h	<none>
svc-node-web-v1	NodePort	10.233.28.43	<none>	80:30001/TCP	81s	app=node-web-v1
svc-node-web-v2	NodePort	10.233.26.164	<none>	80:30002/TCP	78s	app=node-web-v2

```
remote > curl -s http://192.168.100.201:30001
```

You've hit rs-web-v1-wkkk9

```
remote > curl -s http://192.168.100.201:30002
```

You've hit rs-web-v2-gzbqb (Ver2.0)

Hands-On: Ingress Create (Simple Fanout)

ingress-node-web.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress

metadata:
  name: ingress-node-web
  annotations:
    kubernetes.io/ingress.class: "nginx"
    nginx.ingress.kubernetes.io/rewrite-target: /

spec:
  rules:
  - http:
      paths:
      - path: /foo
        pathType: Prefix
        backend:
          service:
            name: svc-node-web-v1
            port:
              number: 80
      - path: /bar
        pathType: Prefix
        backend:
          service:
            name: svc-node-web-v2
            port:
              number: 80
```

```
- path: /bar
  pathType: Prefix
  backend:
    service:
      name: svc-node-web-v2
      port:
        number: 80
```

'host'를 별도로 지정하지 않으면, 들어오는 모든 입력에 적용된다.
즉, IP로 접근을 하더라도 해당 내용이 적용되는 것이다.

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes/04-week/ingress
```

```
remote > kubectl create -f ./ingress-node-web.yaml
```

```
ingress.networking.k8s.io/ingress-node-web created
```

Hands-On: Simple Fanout 1/2

앞에서 생성한 Ingress를 확인해 보자

```
remote > kubectl get ingresses -o wide
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
ingress-node-web	<none>	*	192.168.100.201	80	29s

```
remote > kubectl describe ingress ingress-node-web
```

Name:	ingress-node-web			
Namespace:	default			
Address:	192.168.100.201			
Default backend:	default-http-backend:80 (<error: endpoints "default-http-backend" not found>)			
Rules:				
Host	Path	Backends		
---	---	-----		
*				
	/foo	svc-node-web-v1:80 (10.233.110.33:8080,10.233.110.34:8080,10.233.110.35:8080)		
	/bar	svc-node-web-v2:80 (10.233.110.32:8080,10.233.110.36:8080,10.233.110.37:8080)		
Annotations:	kubernetes.io/ingress.class: nginx			
	nginx.ingress.kubernetes.io/rewrite-target: /			
Events:				
Type	Reason	Age	From	Message
---	---	---	---	-----
Normal	Sync	43s (x2 over 49s)	nginx-ingress-controller	Scheduled for sync

Hands-On: Simple Fanout 2/2

우리가 설치한 Ingress Controller는 NodePort로 설정되어 있으며, 80/433 포트가 이미 등록되어 있다. 확인해 보자.

```
remote > kubectl get services -o wide --namespace ingress-nginx
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
ingress-nginx-controller	NodePort	10.233.45.186	<none>	80:31459/TCP,443:30693/TCP	5d13h	app.kubernetes.io/component=co..
ingress-nginx-controller-admission	ClusterIP	10.233.30.23	<none>	443/TCP	5d13h	app.kubernetes.io/component=co..

그려면, 이제 우리가 생성한 Ingress를 이용해서 접속이 잘되는지 확인해 보자.

```
remote > curl -s http://192.168.100.201:31459/foo
```

```
You've hit rs-web-v1-wkkk9
```

```
remote > curl -s http://192.168.100.201:31459/bar
```

```
You've hit rs-web-v2-gzbqb (Ver2.0)
```

///

Break

돌아오셨으면 채팅창에
복귀!
타이핑하기!

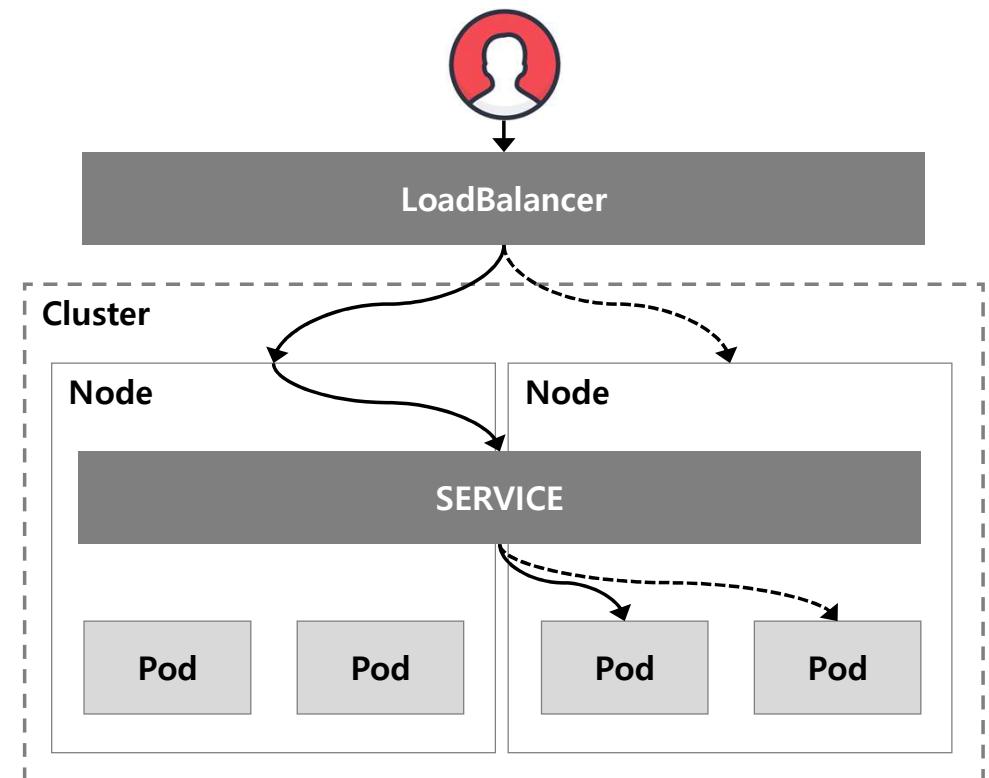
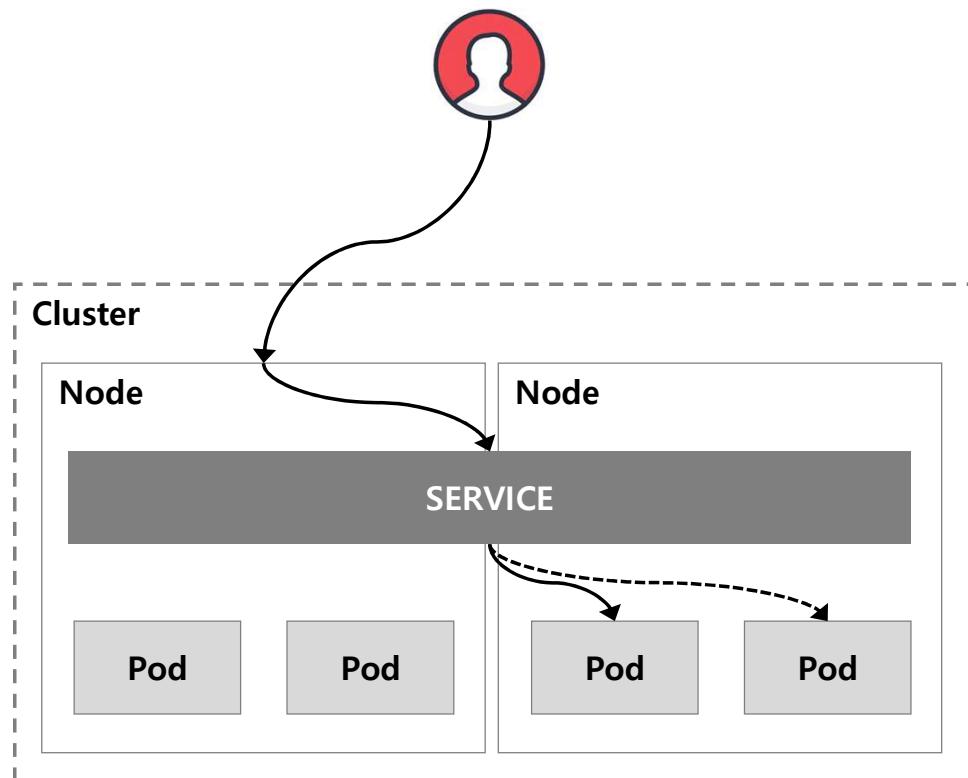
///

Kubernetes

Network - LoadBalancer

Why LoadBalancer ?

NodePort vs LoadBalancer



KCD 2020

Korea Community Day 2020

- <https://kcd2020.festa.io/>

시간	Track 1	Track 2
13:00 ~ 13:25	백 마디 말보다 중요한 것: Backend.AI의 오픈소스 문서 작성과 자동화 경험 신정규 (Backend.AI) 발표자료 / 발표영상	코드 주석으로 GitHub Pages를 만들 수 있다? 박동하 (C++ Korea) 발표자료 / 발표영상
13:25 ~ 13:50	쿠버네티스의 로드밸런서 서비스를 온프레姆(On-Premises) 에서 쓸순 없을까? 조훈 (IT인프라엔지니어그룹) 발표자료 / 발표영상	내가 태국(망해가는 회사)에서 일하는 방법 김태우 (비원이즈) 발표자료 / 발표영상
13:50 ~ 14:15	오픈소스 커뮤니티의 딜레마 김종민 (Elastic) 발표자료 / 발표영상	Win32 API로 다루는 WSL 남정현 (한국WSL 사용자그룹) 발표자료 / 발표영상

[발표자료] : [링크](#)

[강의영상] : <https://www.youtube.com/watch?v=wIQHQbdiYQk>

///

MetalLB

Installation

Requirement - 1/3

- MetalLB requires the following to function:

- ① A Kubernetes cluster, running Kubernetes 1.13.0 or later, that does not already have network load-balancing functionality.
- ② A cluster network configuration that can coexist with MetalLB.
- ③ Some IPv4 addresses for MetalLB to hand out.
- ④ When using the BGP operating mode, you will need one or more routers capable of speaking BGP.
- ⑤ Traffic on port 7946 (TCP & UDP) must be allowed between nodes, as required by memberlist.

① Kubernetes version

```
remote > kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME
master	Ready	control-plane,master	9d	v1.22.5	192.168.100.200	<none>	Ubuntu 20.04.3 LTS	5.4.0-96-generic	containerd://1.5.8
worker1	Ready	<none>	9d	v1.22.5	192.168.100.201	<none>	Ubuntu 20.04.3 LTS	5.4.0-96-generic	containerd://1.5.8
worker2	Ready	<none>	9d	v1.22.5	192.168.100.202	<none>	Ubuntu 20.04.3 LTS	5.4.0-96-generic	containerd://1.5.8

Kubernetes 1.13.0 버전 이상이고, 다른 load-balancing 기능을 갖고 있지 않으면 된다.

※ 참고 : <https://metallb.universe.tf/#requirements>

Requirement - 2/3

② Cluster Network

```
remote > kubectl get pods --namespace kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-5788f6558-g5fl2	1/1	Running	16 (34h ago)	9d
calico-node-67cdg	1/1	Running	13 (34h ago)	9d
calico-node-722h6	1/1	Running	13 (34h ago)	9d
calico-node-jtspn	1/1	Running	14 (34h ago)	9d
coredns-8474476ff8-mk7wr	1/1	Running	1 (34h ago)	2d
coredns-8474476ff8-z96xn	1/1	Running	9 (34h ago)	9d
dns-autoscaler-5ffdc7f89d-hhx9z	1/1	Running	9 (34h ago)	9d
kube-apiserver-master	1/1	Running	11 (34h ago)	9d
kube-controller-manager-master	1/1	Running	11 (34h ago)	9d
...				

Calico 환경에서는 알려진 문제가 있다지만,

BGP를 사용할 경우에만 해당하는 문제이니 일단 무시하고 진행~

Network addon	Compatible
Calico	Mostly (see known issues)
Canal	Yes
Cilium	Yes
Flannel	Yes
Kube-ovn	Yes
Kube-router	Mostly (see known issues)
Weave Net	Mostly (see known issues)

③ IPv4 addresses



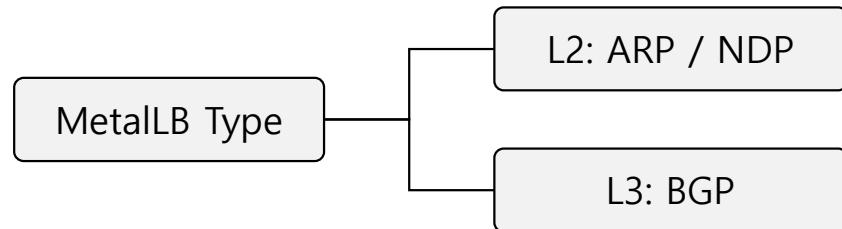
MetallLB에 할당할 IP 대역을 확보해야 한다.

※ 참고 : <https://metallb.universe.tf/installation/network-addons/>

로컬 클라이언트 상태		
MAC 주소	IP 주소	장치명
00:0c:29:00:00:02	192.168.100.102	로컬 클라이언트 2
00:0c:29:00:00:01	192.168.100.107	로컬 클라이언트 1
00:0c:29:00:00:03	192.168.100.101	로컬 클라이언트 3
00:0c:29:00:00:04	192.168.100.200	로컬 클라이언트 4
00:0c:29:00:00:05	192.168.100.201	로컬 클라이언트 5
00:0c:29:00:00:06	192.168.100.202	로컬 클라이언트 6

Requirement - 3/3

④ When using the BGP (Boarder Gateway Protocol)...



MetallLB는 L2/L3 2가지 방식 중 하나를 선택해서 설치할 수 있다.

L3(BGP) 방식을 선택할 경우 추가적으로 고려할 사항들이 있다.

하지만, 여기에서는 일단 L2로 진행할 계획이기에 Skip ...

L2	Data Link	Mac 주소 기반	
L3	Network	IP 주소 기반	
L4	Transport	Port 기반	TCP, UDP
L7	Application	요청(URL) 기반	HTTP, HTTPS

⑤ Traffic on port 7946 (TCP & UDP)

MetallLB가 동작하기 위해서는 Node끼리 port 7946 (TCP/UDP) 통신이 되어야 한다.

보통의 상황에서는 문제가 없겠지만, 방화벽이 있는 환경에서는 해당 port가 열려 있도록 주의해야 한다.

Preparation

- If you're using kube-proxy in IPVS mode, since Kubernetes v1.14.2 you have to enable strict ARP mode.

```
remote > export KUBE_EDITOR=nano  
remote > kubectl edit configmap -n kube-system kube-proxy
```

kube-proxy²를 IPVS 모드에서 사용한다고 하면 strict ARP²를 enable 시켜주어야 한다.

하지만, iptables 모드에서 사용하는 경우에는 Skip ...

```
...  
  iptables:  
    masqueradeAll: false  
    masqueradeBit: 14  
    minSyncPeriod: 0s  
    syncPeriod: 30s  
  ipvs:  
    excludeCIDRs: []  
    minSyncPeriod: 0s  
    scheduler: rr  
    strictARP: false  
    syncPeriod: 30s  
    tcpFinTimeout: 0s  
    tcpTimeout: 0s  
    udpTimeout: 0s  
  kind: KubeProxyConfiguration  
  metricsBindAddress: 127.0.0.1:10249  
  mode: iptables  
  nodePortAddresses: []  
...
```

※ 참고 : <https://metallb.universe.tf/installation/>

Check YAML (skippable)

- MetalLB를 설치하기에 앞서서 설치를 위한 YAML 파일을 다운로드 받은 후 한 번 살펴보기를 권장한다

```
remote > wget https://raw.githubusercontent.com/metallb/metallb/v0.11.0/manifests/namespace.yaml  
remote > wget https://raw.githubusercontent.com/metallb/metallb/v0.11.0/manifests/metallb.yaml
```

namespace.yaml

```
apiVersion: v1  
kind: Namespace  
metadata:  
  name: metallb-system  
  labels:  
    app: metallb
```

꼭 해야 하는 과정은 아니다.

그냥 막 설치하기에 앞서서 미리 한번 살펴보라는 의미.

metallb.yaml

```
apiVersion: policy/v1beta1  
kind: PodSecurityPolicy  
metadata:  
  labels:  
    app: metallb  
  name: controller  
  namespace: metallb-system  
spec:  
  allowPrivilegeEscalation: false  
  allowedCapabilities: []  
  allowedHostPaths: []  
  defaultAddCapabilities: []  
  defaultAllowPrivilegeEscalation: false  
  fsGroup:  
    ranges:  
      - max: 65535  
        min: 1  
      rule: MustRunAs  
  hostIPC: false  
  hostNetwork: false  
...
```

Installation By Manifest

- 이제는 설치를 진행해 보자.

```
remote > kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.11.0/manifests/namespace.yaml
```

```
namespace/metallb-system created
```

```
remote > kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.11.0/manifests/metallb.yaml
```

```
Warning: policy/v1beta1 PodSecurityPolicy is deprecated in v1.21+, unavailable in v1.25+
```

```
podsecuritypolicy.policy/controller created
```

```
podsecuritypolicy.policy/speaker created
```

```
serviceaccount/controller created
```

```
serviceaccount/speaker created
```

```
clusterrole.rbac.authorization.k8s.io/metallb-system:controller created
```

```
clusterrole.rbac.authorization.k8s.io/metallb-system:speaker created
```

```
role.rbac.authorization.k8s.io/config-watcher created
```

```
role.rbac.authorization.k8s.io/pod-lister created
```

```
role.rbac.authorization.k8s.io/controller created
```

```
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:controller created
```

```
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:speaker created
```

```
rolebinding.rbac.authorization.k8s.io/config-watcher created
```

```
rolebinding.rbac.authorization.k8s.io/pod-lister created
```

```
rolebinding.rbac.authorization.k8s.io/controller created
```

```
daemonset.apps/speaker created
```

```
deployment.apps/controller created
```

설치 후에 configmap 작업까지 해줘야 끝난다.

Check & Edit Config

- 설치 잘 되었는지 보고, 설정 마무리를 진행해 보자.

```
remote > kubectl get pods -o wide --namespace metallb-system
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
controller-7dcc8764f4-8m4dv	1/1	Running	0	3m48s	10.233.103.31	worker2	<none>	<none>	
speaker-nfjl8	1/1	Running	0	3m48s	192.168.100.201	worker1	<none>	<none>	
speaker-qjdtg	1/1	Running	0	3m48s	192.168.100.202	worker2	<none>	<none>	
speaker-zw7x7	1/1	Running	0	3m48s	192.168.100.200	master	<none>	<none>	

metallb-config.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  namespace: metallb-system
  name: config
data:
  config: |
    address-pools:
    - name: default
      protocol: layer2
      addresses:
      - 192.168.100.240-192.168.100.250
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes/04-week/metallb
```

```
remote > kubectl create -f ./metallb-config.yaml
```

```
configmap/config created
```

앞에서도 말했지만, 지금은 layer2 모드로 구성하고자 한다.

마찬가지로 앞에서 확인한 부여 가능한 IP 대역에 따라서 config 내역 잡아주고 생성하면 된다.

※ 참고 : <https://metallb.universe.tf/configuration/#layer-2-configuration>

///

Hands-On

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git  
remote > cd advanced-kubernetes/03-week/ReplicaSet
```

```
remote > kubectl create -f ./rs-node-web.yaml
```

```
replicaset.apps/rs-labels created
```

ReplicaSet으로 만들어진 여러 개의 Pod를 준비하고,

```
remote > cd ../../04-week/metallb
```

이를 LoadBalancer로 묶어주는 과정

```
remote > kubectl create -f ./svc-lb-web.yaml
```

→ 독립적인 IP를 할당 받아서 접근 가능

```
service/svc-lb created
```

svc-lb-web.yaml

```
apiVersion: v1  
kind: Service  
metadata:  
  name: svc-lb  
  
spec:  
  type: LoadBalancer  
  
  ports:  
  - name: http  
    port: 80  
    protocol: TCP  
    targetPort: 8080  
  
  selector:  
    app: node-web
```

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	9d	<none>
svc-lb	LoadBalancer	10.233.55.71	192.168.100.240	80:32380/TCP	2m38s	app=node-web

```
remote > curl -s http://192.168.100.240
```

```
You've hit rs-labels-ksxlj
```

```
remote > curl -s http://192.168.100.240
```

```
You've hit rs-labels-wdln2
```

///

Kubernetes

Advanced Network

sessionAffinity (Sticky Session) - 1/2

- Environments (ReplicaSet으로 3쌍의 Pod를 생성하고 이를 LoadBalancer를 이용하여 Service하고 있는 환경을 준비하자)

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
```

```
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./03-week/ReplicaSet/rs-node-web.yaml
```

```
remote > kubectl create -f ./04-week/metallb/svc-lb-web.yaml
```

- 접속할 때마다 바꿔고 있는 Pod를 확인해 보자

```
remote > curl -s http://192.168.100.240
```

```
You've hit rs-labels-ksxlj
```

접속할 때마다 Pod가 바뀌는 것이 바람직할 때도 있겠지만,

```
remote > curl -s http://192.168.100.240
```

```
You've hit rs-labels-wdln2
```

원하지 않는 상황일 때도 있다.

Without Session Stickiness



With Session Stickiness



※ 참고 : <https://www.imperva.com/learn/availability/sticky-session-persistence-and-cookies/>

sessionAffinity (Sticky Session) - 2/2

- 세션이 유지되는 동안 같은 Pod와 연결이 되도록 하려면 다음과 같이 하면 된다

svc-lb-web-sa.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-lb-sa
spec:
  type: LoadBalancer
  ports:
  - name: http
    port: 80
    protocol: TCP
    targetPort: 8080
  selector:
    app: node-web
  sessionAffinity: ClientIP
  sessionAffinityConfig:
    clientIP:
      timeoutSeconds: 3600
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./04-week/advanced/svc-lb-web-sa.yaml
service/svc-lb-sa created

remote > kubectl get services -o wide
NAME         TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)       AGE   SELECTOR
kubernetes   ClusterIP   10.233.0.1    <none>          443/TCP      9d    <none>
svc-lb-sa    LoadBalancer 10.233.39.34  192.168.100.240  80:31032/TCP  5s    app=node-web

remote > curl -s http://192.168.100.240
You've hit rs-labels-98xnl
```

처음 불은 Pod로 계속 접속하게 된다.

```
remote > curl -s http://192.168.100.240
You've hit rs-labels-98xnl
```

'sessionAffinityConfig'

기본값은 10,800sec이다. 즉, 지정하지 않아도 되는 설정이다. (sessionAffinity 기본값은 None이다.)

///

externalTrafficPolicy - 1/3

- Environments (ReplicaSet으로 3쌍의 Pod를 생성하고 이를 NodePort를 이용하여 Service하고 있는 환경을 준비하자)

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./03-week/ReplicaSet/rs-node-web.yaml
remote > kubectl create -f ./03-week/Service/svc-node-web.yaml
```

- NodePort 접속 위치와 Pod 위치를 확인해보자

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	10d	<none>
svc-node	NodePort	10.233.28.186	<none>	80:30123/TCP	12s	app=node-web

```
remote > curl -s http://192.168.100.201:30123
```

```
You've hit rs-labels-z587n
```

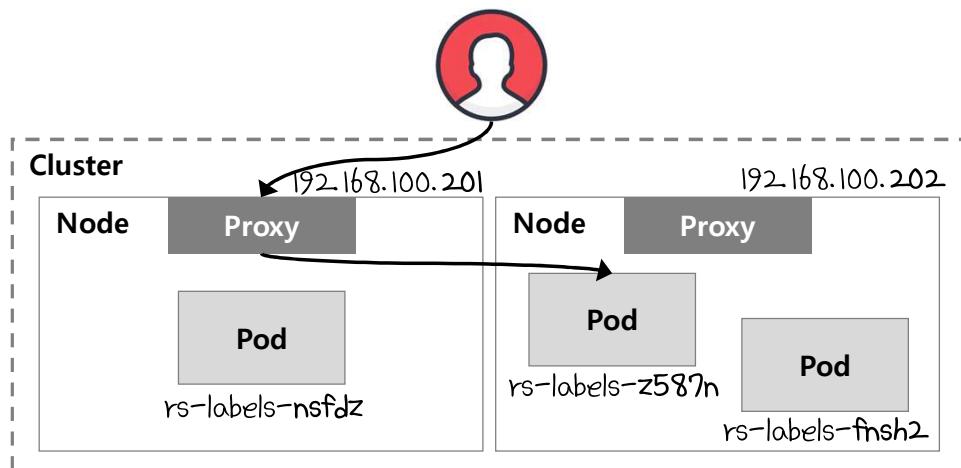
worker2 Node에 있는 Pod를 하나 확인하고,

worker1 Node의 IP로 접근하는 상황을 직접 확인해 보자.

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-labels-fnsh2	1/1	Running	0	64s	10.233.103.35	worker2	<none>	<none>
rs-labels-nsfdz	1/1	Running	0	64s	10.233.110.49	worker1	<none>	<none>
rs-labels-z587n	1/1	Running	0	64s	10.233.103.34	worker2	<none>	<none>

externalTrafficPolicy - 2/3



- 불필요한 네트워크 hop 발생
- 클라이언트 IP가 보존되지 않음

svc-lb-web-sa.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-node-et

spec:
  type: NodePort
  ports:
  - port: 80
    targetPort: 8080
    nodePort: 30123
  selector:
    app: node-web

  externalTrafficPolicy: Local
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl delete services svc-node

service "svc-node" deleted

remote > kubectl create -f ./04-week/advanced/svc-node-web-et.yaml

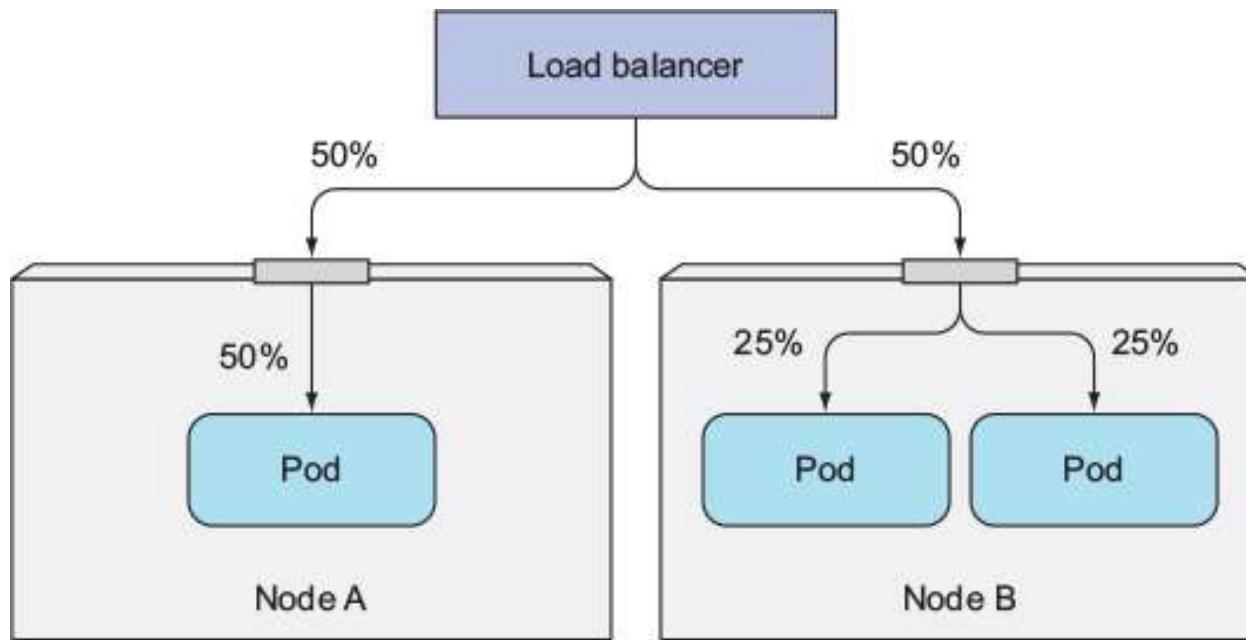
service/svc-node-et created

remote > curl -s http://192.168.100.201:30123
You've hit rs-labels-nsfdz
```

특정 Node로 들어오게 되면,
해당 Node에 있는 Pod를 선택하도록 하기 때문이다.

externalTrafficPolicy - 3/3

- LoadBalancer를 사용하면서 externalTrafficPolicy를 적용하게 되면, 오히려 균등 배부가 되지 않을 수도 있다.



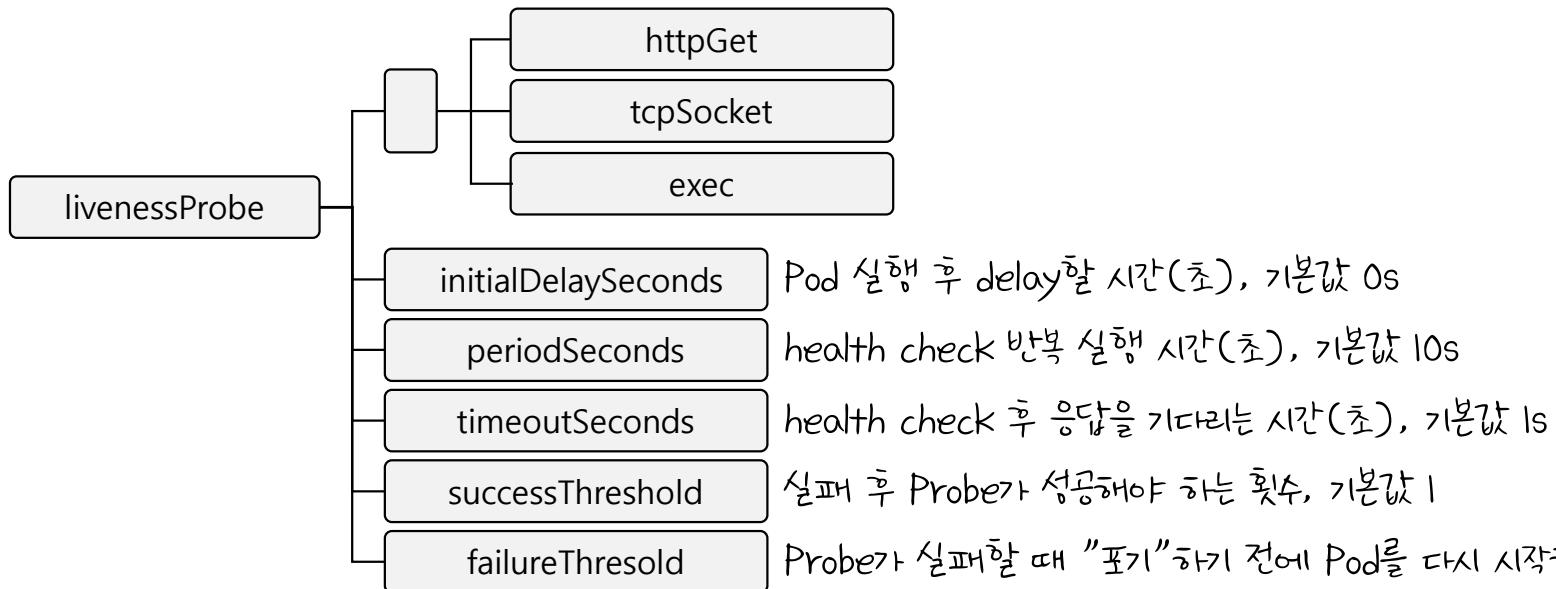
※ "externalTrafficPolicy: Local" 설정에 따른 문제점 : 균등히 배부되지 않을 수 있음

※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-5/230>

///

livenessProbe - 1/4

- 외부로부터 Request를 요청 받을 수 있는 상태인지 판별
- Pod의 spec에 각 container의 livenessProbe를 지정할 수 있음
 - . **HTTP GET Probe**: HTTP GET 요청의 응답 코드를 확인 (2xx, 3xx이면 성공)
 - . **TCP Socket Probe**: TCP 연결 성공 여부
 - . **Exec Probe**: 명령을 실행하고 exit code가 0이면 성공
- Probe가 실패하면 container를 재실행



```
apiVersion: v1
kind: Pod

metadata:
  name: kubia-liveness

spec:
  containers:
    - image: luksa/kubia-unhealthy
      name: kubia

      livenessProbe:
        httpGet:
          path: /
          port: 8080

        initialDelaySeconds: 15
        periodSeconds: 5
        timeoutSeconds: 1

        successThreshold: 1
        failureThresold: 3
```

livenessProbe - 2/4

- `luksa/kubia-unhealthy` 이미지에서 실행되는 node app의 내용을 좀 살펴보고 실습을 해보자.

app.js

```
const http = require('http');
const os = require('os');

console.log("Kubia server starting...");

var requestCount = 0;

var handler = function(request, response) {
    console.log("Received request from " + request.connection.remoteAddress);
    requestCount++;
    if (requestCount > 5) {
        response.writeHead(500);
        response.end("I'm not well. Please restart me!");
        return;
    }
    response.writeHead(200);
    response.end("You've hit " + os.hostname() + "\n");
};

var www = http.createServer(handler);
www.listen(8080);
```

pod-kubia-liveness.yaml

```
apiVersion: v1
kind: Pod

metadata:
  name: kubia-liveness

spec:
  containers:
    - image: luksa/kubia-unhealthy
      name: kubia

    livenessProbe:
      httpGet:
        path: /
        port: 8080

      initialDelaySeconds: 15
      periodSeconds: 5
      timeoutSeconds: 1
      successThreshold: 1
      failureThreshold: 3
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./04-week/advanced/pod-kubia-liveness.yaml
```

livenessProbe - 3/4

- Success가 될 때까지의 여정을 엿보자

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
kubia-liveness	1/1	Running	0	70s	10.233.103.36	worker2	<none>	<none>

```
remote > kubectl describe pod kubia-liveness
```

...

Events:

Type	Reason	Age	From	Message
Normal	Scheduled	2m1s	default-scheduler	Successfully assigned default/kubia-liveness to worker2
Normal	Pulled	92s	kubelet	Successfully pulled image "luksa/kubia-unhealthy" in 28.123275632s
Warning	Unhealthy	41s (x3 over 51s)	kubelet	Liveness probe failed: HTTP probe failed with statuscode: 500
Normal	Killing	41s	kubelet	Container kubia failed liveness probe, will be restarted
Normal	Pulling	11s (x2 over 2m)	kubelet	Pulling image "luksa/kubia-unhealthy"
Normal	Created	9s (x2 over 91s)	kubelet	Created container kubia
Normal	Started	9s (x2 over 91s)	kubelet	Started container kubia
Normal	Pulled	9s	kubelet	Successfully pulled image "luksa/kubia-unhealthy" in 1.984013259s

liveness 확인하고 fail로 판정되게 되면 restart 해버린다.

livenessProbe - 4/4

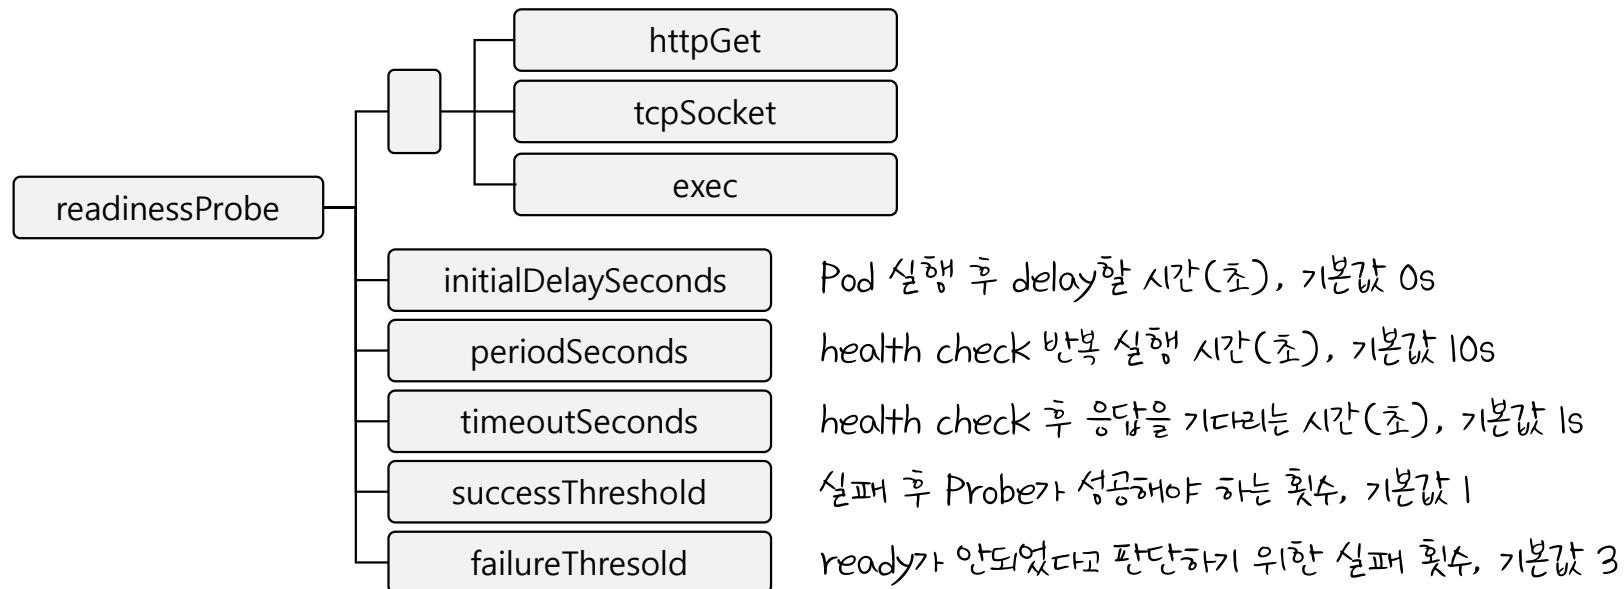
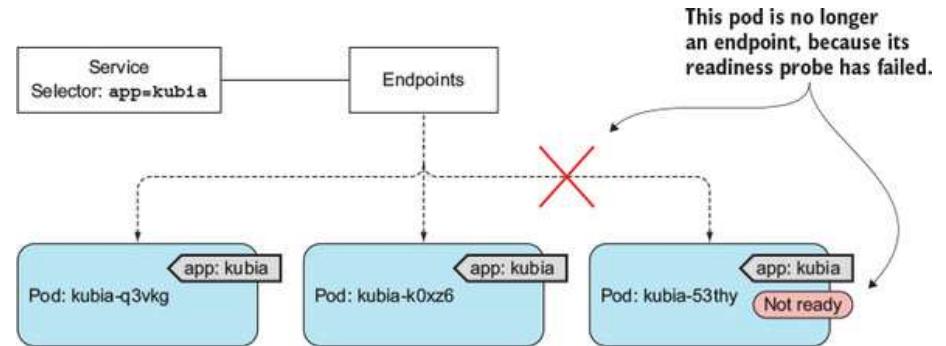
[효과적인 livenessProbe 구성]

- 특정 URL 경로(e.g, /health)에 요청하도록 Probe를 구성
 - . /health에서 모든 주요 구성 요소가 살아 있는지 확인하도록 로직 구성
- Probe는 application의 내부만 체크하고 외부 요인의 영향을 받지 않도록 해야 함
 - . database 장애로 인해 front-end의 Probe가 실패해서는 안된다
- Probe를 가볍게 유지
 - . 1초 내에 완료되어야 한다.
- Probe에 재시도 루프를 구현하지 마라 (=중복)
- container 레벨의 livenessProbe는 Worker Node의 Kubelet이 수행
 - . Master Node는 관여하지 않음
- Node crash로 인해 중단된 모든 Pod를 복구하는 것은 master의 몫

///

readinessProbe - 1/2

- Probe가 실패하면 Service에서 제거, 성공하면 다시 추가
 - . Client들이 healthy pod에만 요청을 할 수 있도록 관리
 - . 실패한다고 해서 Container 종료/재시작 하지 않음
- livenessProbe와 거의 동일한 사용법



※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-5/333>

readinessProbe - 2/2

rs-node-web-readiness.yaml

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-node-web

spec:
  replicas: 3

  selector:
    matchExpressions:
      - key: app
        operator: In
        values:
          - node-web

  template:
    metadata:
      labels:
        app: node-web

    spec:
      containers:
        - name: node-web
          image: whatwant/node-web:1.0
          ports:
            - containerPort: 8080

          readinessProbe:
            exec:
              command:
                - ls
                - /var/ready
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./04-week/advanced/rs-node-web-readiness.yaml

replicaset.apps/rs-node-web created

remote > kubectl get pods

NAME           READY   STATUS    RESTARTS   AGE
rs-node-web-cn8q4  0/1    Running   0          23s
rs-node-web-gj5b8  0/1    Running   0          23s
rs-node-web-nj8dm  0/1    Running   0          23s
```

```
remote > kubectl exec -it rs-node-web-cn8q4 -- touch /var/ready
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-node-web-cn8q4	1/1	Running	0	4m53s	10.233.110.56	worker1	<none>	<none>
rs-node-web-gj5b8	0/1	Running	0	4m53s	10.233.103.38	worker2	<none>	<none>
rs-node-web-nj8dm	0/1	Running	0	4m53s	10.233.103.37	worker2	<none>	<none>

///

Tips

k9s

- Kubernetes CLI To Manage Your Clusters In Style!
- LinuxBrew/HomeBrew Install

```
remote > sudo apt-get install build-essential curl file git
```

```
remote > sh -c "$(curl -fsSL https://raw.githubusercontent.com/Linuxbrew/install/master/install.sh)"
```

```
remote > nano ~/.zshrc
```

```
...
```

```
remote > source ~/.zshrc
```

```
~/.zshrc
```

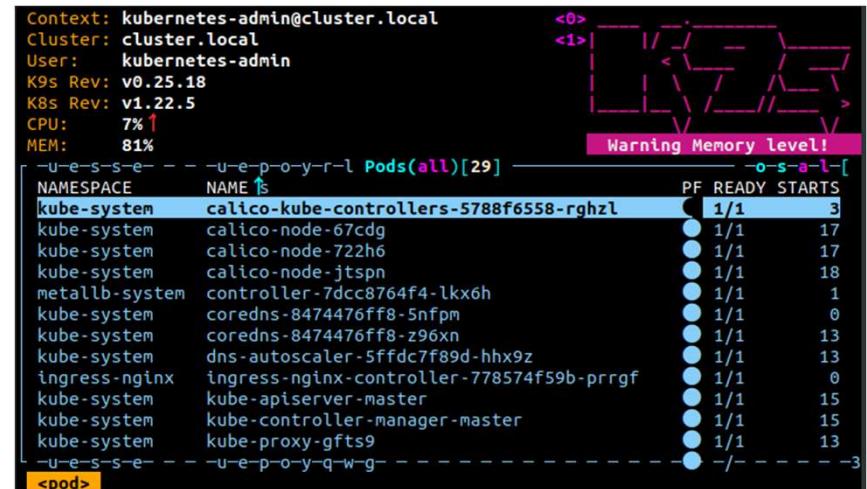
```
export PATH="/home/linuxbrew/.linuxbrew/bin:$PATH"
export MANPATH="/home/linuxbrew/.linuxbrew/share/man:$MANPATH"
export INFOPATH="/home/linuxbrew/.linuxbrew/share/info:$INFOPATH"
```

- k9s Install

```
remote > brew install derailed/k9s/k9s
```

```
remote > k9s
```

※ 참고 : <https://k9scli.io/>



///

<https://kahoot.it/>

[Score]

이민준 (6)

박남준 (3)

이혜정 (3)

김상호 (2)

정현찬 (1)

김정은 (1)

김남형 (1)