

**5<sup>th</sup>**  
**Week**

# 다섯번째 봅겠습니다 ?!

## ▷ 출석 체크도 한 번 해보시면 어떠세요?!

- <https://modulabs.co.kr/>
- 모두연 홈페이지 → 로그인 → 마이페이지 → 참여한 랩·풀잎 → 자세히 보기 → 내 풀잎스쿨 출석 확인하기

## ▷ Ground Rule

- 가급적 지각/결석 하지 않기
- 가급적 Camera 켜 놓고 수업 참여하기
- 가급적 적극적으로 참여하기
- 3시간이 넘더라도 배고프다고 화내지 않기
- Slack 잊지 않기
- 꼭 끝까지 함께하기

잡담 &  
지난 수업 관련 이야기

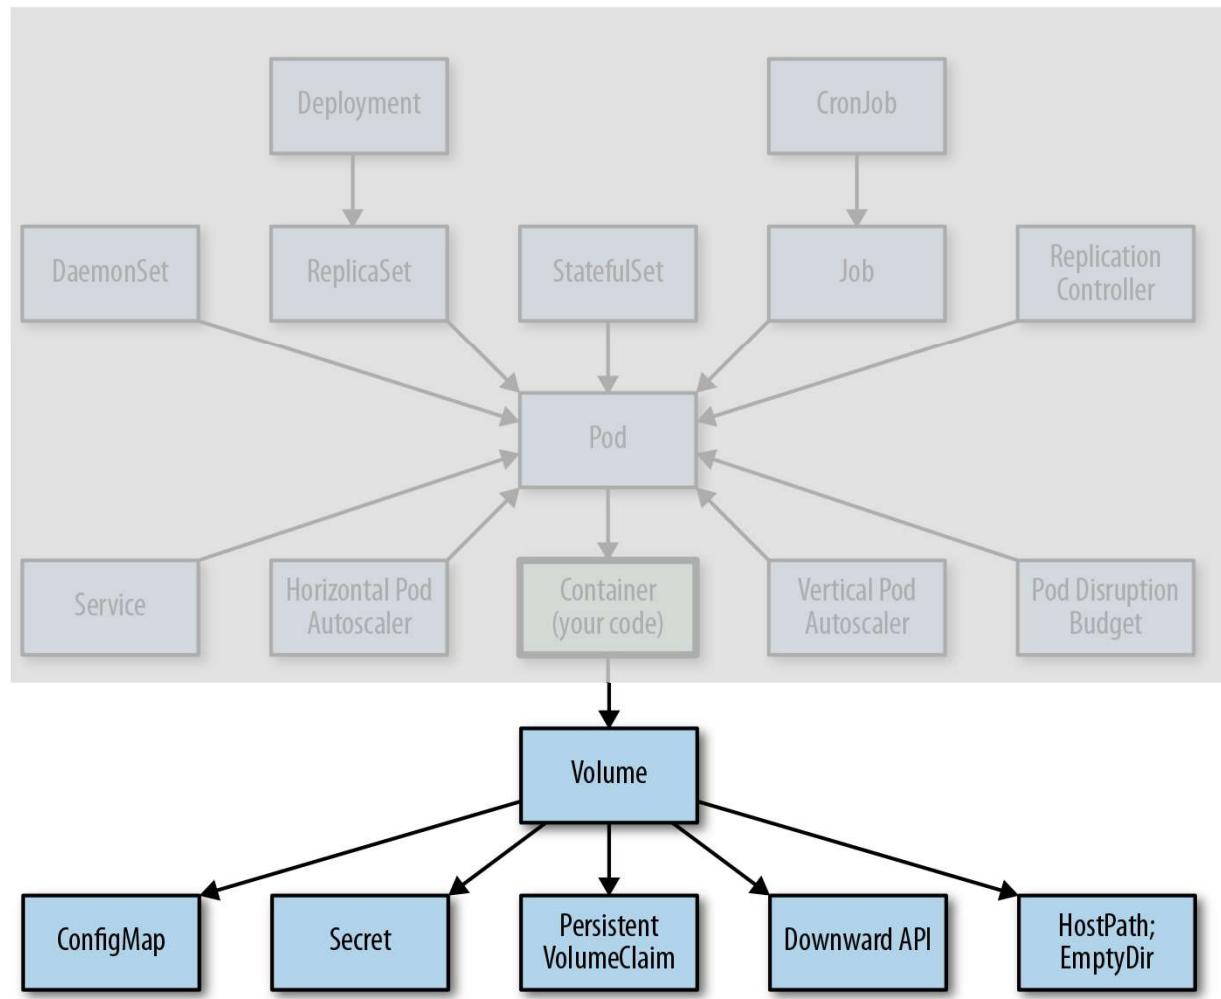
///

# Agenda

- ▷ [10m] make-friendship
- ▷ [30m] Flip - Volume (emptyDir/hostPath/PersistentVolume)
- ▷ [30m] Wrap-Up & Hands-On
- ▷ [10m] Break-Time
  
- ▷ [30m] Flip - Volume (configMap/Secret/downwardAPI)
- ▷ [30m] Wrap-Up & Hands-On
- ▷ [20m] REST API
  
- ▷ [20m] Quiz

# **Kubernetes**

## **Volume - Overview**



# kinds of volume

- "Kubernetes In Action" 책에서 설명하는 범위

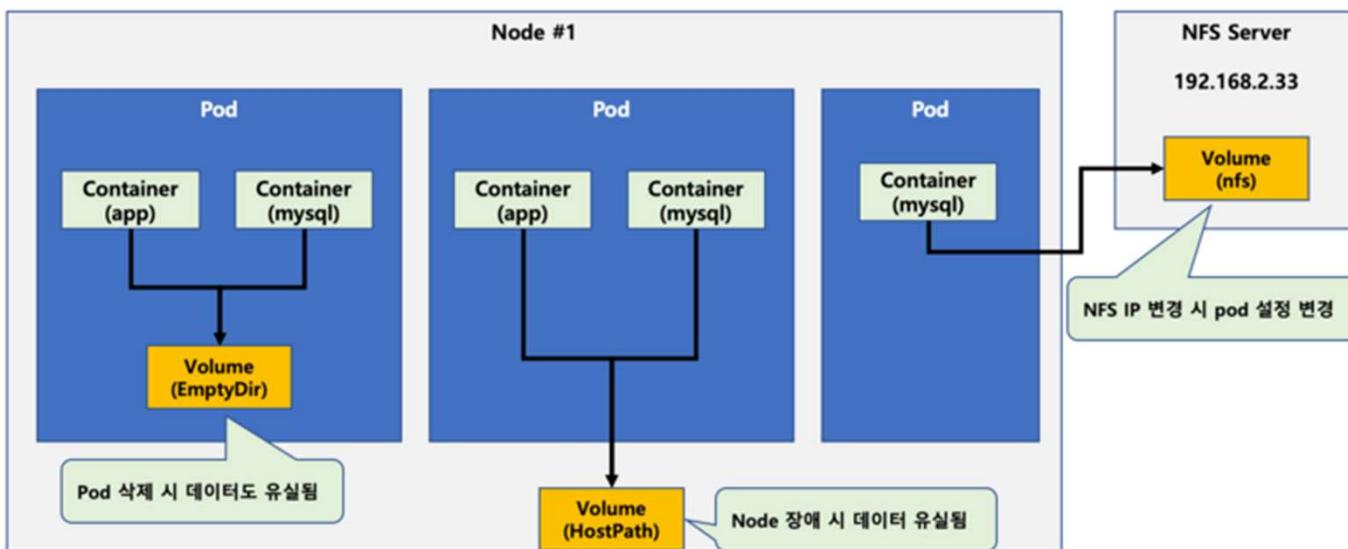
On-Premise 환경에서

별도 구축할 수 있는 solution

- 
- awsElasticBlockStore
  - azureDisk
  - azureFile
  - cephfs
  - cinder
  - configMap
  - csi
  - downwardAPI
  - emptyDir
  - fc (파이버 채널)
  - flexVolume
  - flocker
  - gcePersistentDisk
  - gitRepo (사용중단(deprecated))
  - glusterfs
  - hostPath
  - iscsi
  - local
  - nfs
  - persistentVolumeClaim
  - projected
  - portworxVolume
  - quobyte
  - rbd
  - scaleIO
  - secret
  - storageos
  - vsphereVolume

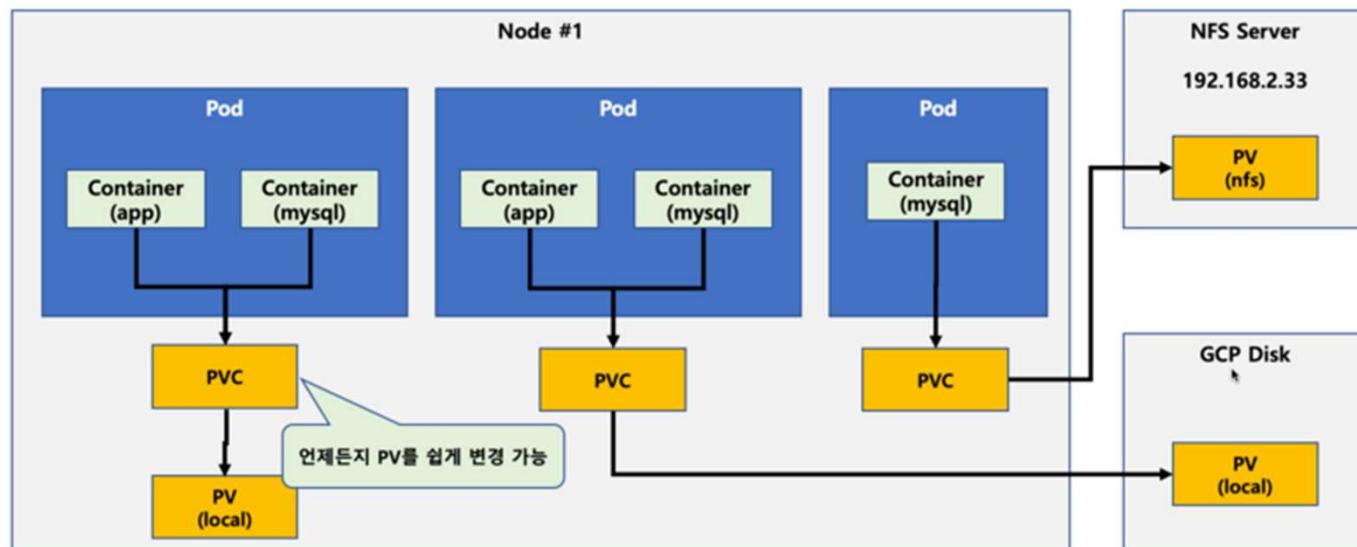
※ 참고 : <https://kubernetes.io/ko/docs/concepts/storage/volumes/>

# Pod에서 직접 Volume을 지정하는 방식



※ 참고 : <https://m.blog.naver.com/freepsw/222005161870>

# Pod에서 Persistent Volume을 활용하는 방식



※ 참고 : <https://m.blog.naver.com/freepsw/222005161870>

///

# **Flip Learning**

**(Volume - emptyDir/hostPath/PV)**

**이원준님**

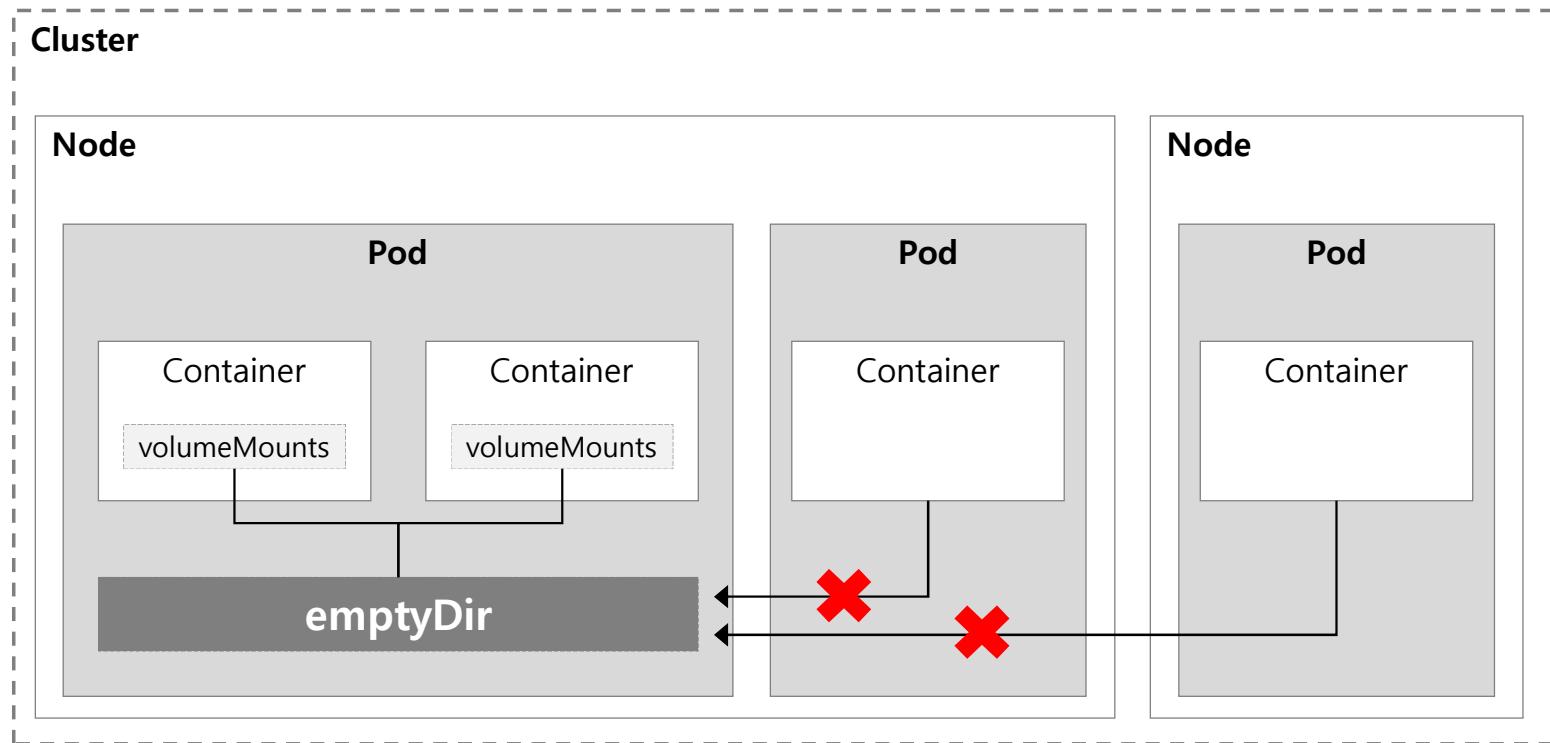
///

# Kubernetes

## Volume - emptyDir

# emptyDir

- Pod 안에서 Container끼리 공유 또는 Container가 재시작 하더라도 저장된 파일 유지
- Pod가 재시작 하는 경우에는 저장된 파일 유실



# Hands-On - 1/4

rs-emptydir.yaml

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-emptydir

spec:
  replicas: 1

  selector:
    matchLabels:
      app: ubuntu

  template:
    metadata:
      labels:
        app: ubuntu

    spec:
      containers:
        - image: ubuntu:20.04
          name: worker1
          command: ["/bin/sleep", "3650d"]

      volumeMounts:
        - name: emptydir-demo
          mountPath: /data/emptydir1
```

```
- image: ubuntu:20.04
  name: worker2
  command: ["/bin/sleep", "3650d"]

  volumeMounts:
    - name: emptydir-demo
      mountPath: /data/emptydir2

  volumes:
    - name: emptydir-demo
      emptyDir: {}
```

1개의 Pod 안에 2개의 container를 구성하고  
각 container에서 emptyDir을 mount 하도록 구성

# Hands-On - 2/4

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
```

```
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./05-week/emptyDir/rs-emptyDir.yaml
```

```
replicaset.apps/rs-emptydir created
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
rs-emptydir-wgfwx	2/2	Running	0	22s	10.233.103.40	worker2	<none>	<none>	

```
remote > kubectl exec -it rs-emptydir-wgfwx -c worker1 -- ls -al /data
```

```
total 12
drwxr-xr-x 3 root root 4096 Feb  2 15:44 .
drwxr-xr-x 1 root root 4096 Feb  2 15:44 ..
drwxrwxrwx 2 root root 4096 Feb  2 15:44 emptydir1
```

생성된 각 container에

mount된 directory를 확인!

```
remote > kubectl exec -it rs-emptydir-wgfwx -c worker2 -- ls -al /data
```

```
total 12
drwxr-xr-x 3 root root 4096 Feb  2 15:44 .
drwxr-xr-x 1 root root 4096 Feb  2 15:44 ..
drwxrwxrwx 2 root root 4096 Feb  2 15:44 emptydir2
```

# Hands-On - 3/4

```
remote > kubectl exec -it rs-emptydir-wgfwx -c worker1 -- ls -al /data/emptydir1
```

```
total 8  
drwxrwxrwx 2 root root 4096 Feb  2 14:44 .  
drwxr-xr-x 3 root root 4096 Feb  2 14:44 ..
```

```
remote > kubectl exec -it rs-emptydir-wgfwx -c worker2 -- ls -al /data/emptydir2
```

```
total 8  
drwxrwxrwx 2 root root 4096 Feb  2 14:44 .  
drwxr-xr-x 3 root root 4096 Feb  2 14:44 ..
```

```
remote > kubectl exec -it rs-emptydir-wgfwx -c worker1 -- touch /data/emptydir1/wow
```

```
remote > kubectl exec -it rs-emptydir-wgfwx -c worker1 -- ls -al /data/emptydir1
```

```
total 8  
drwxrwxrwx 2 root root 4096 Feb  2 15:21 .  
drwxr-xr-x 3 root root 4096 Feb  2 14:44 ..  
-rw-r--r-- 1 root root    0 Feb  2 15:21 wow
```

container 1번에서 생성한 파일이

container 2번에서도 공유되고 있음을 확인

```
remote > kubectl exec -it rs-emptydir-wgfwx -c worker2 -- ls -al /data/emptydir2
```

```
total 8  
drwxrwxrwx 2 root root 4096 Feb  2 15:21 .  
drwxr-xr-x 3 root root 4096 Feb  2 14:44 ..  
-rw-r--r-- 1 root root    0 Feb  2 15:21 wow
```

# Hands-On - 4/4

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-emptydir-wgfwx	2/2	Running	0	10m	10.233.103.40	worker2	<none>	<none>

```
remote > kubectl exec -it rs-emptydir-wgfwx -c worker1 -- ls -al /data/emptydir1
```

```
total 8
drwxrwxrwx 2 root root 4096 Feb  2 15:53 .
drwxr-xr-x 3 root root 4096 Feb  2 15:44 ..
-rw-r--r-- 1 root root    0 Feb  2 15:53 wow
```

Pod가 재시작 되었을 때,

```
remote > kubectl delete pods rs-emptydir-wgfwx
```

emptyDir 내용이 유지가 되는지를 보기 위한 실습이다.

```
pod "rs-emptydir-wgfwx" deleted
```

당연히, 유지가 안된다.

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-emptydir-5s8l8	2/2	Running	0	61s	10.233.103.41	worker2	<none>	<none>

```
remote > kubectl exec -it rs-emptydir-5s8l8 -c worker1 -- ls -al /data/emptydir1
```

```
total 8
drwxrwxrwx 2 root root 4096 Feb  2 15:55 .
drwxr-xr-x 3 root root 4096 Feb  2 15:55 ..
```

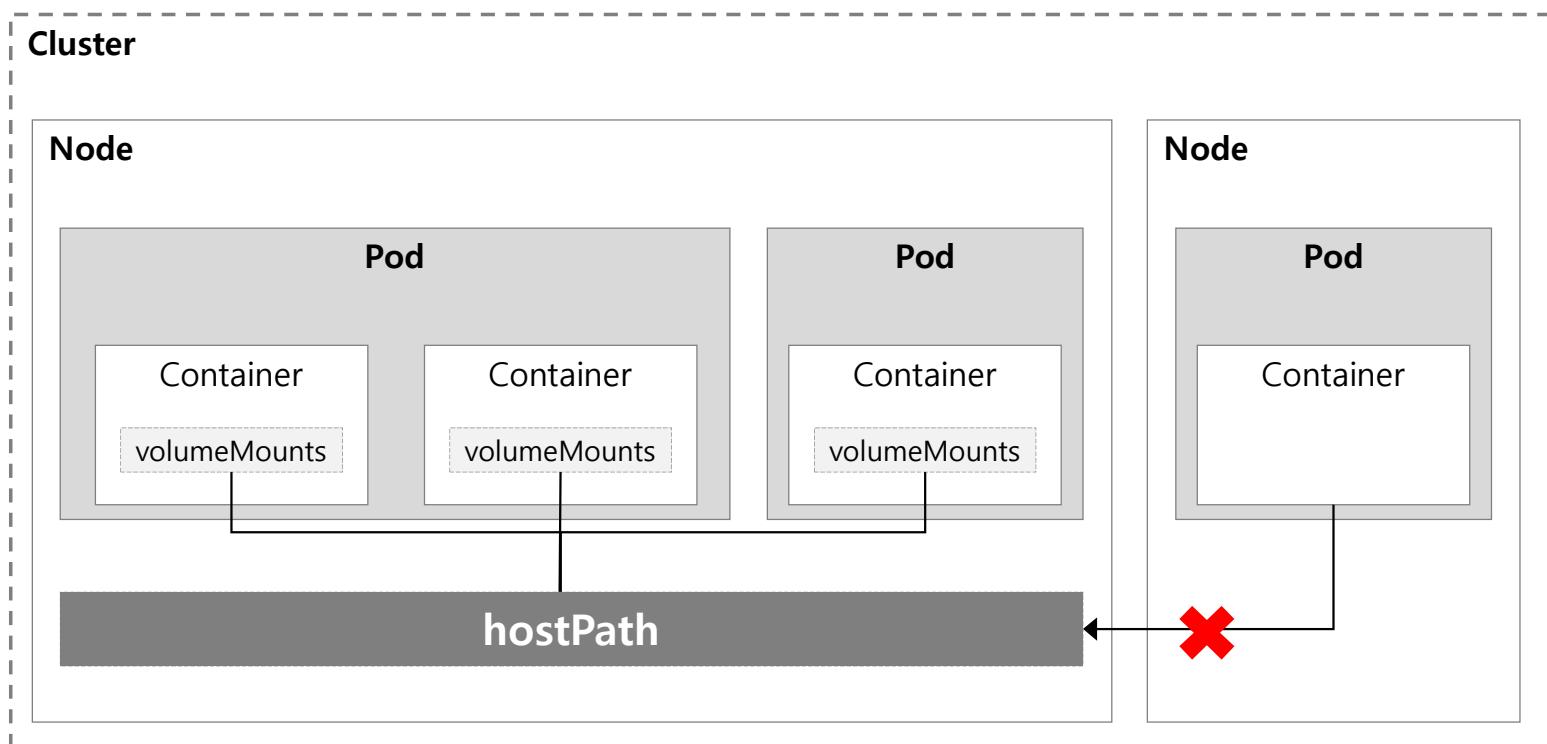
///

# **Kubernetes**

## **Volume - hostPath**

# hostPath

- 같은 Node에 있는 Pod끼리 서로 공유할 수 있고, 다른 Node에서는 참조할 수 없다
- Pod가 재시작 되더라도 hostPath에 저장된 내용은 유지 된다



# Type of `hostPath`

값	행동
	빈 문자열 (기본값)은 이전 버전과의 호환성을 위한 것으로, hostPath 볼륨은 마운트 하기 전에 아무런 검사도 수행되지 않는다.
DirectoryOrCreate	만약 주어진 경로에 아무것도 없다면, 필요에 따라 Kubelet이 가지고 있는 동일한 그룹과 소유권, 권한을 0755로 설정한 빈 디렉터리를 생성한다.
Directory	주어진 경로에 디렉터리가 있어야 함
FileOrCreate	만약 주어진 경로에 아무것도 없다면, 필요에 따라 Kubelet이 가지고 있는 동일한 그룹과 소유권, 권한을 0644로 설정한 빈 디렉터리를 생성한다.
File	주어진 경로에 파일이 있어야 함
Socket	주어진 경로에 UNIX 소켓이 있어야 함
CharDevice	주어진 경로에 문자 디바이스가 있어야 함
BlockDevice	주어진 경로에 블록 디바이스가 있어야 함

※ 참고 : <https://kubernetes.io/ko/docs/concepts/storage/volumes/>

# Hands-On - 1/5

rs-hostPath.yaml

```
apiVersion: apps/v1
kind: ReplicaSet

metadata:
  name: rs-hostpath

spec:
  replicas: 3

  selector:
    matchLabels:
      app: ubuntu

  template:
    metadata:
      labels:
        app: ubuntu

    spec:
      containers:
        - image: ubuntu:20.04
          name: worker1
          command: ["/bin/sleep", "3650d"]

      volumeMounts:
        - name: hostpath-demo
          mountPath: /data/hostpath
```

```
volumes:
  - name: hostpath-demo
    hostPath:
      path: /tmp/hostpath-demo
      type: DirectoryOrCreate
```

hostPath을 mount 하도록 구성한 Pod 3개를 생성

# Hands-On - 2/5

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git  
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./05-week/hostPath/rs-hostPath.yaml
```

```
replicaset.apps/rs-hostpath created
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
rs-hostpath-2xpmw	1/1	Running	0	13s	10.233.110.57	worker1	<none>	<none>	
rs-hostpath-nlgdh	1/1	Running	0	13s	10.233.103.42	worker2	<none>	<none>	
rs-hostpath-pddz5	1/1	Running	0	13s	10.233.103.43	worker2	<none>	<none>	

```
remote > kubectl exec -it rs-hostpath-2xpmw -- ls -al /data
```

```
total 12  
drwxr-xr-x 3 root root 4096 Feb  2 16:45 .  
drwxr-xr-x 1 root root 4096 Feb  2 16:45 ..  
drwxr-xr-x 2 root root 4096 Feb  2 16:45 hostpath
```

서로 다른 Node에 있는 Pod 모두

일단 hostPath가 mount되어 있는 것을 확인할 수 있다

```
remote > kubectl exec -it rs-hostpath-nlgdh -- ls -al /data
```

```
total 12  
drwxr-xr-x 3 root root 4096 Feb  2 16:45 .  
drwxr-xr-x 1 root root 4096 Feb  2 16:45 ..  
drwxr-xr-x 2 root root 4096 Feb  2 16:45 hostpath
```

# Hands-On - 3/5

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-hostpath-2xpmw	1/1	Running	0	13s	10.233.110.57	worker1	<none>	<none>
rs-hostpath-nlgdh	1/1	Running	0	13s	10.233.103.42	worker2	<none>	<none>
rs-hostpath-pddz5	1/1	Running	0	13s	10.233.103.43	worker2	<none>	<none>

```
remote > kubectl exec -it rs-hostpath-nlgdh -- touch /data/hostpath-wow
```

Worker2에 있는 Pod에서 파일을 하나 생성했고,  
실제로 확인해보면 잘 생성되어 있다.

```
remote > kubectl exec -it rs-hostpath-nlgdh -- ls -al /data/hostpath
```

```
total 8
drwxr-xr-x 2 root root 4096 Feb  2 16:57 .
drwxr-xr-x 3 root root 4096 Feb  2 16:45 ..
-rw-r--r-- 1 root root    0 Feb  2 16:57 wow
```

```
remote > kubectl exec -it rs-hostpath-2xpmw -- ls -al /data/hostpath
```

Worker1에 있는 Pod에서는

파일이 안보이지만

Worker2에 있는 다른 Pod에서는

파일이 잘 보인다.

```
total 8
drwxr-xr-x 2 root root 4096 Feb  2 16:57 .
drwxr-xr-x 3 root root 4096 Feb  2 16:45 ..
-rw-r--r-- 1 root root    0 Feb  2 16:57 wow
```

```
remote > kubectl exec -it rs-hostpath-pddz5 -- ls -al /data/hostpath
```

```
total 8
drwxr-xr-x 2 root root 4096 Feb  2 16:57 .
drwxr-xr-x 3 root root 4096 Feb  2 16:45 ..
-rw-r--r-- 1 root root    0 Feb  2 16:57 wow
```

# Hands-On - 4/5

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-hostpath-2xpmw	1/1	Running	0	30m	10.233.110.57	worker1	<none>	<none>
rs-hostpath-nlgdh	1/1	Running	0	30m	10.233.103.42	worker2	<none>	<none>
rs-hostpath-pddz5	1/1	Running	0	30m	10.233.103.43	worker2	<none>	<none>

```
remote > kubectl exec -it rs-hostpath-pddz5 -- ls -al /data/hostpath
```

```
total 8
drwxr-xr-x 2 root root 4096 Feb  2 16:57 .
drwxr-xr-x 3 root root 4096 Feb  2 16:45 ..
-rw-r--r-- 1 root root    0 Feb  2 16:57 wow
```

기존 Pod가 삭제되고, 다른 Pod가 생성되었음에도

```
remote > kubectl delete pods rs-hostpath-pddz5
```

hostPath 내용이 유지가 되는지를 보기 위한 실습이다.

```
pod "rs-hostpath-pddz5" deleted
```

당연히, 유지가 된다.

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
rs-hostpath-2xpmw	1/1	Running	0	31m	10.233.110.57	worker1	<none>	<none>
rs-hostpath-5v9r8	1/1	Running	0	42s	10.233.103.44	worker2	<none>	<none>
rs-hostpath-nlgdh	1/1	Running	0	31m	10.233.103.42	worker2	<none>	<none>

```
remote > kubectl exec -it rs-hostpath-5v9r8 -- ls -al /data/hostpath
```

```
total 8
drwxr-xr-x 2 root root 4096 Feb  2 16:57 .
drwxr-xr-x 3 root root 4096 Feb  2 17:16 ..
-rw-r--r-- 1 root root    0 Feb  2 16:57 wow
```

# Hands-On - 5/5

```
remote > ssh vagrant@192.168.100.201
```

```
worker1 > ls -al /tmp/hostPath-demo
```

```
total 8  
drwxr-xr-x 2 root root 4096 2월  3 01:57 .  
drwxrwxrwt 12 root root 4096 2월  3 02:25 ..
```

worker1 Node에서도 hostPath로 지정한 디렉토리가 생성이 되었다.

하지만, 파일 내용은 없다.

```
worker1 > exit
```

```
remote > ssh vagrant@192.168.100.202
```

```
worker2 > ls -al /tmp/hostpath-demo
```

```
total 8  
drwxr-xr-x 2 root root 4096 2월  3 01:57 .  
drwxrwxrwt 12 root root 4096 2월  3 02:26 ..  
-rw-r--r-- 1 root root    0 2월  3 01:57 wow
```

worker2 Node에서도 hostPath로 지정한 디렉토리가 생성이 되었다.

worker2 Node에 있는 Pod에서 파일을 생성했기에

여기에서는 파일이 보인다.

hostPath의 내용은 해당 Node host에서 그대로 확인이 가능하다.

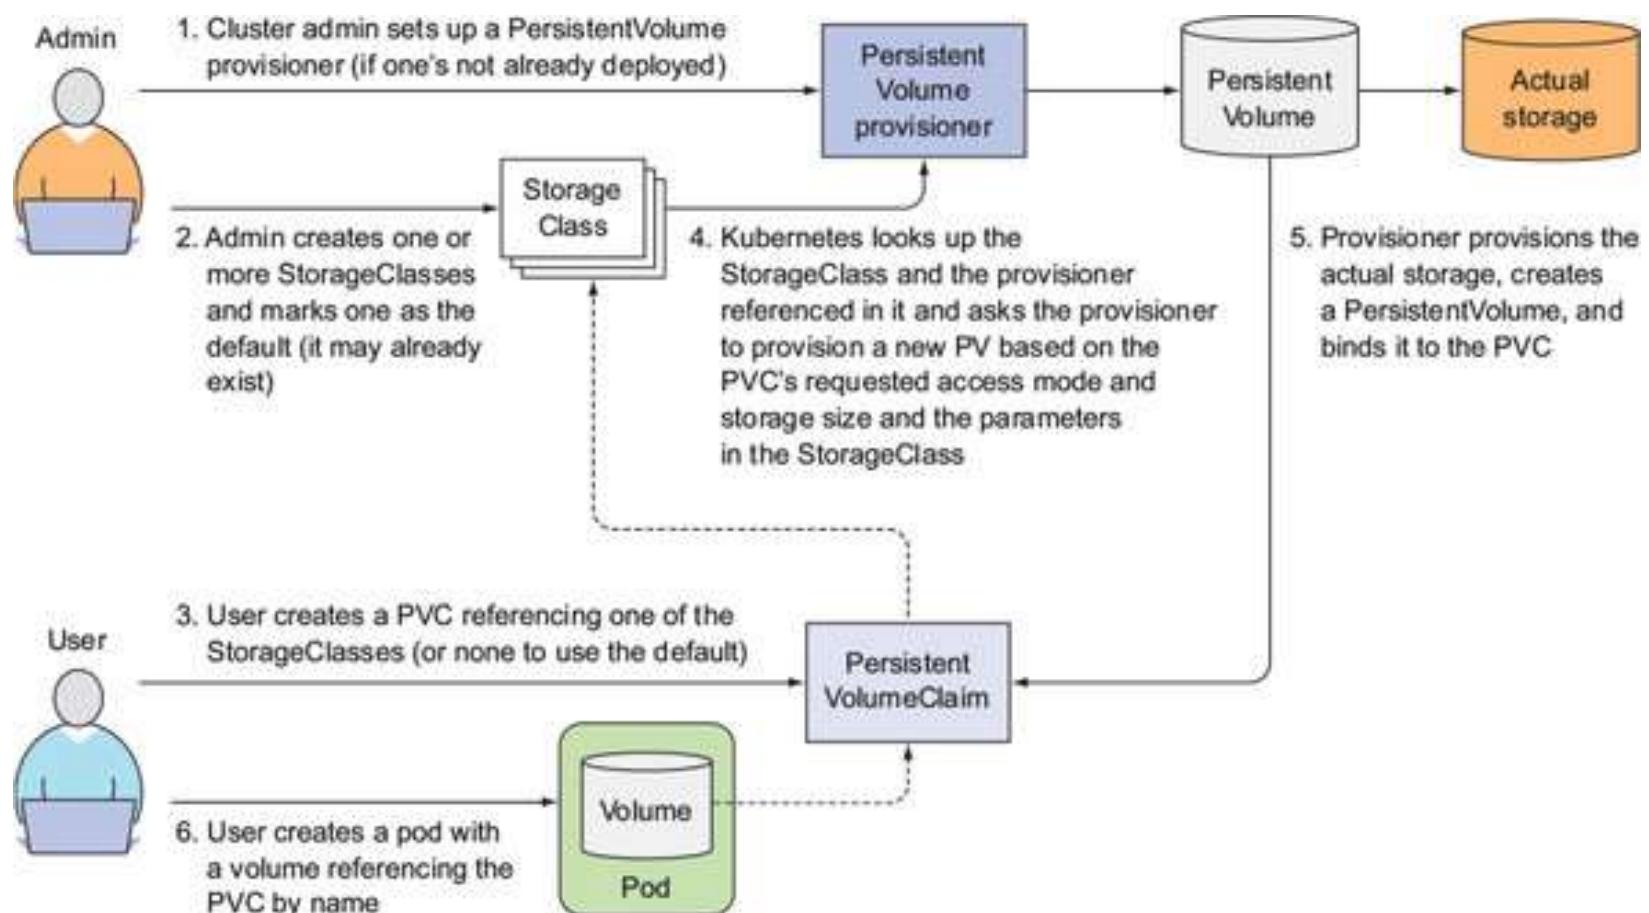
다르게 생각해보면, host에서 Pod로 파일을 공유할 수 있는 방법이 될 수도 있다.

반면, 쓸모 없어진 hostPath 내용이 자동으로 삭제되지도 않고, 이름이 중복되어 잘못 노출될 수 있는 위험도 있고,

같은 ReplicaSet에 속한 Pod인데도 위치한 Node에 따라 다른 내용이 저장된다는 문제도 있다.

///

# Provisioning of PersistentVolumes



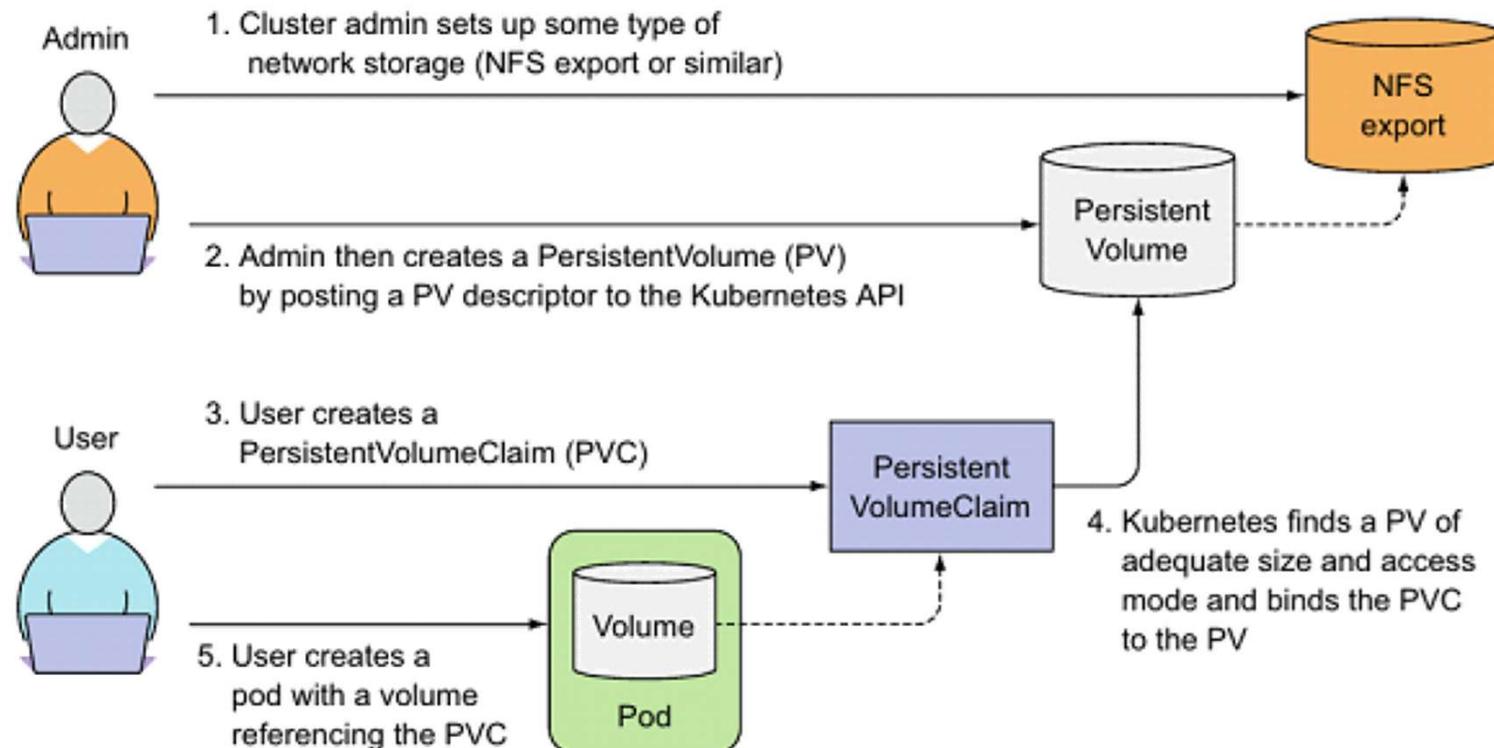
※ 참고 : <https://livebook.manning.com/concept/kubernetes/persistentvolumes>

///

# **Kubernetes**

## **Volume - Static Provisioning (PV & PVC)**

# PersistentVolume(PV) & PersistentVolumeClaim(PVC)



▲ 그림 6.6 클러스터 관리자가 퍼시스턴트볼륨을 프로비저닝하면 파드는 퍼시스턴트볼륨클레임을 통해 이를 사용 한다.

※ 참고 : <https://thenewstack.io/strategies-running-stateful-applications-kubernetes-persistent-volumes-claims/>

# PersistentVolume (hostPath)

- Volume 사용을 위해서는 물리적인 저장 공간이 필요 → 지금은 가장 기본적인 hostPath를 이용해서 살펴보자

pv-hostpath.yaml

```
apiVersion: v1
kind: PersistentVolume

metadata:
  name: pv-hostpath
  labels:
    type: local

spec:
  storageClassName: manual
  persistentVolumeReclaimPolicy: Retain

  capacity:
    storage: 100Mi

  accessModes:
    - ReadWriteOnce

  hostPath:
    path: "/tmp/pv-data"
    type: DirectoryOrCreate
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./05-week/persistentvolume/pv-hostpath.yaml
persistentvolume/pv-hostpath created
```

```
remote > kubectl get persistentvolumes -o wide
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE	VOLUME MODE
pv-hostpath	100Mi	RWO	Retain	Available		manual		13s	Filesystem

Reclaim Policy (반환 정책)

구분	설명
Retain	수동 반환 (default)
Delete	삭제
Recycle	Deprecated

accessModes

구분	설명
ReadWriteOnce	RWO
ReadOnlyMany	ROX
ReadWriteMany	RWX

※ 참고 : <https://kubernetes.io/ko/docs/tasks/configure-pod-container/configure-persistent-volume-storage/>

# PersistentVolumeClaim

- PersistentVolume을 요청하는 주문서라고 생각하면 된다

pvc-50mi.yaml

```
apiVersion: v1
kind: PersistentVolumeClaim

metadata:
  name: pvc-50mi

spec:
  storageClassName: manual

  accessModes:
    - ReadWriteOnce

  resources:
    requests:
      storage: 50Mi
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./05-week/persistentvolume/pvc-50mi.yaml
```

```
persistentvolumeclaim/pvc-50mi created
```

```
remote > kubectl get persistentvolumeclaims -o wide
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE	VOLUMEMODE
pvc-50mi	Bound	pv-hostpath	100Mi	RWO	manual	6s	Filesystem

```
remote > kubectl get persistentvolumes -o wide
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE	VOLUMEMODE
pv-hostpath	100Mi	RWO	Retain	Bound	default/pvc-50mi	manual		6h15m	Filesystem

50Mi 용량으로 주문을 했는데, 주문 내용과 맞는 PV를 찾아보니 100Mi 용량짜리가 있어서 그걸로 matching 된 것이다.

※ 참고 : <https://kubernetes.io/ko/docs/tasks/configure-pod-container/configure-persistent-volume-storage/>

# volumeMounts

rs-pvc-static.yaml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-pvc-static

spec:
  replicas: 1
  selector:
    matchLabels:
      app: ubuntu

  template:
    metadata:
      labels:
        app: ubuntu
    spec:
      containers:
        - image: ubuntu:20.04
          name: ubuntu
          command: ["/bin/sleep", "3650d"]

        volumeMounts:
          - name: pv-claim
            mountPath: /data/pv

  volumes:
    - name: pv-claim
      persistentVolumeClaim:
        claimName: pvc-50mi
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./05-week/persistentvolume/rs-pvc-static.yaml
replicaset.apps/rs-pvc-static created

remote > kubectl get replicsets -o wide
NAME           DESIRED   CURRENT   READY   AGE   CONTAINERS   IMAGES   SELECTOR
rs-pvc-static   1         1         1       15s   ubuntu        ubuntu:20.04   app=ubuntu

remote > kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
rs-pvc-static-7xv6m   1/1     Running   0          23s   10.233.103.47   worker2   <none>   <none>

remote > kubectl exec -it rs-pvc-static-7xv6m -- touch /data/pv-wow

remote > ssh whatwant@192.168.100.202
worker2 > ls -al /tmp/pv-data
total 8
drwxr-xr-x  2 root root 4096  2월  6 21:50 .
drwxrwxrwt 13 root root 4096  2월  6 21:50 ..
-rw-r--r--  1 root root    0  2월  6 21:50 wow
```

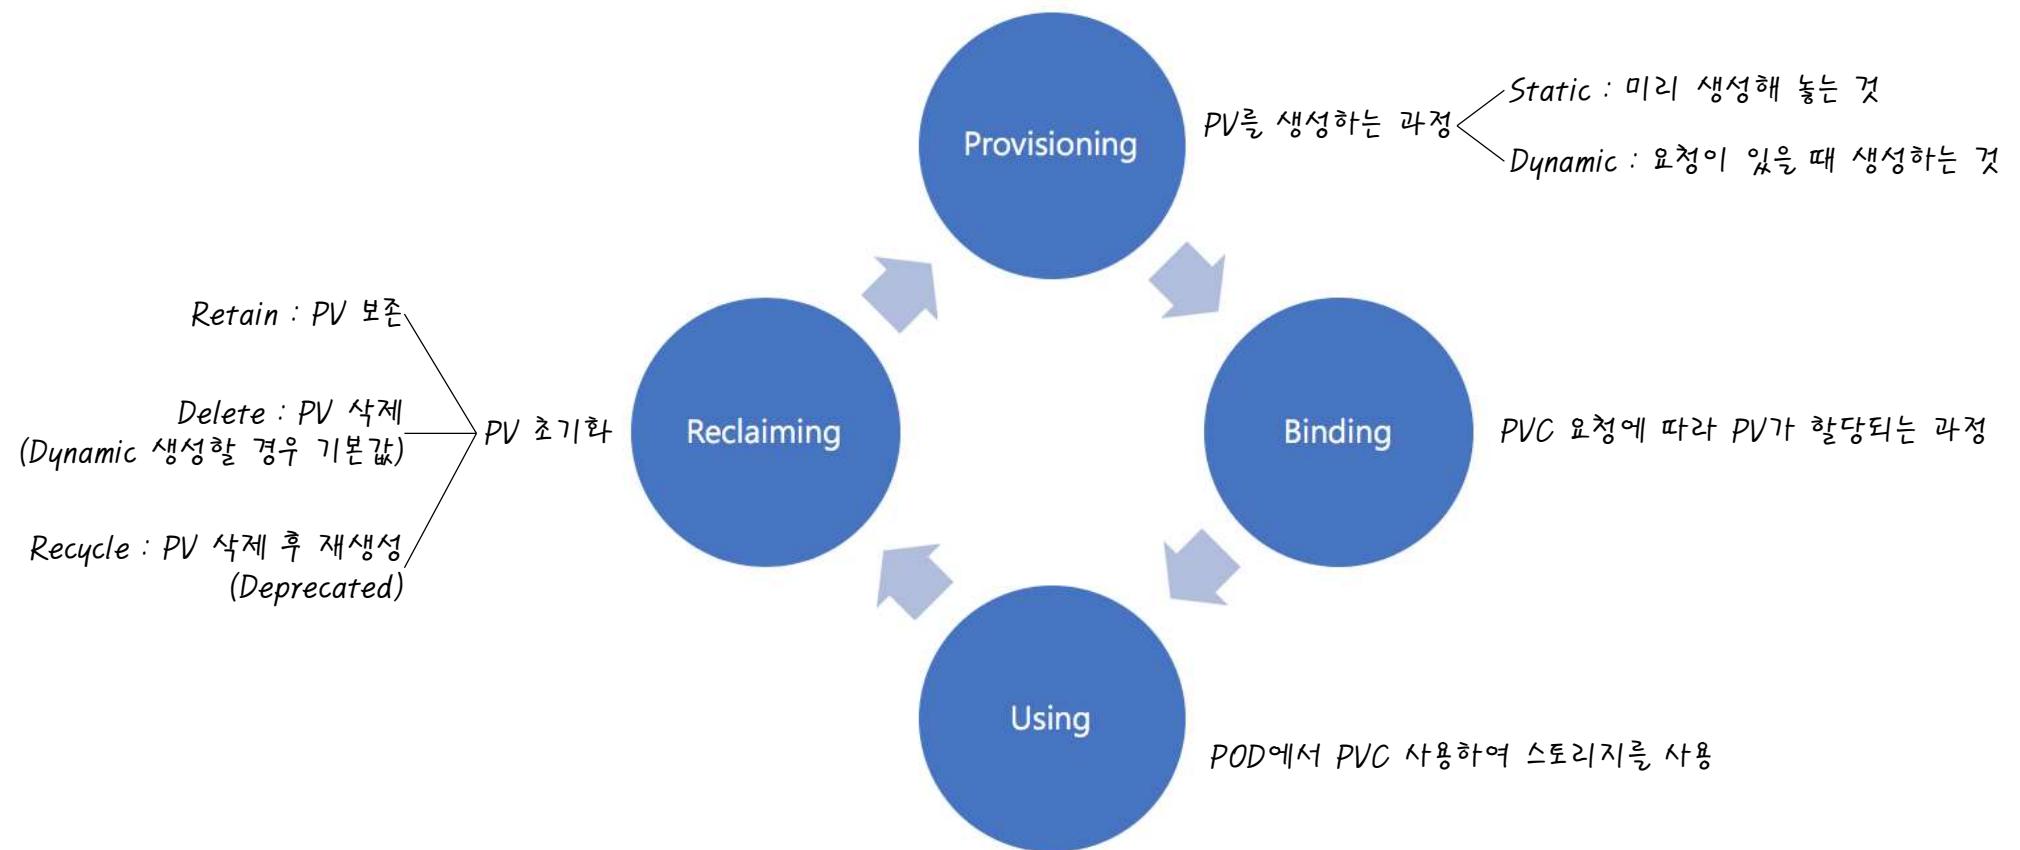
claimName 지정만으로 PV와 연결이 된다.

///

# **Kubernetes**

**Volume -  
Dynamic Provisioning  
(StorageClass)**

# “PV & PVC” Life Cycle (생명주기)



※ 참고 : <https://arisu1000.tistory.com/27849>

# StorageClass (local) & Ready Storage

- local storage를 사용하는 경우, provisioner를 no-provisioner로 지정하면 된다

sc-local.yaml

```
apiVersion: storage.k8s.io/v1
kind: StorageClass

metadata:
  name: sc-local

provisioner: kubernetes.io/no-provisioner

reclaimPolicy: Retain

volumeBindingMode: WaitForFirstConsumer
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./05-week/persistentvolume/sc-local.yaml

storageclass.storage.k8s.io/sc-local created

remote > kubectl get storageclasses -o wide
NAME      PROVISIONER          RECLAIMPOLICY   VOLUMEBINDINGMODE   ALLOWVOLUMEEXPANSION   AGE
sc-local  kubernetes.io/no-provisioner  Retain        WaitForFirstConsumer  false                  32s
```

- 실제 파일 저장을 할 물리적인 공간을 확보해야 한다. 여기에서는 worker1 Node에 준비를 해봤다.

```
remote > ssh whatwant@192.168.100.201
worker1 > mkdir -p /tmp/pv-local
worker1 > chmod 777 /tmp/pv-local
```

자꾸 /tmp 디렉토리를 이용하는 이유는

해당 Node가 재시작 하면 자동으로 해당 내용들을 삭제해주기 때문이다.

만약, Node가 재시작 되더라도 유지되게 하고 싶다면, 다른 디렉토리를 이용해야 한다.

# PersistentVolume (local)

- 지금 만들어야 하는 PV는 앞에서 선언한 StorageClass와 미리 확보해 놓은 저장 공간을 사용하도록 해야 한다

pv-local.yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-local

spec:
  capacity:
    storage: 100Mi

  accessModes:
  - ReadWriteOnce

  persistentVolumeReclaimPolicy: Retain

  storageClassName: sc-local

  local:
    path: /tmp/pv-local

  nodeAffinity:
    required:
      nodeSelectorTerms:
      - matchExpressions:
        - key: kubernetes.io/hostname
          operator: In
          values:
          - worker1
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./05-week/persistentvolume/pv-local.yaml
persistentvolume/pv-local created

remote > kubectl get persistentvolumes -o wide
NAME      CAPACITY   ACCESS MODES  RECLAIM POLICY  STATUS    CLAIM     STORAGECLASS  REASON   AGE   VOLUMEMODE
pv-local  100Mi      RWO          Retain         Available  sc-local
                                     sc-local
                                     8s   Filesystem

remote > kubectl get storageclasses -o wide
NAME      PROVISIONER           RECLAIMPOLICY  VOLUMEBINDINGMODE  ALLOWVOLUMEEXPANSION  AGE
sc-local  kubernetes.io/no-provisioner  Retain        WaitForFirstConsumer  false
                                                 sc-local
                                                 31m
```

앞에서 미리 만들어 놓은 저장 공간이

Worker1 Node의 '/tmp/pv-local'이었으므로

'nodeAffinity'를 이용해서 생성 Node를 제한한다.

PV는 아직 생성되지 않았다.

# PersistentVolumeClaim

- Static Provisioning과는 분명한 차이가 있다

pvc-local.yaml

```
apiVersion: v1
kind: PersistentVolumeClaim

metadata:
  name: pvc-local

spec:
  storageClassName: sc-local

  accessModes:
    - ReadWriteOnce

resources:
  requests:
    storage: 50Mi
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./05-week/persistentvolume/pvc-local.yaml
```

```
persistentvolumeclaim/pvc-local created
```

```
remote > kubectl get persistentvolumeclaims -o wide
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE	VOLUMEMODE
pvc-local	Pending				sc-local	17s	Filesystem

```
remote > kubectl get persistentvolumes -o wide
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE	VOLUMEMODE
pv-local	100Mi	RWO	Retain	Available		sc-local		11m	Filesystem

PVC를 생성했음에도 PV와 Binding 되지 않았다.

PV는 아직도 Available 상태이고, PVC는 Pending 상태이다.

# volumeMounts

rs-pvc-dynamic.yaml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-pvc-dynamic

spec:
  replicas: 1
  selector:
    matchLabels:
      app: ubuntu

  template:
    metadata:
      labels:
        app: ubuntu

    spec:
      containers:
        - image: ubuntu:20.04
          name: ubuntu
          command: ["/bin/sleep", "3650d"]

      volumeMounts:
        - name: pv-claim
          mountPath: /data/pv

  volumes:
    - name: pv-claim
      persistentVolumeClaim:
        claimName: pvc-local
```

```
remote > kubectl create -f ./05-week/persistentvolume/rs-pvc-dynamic.yaml
replicaset.apps/rs-pvc-dynamic created

remote > kubectl get persistentvolumeclaims -o wide
NAME           STATUS  VOLUME   CAPACITY  ACCESS MODES  STORAGECLASS  AGE  VOLUMEMODE
pvc-local      Bound   pv-local  100Mi     RWO          sc-local     12m  Filesystem

remote > kubectl get persistentvolumes -o wide
NAME       CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM           STORAGECLASS  REASON  AGE  VOLUMEMODE
pv-local   100Mi     RWO          Retain         Bound   default/pvc-local  sc-local     23m  Filesystem

remote > kubectl get pods -o wide
NAME           READY  STATUS    RESTARTS  AGE   IP           NODE  NOMINATED NODE  READINESS GATES
rs-pvc-dynamic-htrvf  1/1   Running   0          97s  10.233.110.63  worker1  <none>        <none>

remote > kubectl exec -it rs-pvc-dynamic-htrvf -- touch /data/pv-wow
remote > ssh whatwant@192.168.100.201
worker1 > ls -al /tmp/pv-local
total 8
drwxrwxrwx  2 whatwant whatwant 4096 2월  6 23:45 .
drwxrwxrwt 13 root    root     4096 2월  6 23:48 ..
-rw-r--r--  1 root    root      0  2월  6 23:45 wow
```

Pod가 생성되어서 volume을 요청하는 시점에  
PVC ~ PV Binding이 이루어진다.

///

# **Break**

///

# **Flip Learning**

**(Volume - ConfigMap/Secret/downwardAPI)**

**김남형님**

///

# **Kubernetes**

## **Environment Variables & Arguments**

# Environment Variables

- 환경 변수를 Container에 선언하기 위해서는 `env:`를 사용하면 된다.

pod-env.yaml

```
apiVersion: v1
kind: Pod

metadata:
  name: pod-env

spec:
  containers:
    - name: print-env
      image: bash

      env:
        - name: FIRST
          value: "One"
        - name: SECOND
          value: "Two"
        - name: THIRD
          value: "Three"

      command: ["echo"]
      args: ["$(FIRST)", $(SECOND), $(THIRD)"]
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./05-week/env/pod-env.yaml
pod/pod-env created

remote > kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS
GATES
pod-env   0/1     Completed   0          3s     10.233.103.33   worker2   <none>       <none>
```

```
remote > kubectl logs pod-env
```

```
One, Two, Three
```

Dockerfile에서 환경 변수를 정의할 수 있지만,

그렇게 되면 환경 변수 값이 변경될 때 이미지를 변경해야 하기에 추천하지 않는다.

※ 참고 : <https://kubernetes.io/ko/docs/tasks/inject-data-application/define-environment-variable-container/>

# Arguments in Docker

- Container에서 arguments를 어떻게 사용하는지 먼저 살펴보고, Image를 만들어 보자.

fortuneloop.sh

```
#!/bin/bash
trap "exit" SIGINT
INTERVAL=$1
echo generate every $INTERVAL seconds
mkdir -p /var/htdocs

while :
do
    echo $(date) Writing to /var/htdocs/index.html
    /usr/games/fortune > /var/htdocs/index.html
    sleep $INTERVAL
done
```

Dockerfile

```
FROM ubuntu:latest

RUN apt-get update
RUN apt-get -y install fortune

ADD fortuneloop.sh /bin/fortuneloop.sh
RUN chmod +x /bin/fortuneloop.sh

ENTRYPOINT ["/bin/fortuneloop.sh"]
CMD ["10"]
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes/05-week/arguments

remote > docker build -t fortune:v1.0 .

...
remote > docker run -it --name fortune fortune:v1.0

generate every 10 seconds
Sat Feb 12 18:28:10 UTC 2022 Writing to /var/htdocs/index.html
Sat Feb 12 18:28:20 UTC 2022 Writing to /var/htdocs/index.html
Sat Feb 12 18:28:30 UTC 2022 Writing to /var/htdocs/index.html
^C
```

10초 간격으로 동작하는 것을  
볼 수 있다.

( DockerHub에서 fortune repository 생성 )

```
remote > docker login

remote > docker tag fortune:v1.0 whatwant/fortune:v1.0

remote > docker push whatwant/fortune:v1.0
```

※ 참고 : <https://github.com/luksa/kubernetes-in-action/tree/master/Chapter07>

# Arguments in Kubernetes - 1/2

- 이번에는 Kubernetes에서 arguments를 어떻게 사용하는지 살펴보자.

pod-fortune.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-fortune
  labels:
    app: fortune
spec:
  containers:
    - name: html-generator
      image: whatwant/fortune:v1.0
      args: ["5"]
      volumeMounts:
        - name: web-fortune
          mountPath: /var/htdocs
    - name: web-server
      image: nginx:alpine
      volumeMounts:
        - name: web-fortune
          mountPath: /usr/share/nginx/html
          readOnly: true
      ports:
        - containerPort: 80
  volumes:
    - name: web-fortune
      emptyDir: {}
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./05-week/arguments/pod-fortune.yaml
```

```
pod/pod-fortune created
```

```
remote > kubectl logs pod-fortune -c html-generator
```

```
generate every 5 seconds
Sat Feb 12 19:31:43 UTC 2022 Writing to /var/htdocs/index.html
Sat Feb 12 19:31:48 UTC 2022 Writing to /var/htdocs/index.html
Sat Feb 12 19:31:53 UTC 2022 Writing to /var/htdocs/index.html
```

5초 간격으로

변경된 것을 확인할 수 있다.

설명	도커 필드 이름	쿠버네티스 필드 이름
컨테이너에서 실행되는 커맨드	Entrypoint	command
커맨드에 전달되는 인자들	Cmd	arg

command 값과 args 값을 동시에 정의한다면,

Docker 이미지 안에 정의된 기본 EntryPoint 값과

기본 Cmd 값이 덮어 쓰여진다.

※ 참고 : <https://kubernetes.io/ko/docs/tasks/inject-data-application/define-command-argument-container/>

# Arguments in Kubernetes - 2/2

- 이왕 만들어진 App이니, Service까지 해서 어떤 결과로 나오는지 확인해보자.

svc-fortune.yaml

```
apiVersion: v1
kind: Service

metadata:
  name: svc-fortune

spec:
  type: NodePort

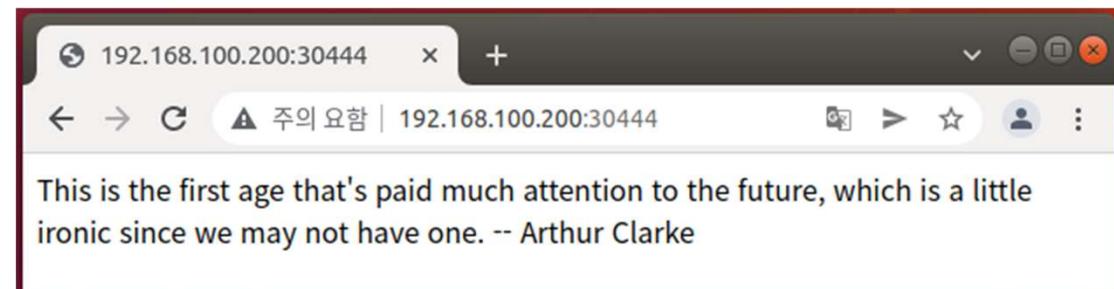
  ports:
  - port: 80
    targetPort: 80
    nodePort: 30444

  selector:
    app: fortune
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./05-week/arguments/svc-fortune.yaml
service/svc-fortune created

remote > kubectl get services -o wide
NAME         TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE   SELECTOR
kubernetes   ClusterIP   10.233.0.1    <none>          443/TCP       22d   <none>
svc-fortune  NodePort    10.233.0.134  <none>          80:30444/TCP  17h   app=fortune
```



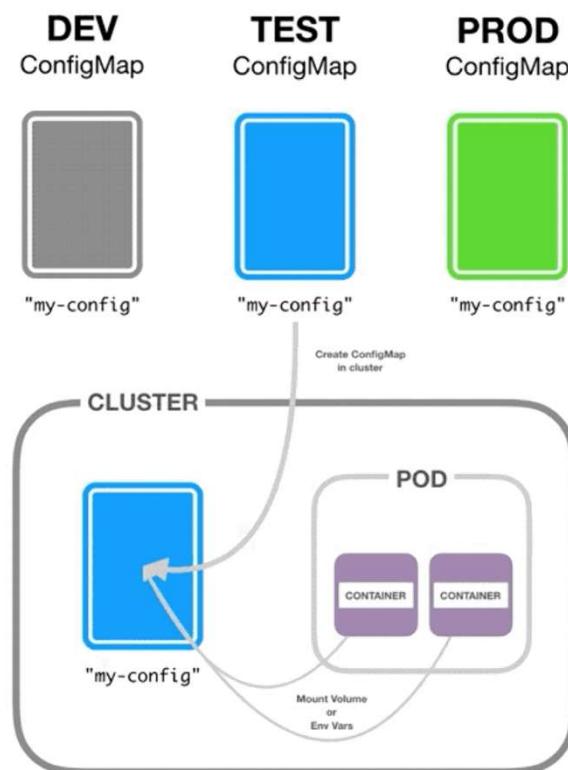
///

# **Kubernetes**

## **Volume - configMap**

# Why configMap ?

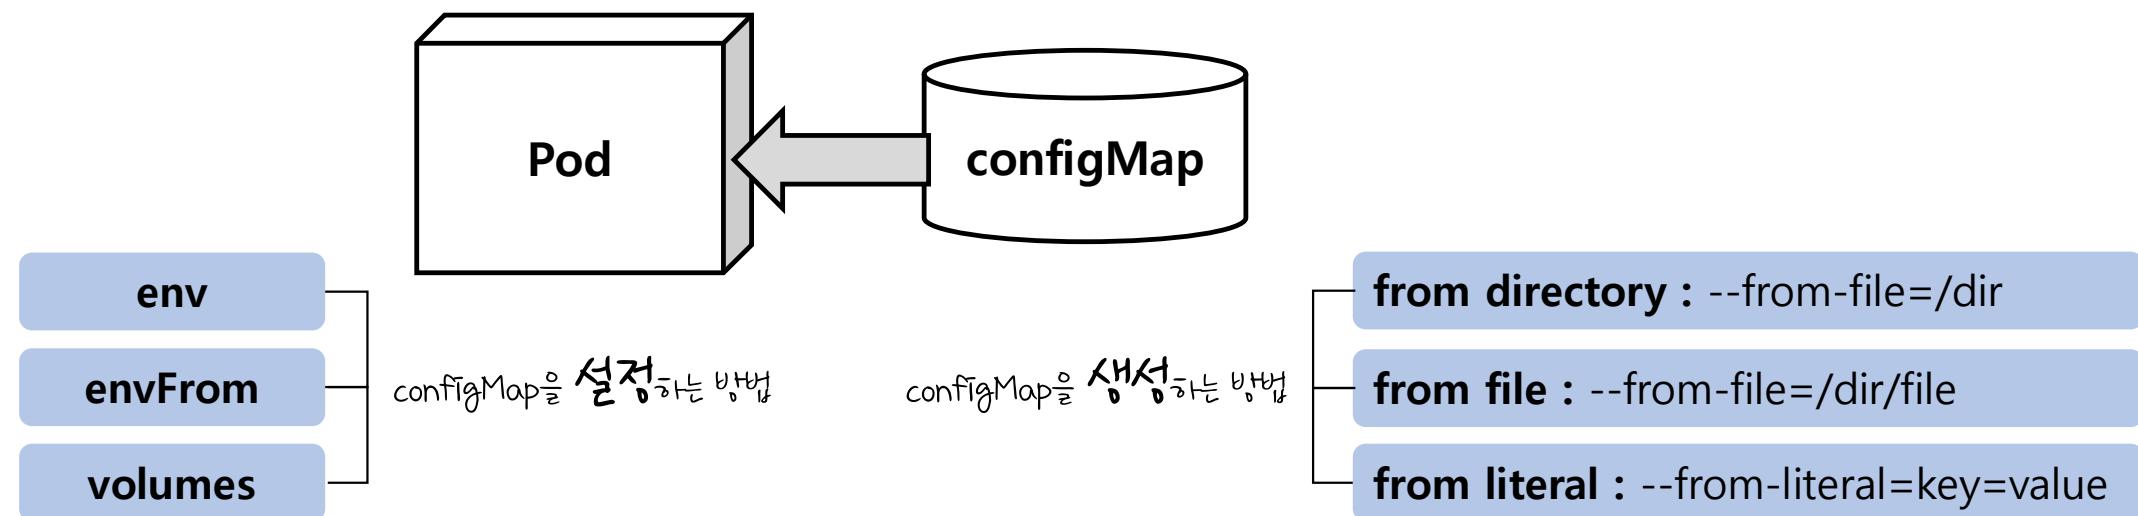
- 독립적인 리소스로 동일한 Pod라도 여러 개의 configMap을 선택하여 사용 가능
  - . configMap만 변경하여 Pod를 손쉽게 테스트



※ 참고 : <https://timewizhan.tistory.com/entry/Kubernetes-ConfigMap>

# How configMap ?

- configMap을 생성해 놓으면, Pod에서 가져다 사용하는 방식



# Hands-On : YAML - 1/2

- YAML을 이용한 configMap 사용법을 알아보자

envFrom & volumes 의 2가지 방식으로  
configMap을 사용하는 방법을  
살펴보고자 한다.

cm-envfrom.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: cm-envfrom
data:
  map-hash-bucket-size: "128"
  ssl-protocols: SSLv2
```

pod-configmap.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-configmap
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
  envFrom:
    - configMapRef:
        name: cm-envFrom
  volumeMounts:
    - name: nginx-volume
      subPath: index.html
      mountPath: /usr/share/nginx/html/index.html
```

cm-volume.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: cm-volume
data:
  index.html: |->
    <html>
    <h1>Hello from ConfigMap</h1>
    </html>
```

```
volumes:
  - name: nginx-volume
    configMap:
      name: cm-volume
```

svc-configmap.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: svc-configmap
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: 80
      nodePort: 30501
  selector:
    app: nginx
```

# Hands-On : YAML - 2/2

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./05-week/configMap/cm-envfrom.yaml
remote > kubectl create -f ./05-week/configMap/cm-volume.yaml
remote > kubectl create -f ./05-week/configMap/pod-configmap.yaml
remote > kubectl create -f ./05-week/configMap/svc-configmap.yaml
```

```
remote > kubectl get pods -o wide
```

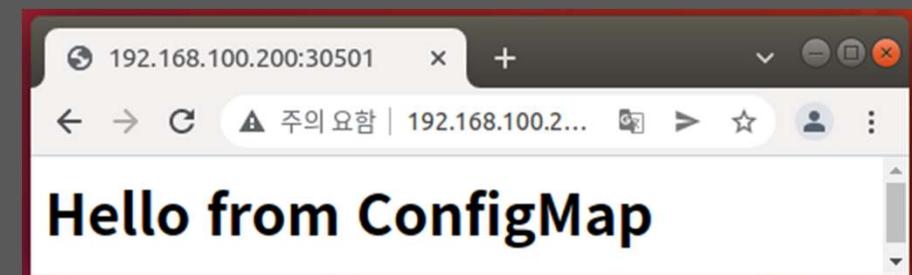
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pod-configmap	1/1	Running	0	4m5s	10.233.103.35	worker2	<none>	<none>

```
remote > kubectl exec -it pod-configmap -- printenv
```

```
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=pod-configmap
NGINX_VERSION=1.14.2-1~stretch
NJS_VERSION=1.14.2.0.2.6-1~stretch
map-hash-bucket-size=128
ssl-protocols=SSLv2
...
...
```

```
remote > kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.233.0.1	<none>	443/TCP	22d	<none>
svc-configmap	NodePort	10.233.37.4	<none>	80:30501/TCP	6m20s	app=nginx



# Hands-On : CLI

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git  
remote > cd advanced-kubernetes/05-week/configMap
```

```
remote > kubectl get configmaps
```

NAME	DATA	AGE
kube-root-ca.crt	1	22d

포트에서 작동하는

모든 것을 지원하는 상태

```
remote > cat ./cm-env/map-hash-bucket-size
```

128

```
remote > cat ./cm-env/ssl-protocols
```

SSLv2

```
remote > cat ./cm-vol/index.html
```

```
<html>  
<h1>Hello from ConfigMap</h1>  
</html>
```

```
remote > kubectl create configmap cm-envfrom --from-file=./cm-env/
```

configmap/cm-envfrom created

```
remote > kubectl create configmap cm-volume --from-file=./cm-vol/index.html
```

configmap/cm-volume created

```
remote > kubectl get configmaps
```

NAME	DATA	AGE
cm-envfrom	2	4m18s
cm-volume	1	7s
kube-root-ca.crt	1	22d

```
remote > kubectl create -f ./pod-configmap.yaml
```

pod/pod-configmap created

```
remote > kubectl create -f ./svc-configmap.yaml
```

service/svc-configmap created

```
remote > kubectl exec -it pod-configmap -- printenv
```

...  
map-hash-bucket-size=128  
ssl-protocols=SSLv2  
...



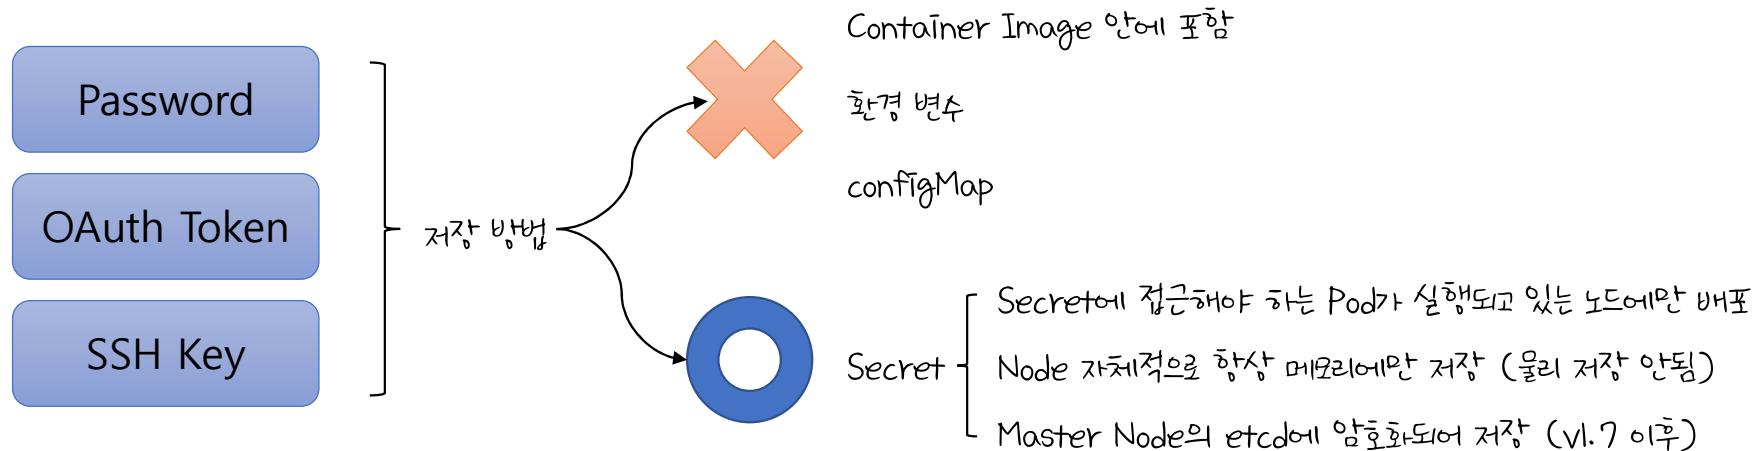
///

# **Kubernetes**

## **Volume - Secret**

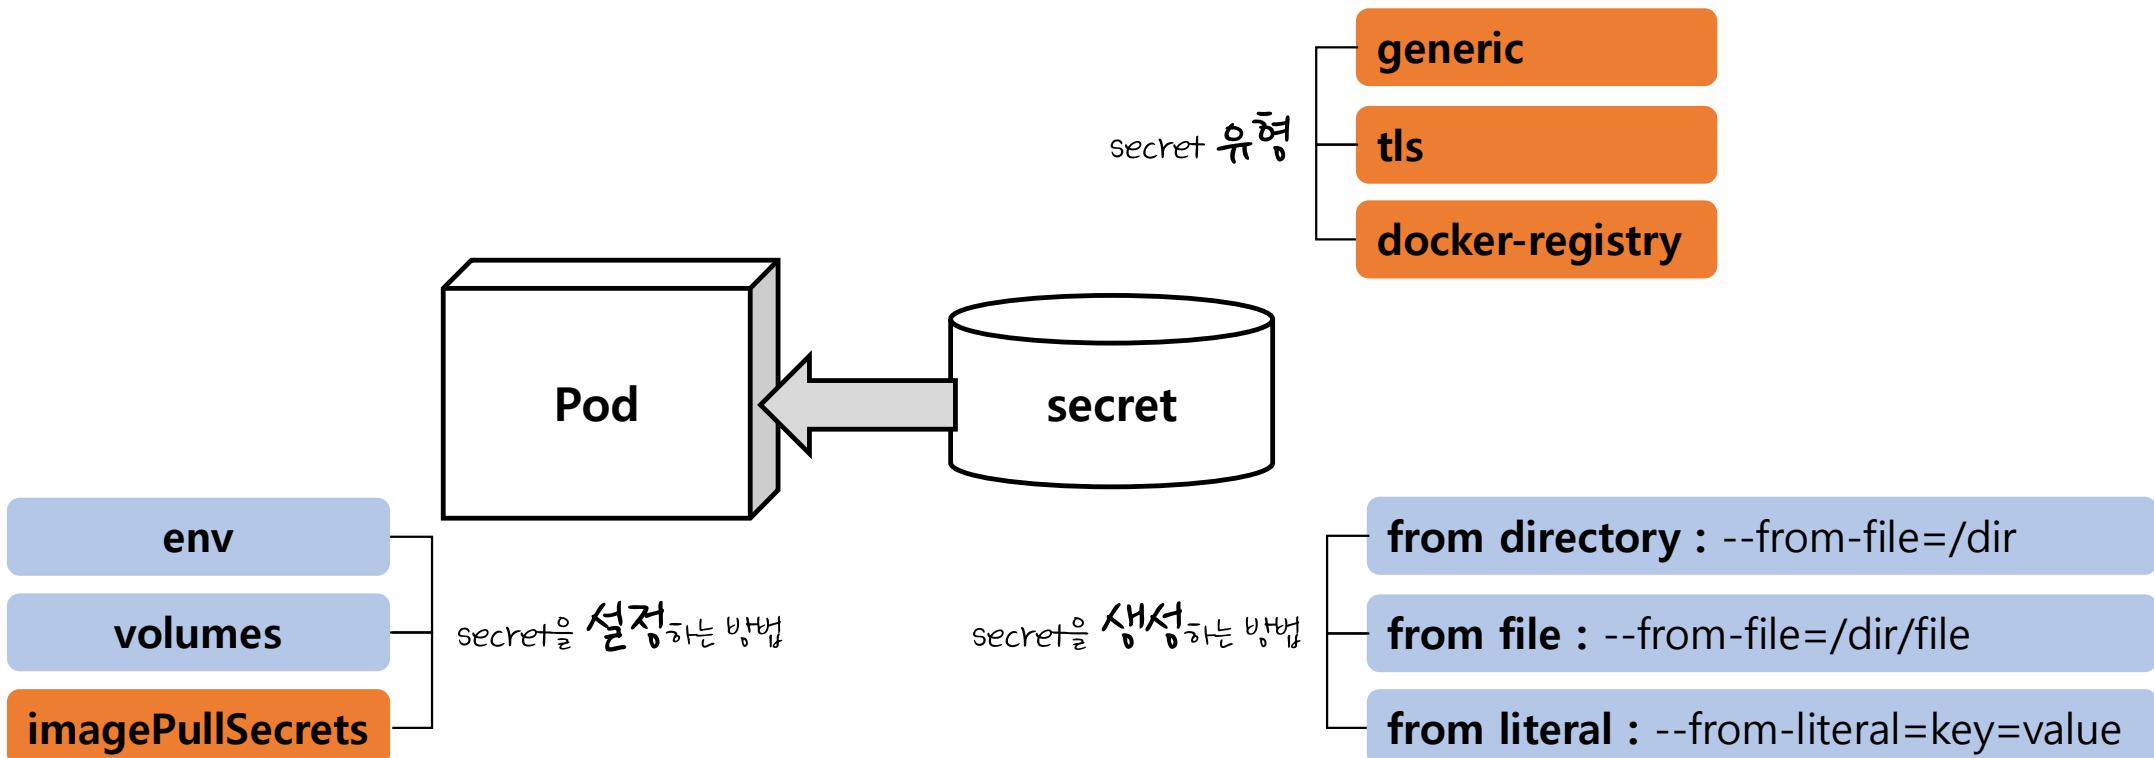
# Why secret ?

- 노출되면 안되는 정보를 저장하기 위한 방법 필요



# How secret ?

- configMap 사용법과 유사



# default token

- 기본적으로 생성되어 있는 secret을 확인해보자

```
remote > kubectl get secrets -o wide
```

NAME	TYPE	DATA	AGE	
default-token-xf884	kubernetes.io/service-account-token	3	23d	3개의 secret 데이터 존재

```
remote > kubectl describe secrets default-token-xf884
```

```
Name:           default-token-xf884
Namespace:      default
Labels:         <none>
Annotations:   kubernetes.io/service-account.name: default
               kubernetes.io/service-account.uid: e72b51f8-e9f6-44b4-abf1-4f38867c35ee
```

```
Type:  kubernetes.io/service-account-token
```

```
Data
```

```
====
```

```
ca.crt: 1099 bytes      Kubernetes API 통신을 위한 3종 데이터 정보
namespace: 7 bytes
token:
```

```
eyJhbGciOiJSUzI1NiIsImtpZCI6IjJHUXBHX0NzeUl0a0xVSDh5Nk9CRVNleVBKcmxWdThRTVFubWR4d0pBc0UiFQ.eyJpc3Mi0iJrdWJlcm5ldGVzL3NlcnZpY2VhY2NvdW50Iiwia3ViZXJuZXRLc
y5pb9zZXJ2aNlYWNjb3VudC9uYW1lc3BhY2Ui0iJkZWZhdWx0Iiwia3ViZXJuZXRLcy5pb9zZXJ2aNlYWNjb3VudC9zZWNyZXQubmFtZSI6ImRlZmF1bHQtdG9rZW4teGY40DQiLCJrdWJlcm5ld
GVzLmlvL3NlcnZpY2VhY2NvdW50L3NlcnZpY2UtYWNjb3VudC5uYW1lIjoicGVmYXVsdcIsImt1YmVybmv0ZXMuaw8vc2VydmljZWFjY291bnQvc2VydmljZS1hY2NvdW50LnVpZCI6ImU3MmI1MWY4L
WU5ZjYtNDRIiNC1hYmYxLTRmMzg4NjdjMzVlZSiInN1YiI6InN5c3R1bTpzZXJ2aNlYWNjb3VudDpkZWZhdWx00mRlZmF1bHQifQ.lmbn_Ucr-
gKqrxxkU4zYQgqxeABzFac4ml21HPz4GexH1AHknGPg8--pnlihCQq32zrTmRgL-Y6CYd6_Wc4XTEu_fv14Yqp71wiN9k4IvAlqxY-X591-0Vr_LMIyG8S13-DZnwKrwciWKhu-
SG557X1U29hjP9gbsQryBoI69Rg0D0NgQNdpPeKQZwHdf5uh0F8vpd5rDzE2tK_dHvTz8zpofKKf9g1cDqMnJB00Nag2Rm2xSwDvKdRsWUYdzuhGgrx3iYxd0qUQuN064u-
jE4pb4KLMXJyR42lTX8iw5x0ecTav1CXYlJ_PBFqUqGmmHPbD2nLDV7oTvS4qgZA71A
```

# default token in Pod

- Pod를 생성하면 기본적으로 API Server와 통신하기 위한 secret이 mount 된다.

```
remote > kubectl run web --image=nginx:latest
```

pod/web created

CLI 활용한 Pod 생성

```
remote > kubectl describe pods web
```

Name: web  
Namespace: default  
Priority: 0  
Node: worker2/192.168.100.202  
Start Time: Tue, 15 Feb 2022 00:11:11 +0900  
Labels: run=web  
Annotations: cni.projectcalico.org/containerID: 6...  
cni.projectcalico.org/podIP: 10.233....  
cni.projectcalico.org/podIPs: 10.23....  
Status: Running  
IP: 10.233.103.39  
IPs:  
 IP: 10.233.103.39  
Containers:  
 web:  
 Container ID: containerd://48f16fc7f85e54703...  
 Image: nginx:latest  
 Image ID: docker.io/library/nginx@sha256...  
 Port: <none>  
 Host Port: <none>

...

```
Environment: <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-ggr4n (ro)
Conditions:
  Type        Status
  Initialized  True
  Ready       True
  ContainersReady  True
  PodScheduled  True
Volumes:
  kube-api-access-ggr4n:
    Type:      Projected (a volume that contains injected data from multiple ... )
    TokenExpirationSeconds: 3607
    ConfigMapName:          kube-root-ca.crt
    ConfigMapOptional:      <nil>
    DownwardAPI:            true
QoS Class:      BestEffort
Node-Selectors: <none>
Tolerations:
  node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
  node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type  Reason  Age   From           Message
  ----  -----  ---   ----           -----
  Normal Scheduled  17s  default-scheduler  Successfully assigned default/web to worker2
  Normal Pulling   17s  kubelet         Pulling image "nginx:latest"
  Normal Pulled    7s   kubelet         Successfully pulled image "nginx:latest" in 9.1...
  Normal Created   7s   kubelet         Created container web
  Normal Started   7s   kubelet         Started container web
```

///

# docker registry : Private Image

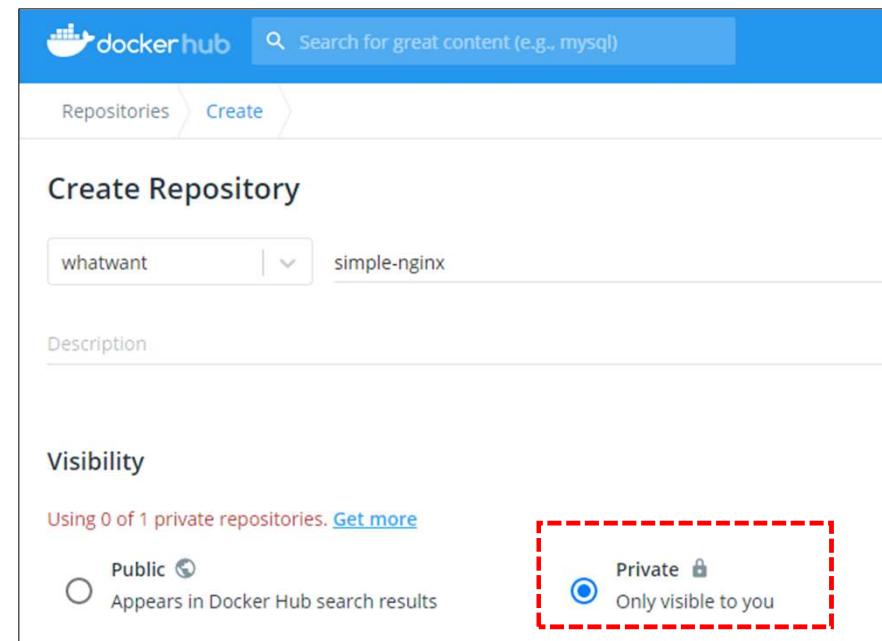
- Private Image를 하나 만들어서 Docker Hub에 업로드 해놓자.

index.html

```
<!doctype html>
<html lang="en">
<head>
  <meta charset="utf-8">
  <title>Docker Nginx</title>
</head>
<body>
  <h2>Hello from Nginx container</h2>
</body>
</html>
```

Dockerfile

```
FROM nginx:latest
COPY ./04-index.html /usr/share/nginx/html/index.html
```



```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes/05-week/secret
```

```
remote > docker login
```

```
remote > docker build -t whatwant/simple-nginx:v0.1 -f 04-Dockerfile .
```

local에 build 한 뒤, tagging하고 push를 해왔는데,

```
remote > docker push whatwant/simple-nginx:v0.1
```

이렇게 바로 DockerHub 用 tagging으로 build를 바로 해버릴 수도 있다.

# docker registry : Failed to pull image

- 앞에서 등록한 Private Image를 그냥 내려 받으면 어떻게 되는지 확인해보자.

pod-private-fail.yaml

```
apiVersion: v1
kind: Pod

metadata:
  name: pod-private
  labels:
    app: nginx

spec:
  containers:
    - name: nginx
      image: whatwant/simple-nginx:v0.1
```

```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes

remote > kubectl create -f ./05-week/secret/pod-private-fail.yaml
pod/pod-private created

remote > kubectl get pods -o wide
NAME        READY   STATUS     RESTARTS   AGE       IP          NODE   NOMINATED NODE   READINESS GATES
pod-private  0/1     ErrImagePull  0          16s      10.233.103.42  worker2  <none>        <none>
```

```
remote > kubectl describe pods pod-private
```

...

Events:

Type	Reason	Age	From	Message
---	-----	----	-----	-----
Normal	Scheduled	2m10s	default-scheduler	Successfully assigned default/pod-private to worker2
Normal	Pulling	31s (x4 over 2m9s)	kubelet	Pulling image "whatwant/simple-nginx:v0.1"
Warning	Failed	28s (x4 over 2m5s)	kubelet	Failed to pull image "whatwant/simple-nginx:v0.1": rpc error: code = Unknown desc = failed to pull and unpack image "docker.io/whatwant/simple-nginx:v0.1": failed to resolve reference "docker.io/whatwant/simple-nginx:v0.1": pull access denied, repository does not exist or may require authorization: server message: insufficient_scope: authorization failed
Warning	Failed	28s (x4 over 2m5s)	kubelet	Error: ErrImagePull
Warning	Failed	16s (x6 over 2m5s)	kubelet	Error: ImagePullBackOff
Normal	BackOff	4s (x7 over 2m5s)	kubelet	Back-off pulling image "whatwant/simple-nginx:v0.1"

# docker registry : Create secret

- Registry Server 인증 정보를 secret으로 등록하자

pod-private-success.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-private
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: whatwant/simple-nginx:v0.1
  imagePullSecrets:
    - name: my-docker-hub
```

```
kubectl create secret docker-registry <secret-name> \
  [ --docker-server=<your-registry-server> \ ]
  --docker-username=<your-name> \
  --docker-password=<your-password> \
  --docker-email=<your-email>
```

```
remote > kubectl create secret docker-registry my-docker-hub \
  --docker-username=whatwant \
  --docker-password='xxx' \
  --docker-email='whatwant@gmail.com'
```

secret/my-docker-hub created

```
remote > kubectl get secrets -o wide
```

NAME	TYPE	DATA	AGE
default-token-xf884	kubernetes.io/service-account-token	3	24d
my-docker-hub	kubernetes.io/dockerconfigjson	1	19s

```
remote > kubectl create -f ./05-week/secret/pod-private-success.yaml
pod/pod-private created
```

```
remote > kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
pod-private	1/1	Running	0	3m4s

# Type of Secret

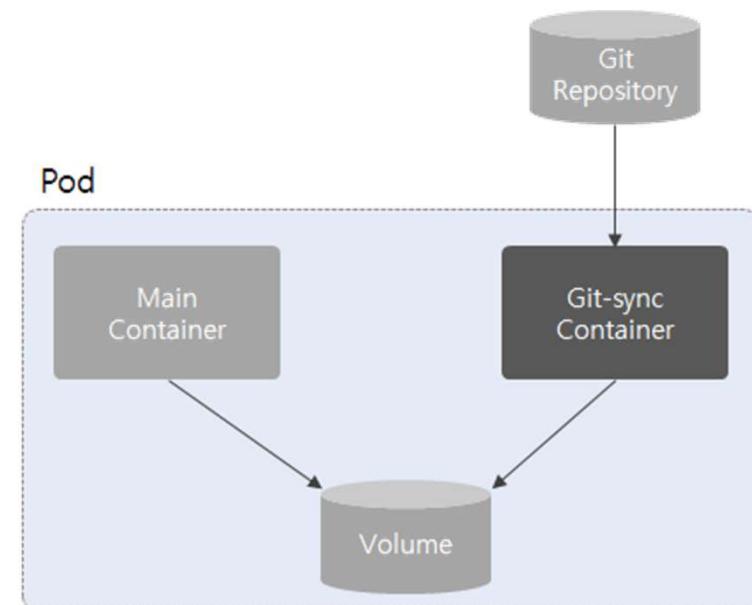
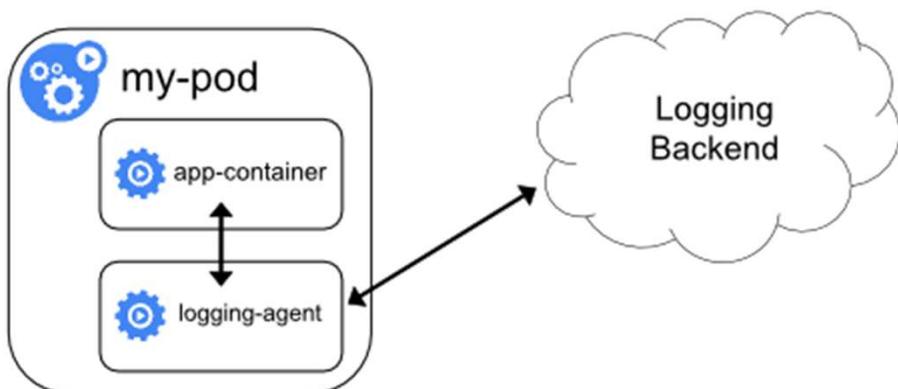
빌트인 타입	사용처
Opaque	임의의 사용자 정의 데이터
kubernetes.io/service-account-token	서비스 어카운트 토큰
kubernetes.io/dockercfg	직렬화 된(serialized) ~/.dockercfg 파일
 kubernetes.io/dockerconfigjson	직렬화 된 ~/.docker/config.json 파일
kubernetes.io/basic-auth	기본 인증을 위한 자격 증명(credential)
kubernetes.io/ssh-auth	SSH를 위한 자격 증명
kubernetes.io/tls	TLS 클라이언트나 서버를 위한 데이터
bootstrap.kubernetes.io/token	부트스트랩 토큰 데이터

※ 참고 : <https://kubernetes.io/ko/docs/concepts/configuration/secret/>

///

# volume gitRepo: Deprecated → 'git-sync' Sidecar Container

- 사이드카 패턴 (Sidecar Pattern)
  - . 기본 컨테이너의 기능을 확장하거나 강화하는 용도의 컨테이너를 추가하는 패턴
  - . 기본 컨테이너에는 원래 목적의 기능에만 충실하고 나머지 부가적인 공통 기능들은 사이드카 컨테이너를 추가해서 사용
- deprecate 된 gitRepo volume 대신 git-sync의 sidecar container를 구성하는 것으로 구현해야 한다



※ 참고 : <https://arisu1000.tistory.com/27863>

※ 참고 : <https://github.com/kubernetes/git-sync>

# secret with SSH

- ssh-key 정보와 known\_hosts 정보를 secret으로 등록해보자
  - . ssh-key를 아직 생성하지 않았다면, `ssh-keygen`으로 생성 후 진행하면 된다.

```
remote > ssh-keyscan github.com > ./known_hosts
```

```
# github.com:22 SSH-2.0-babeld-e1420b26  
# github.com:22 SSH-2.0-babeld-e1420b26  
# github.com:22 SSH-2.0-babeld-e1420b26
```

github.com 접근을 위해 미리 known\_hosts에 등록해야 한다.

그렇지 않으면, 처음 접근하는 서버의 등록 과정으로 문제가 발생하게 된다.

```
remote > kubectl create secret generic my-ssh --from-file=ssh=$HOME/.ssh/id_rsa --from-file=known_hosts=./known_hosts
```

```
secret/my-ssh created
```

```
remote > kubectl get secrets -o wide
```

NAME	TYPE	DATA	AGE
default-token-xf884	kubernetes.io/service-account-token	3	24d
my-ssh	Opaque	2	12s

# 'git-sync' Sidecar Container - 1/2

pod-git-sync.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: git-sync

spec:
  containers:
    - name: ubuntu
      image: ubuntu:latest
      args:
        - sleep
        - "370000"

      volumeMounts:
        - name: myrepo
          mountPath: "/repo"

        - name: git-sync
          image: k8s.gcr.io/git-sync/git-sync:v3.2.2
          args:
            - -ssh
            - -repo=git@github.com:whatwant-school/private-repo.git
            - -root=/repo
            - -dest=private-repo
            - -branch=main
            - -depth=1
```

```
volumeMounts:
  - name: myrepo
    mountPath: "/repo"

  - name: git-secret
    mountPath: /etc/git-secret
    readOnly: true

securityContext:
  runAsUser: 65533

securityContext:
  fsGroup: 65533

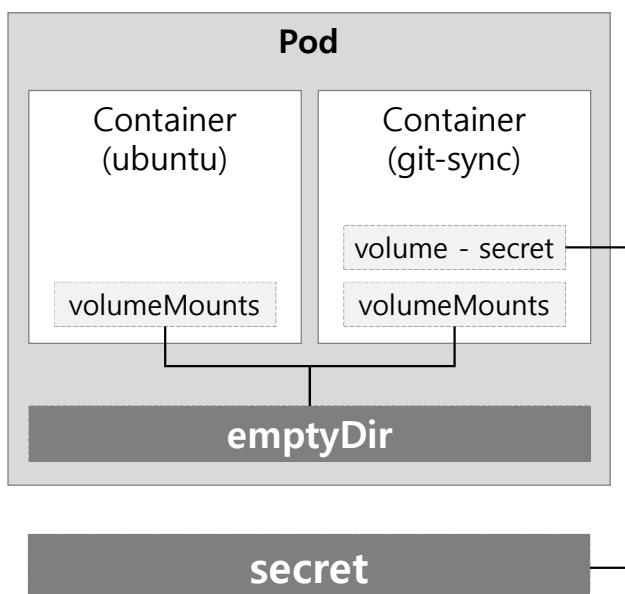
restartPolicy: Never

volumes:
  - name: myrepo
    emptyDir: {}

  - name: git-secret
    secret:
      secretName: my-ssh
      defaultMode: 0400
```

※ 참고 : <https://github.com/kubernetes/git-sync/blob/release-3.x/docs/ssh.md>

# 'git-sync' Sidecar Container - 2/2



```
remote > git clone https://github.com/whatwant-school/advanced-kubernetes.git
remote > cd advanced-kubernetes
```

```
remote > kubectl create -f ./05-week/secret/pod-git-sync.yaml
```

```
pod/git-sync created
```

```
remote > kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
git-sync	2/2	Running	0	32s	10.233.103.50	worker2	<none>	<none>

```
remote > kubectl exec -it git-sync -c ubuntu -- /bin/bash
```

```
groups: cannot find name for group ID 65533
```

```
root@git-sync:/# ls -al /repo
```

```
total 16
```

```
drwxrwsrwx 4 root 65533 4096 Feb 15 17:56 .
```

```
drwxr-xr-x 1 root root 4096 Feb 15 17:56 ..
```

```
drwxr-sr-x 9 65533 65533 4096 Feb 15 17:56 .git
```

```
lrwxrwxrwx 1 65533 65533 44 Feb 15 17:56 private-repo -> rev-56af7aa4197fee5da30ccdd562cc2ea5e751af76
```

```
drwxr-sr-x 2 65533 65533 4096 Feb 15 17:56 rev-56af7aa4197fee5da30ccdd562cc2ea5e751af76
```

```
root@git-sync:/# ls -al /repo/private-repo/
```

```
total 16
```

```
drwxr-sr-x 2 65533 65533 4096 Feb 15 17:56 .
```

```
drwxrwsrwx 4 root 65533 4096 Feb 15 17:56 ..
```

```
-rw-r--r-- 1 65533 65533 71 Feb 15 17:56 .git
```

```
-rw-r--r-- 1 65533 65533 60 Feb 15 17:56 README.md
```

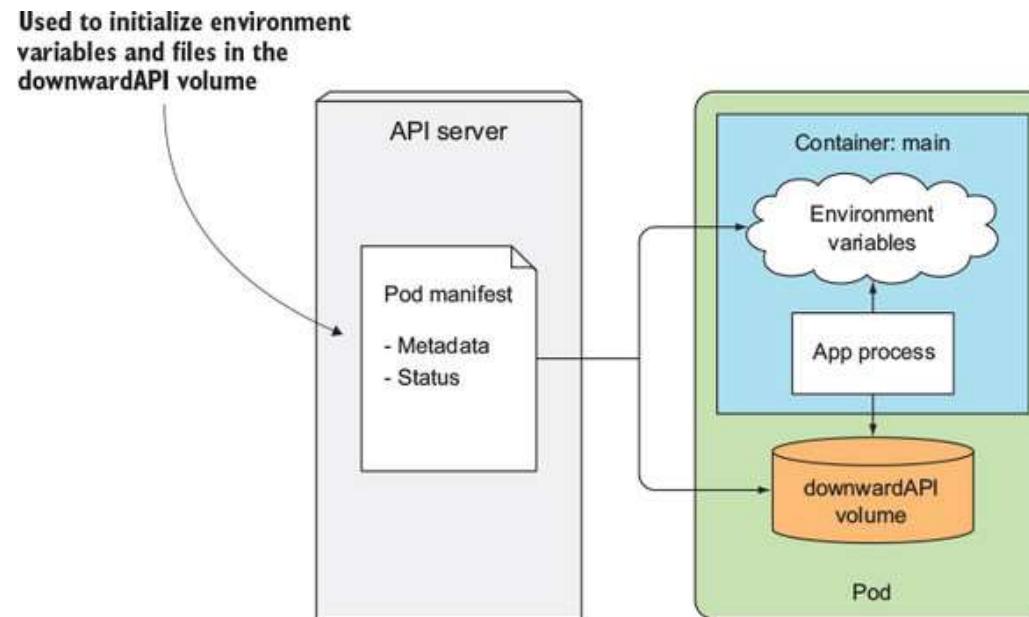
///

# **Kubernetes**

## **Volume - downwardAPI**

# Why downwardAPI

- Pod의 IP, 호스트 노드 이름, Pod 자체의 이름과 같이 실행 시점까지 알려지지 않은 데이터를 얻기 위한 방법
  - . downwardAPI는 애플리케이션이 호출해서 데이터를 가져오는 REST Endpoint와는 다르다.



※ 참고 : <https://livebook.manning.com/book/kubernetes-in-action/chapter-8/11>

# capabilities of downwardAPI

Information available via	item	comment
fieldRef:	metadata.name	the pod's name
	metadata.namespace	the pod's namespace
	metadata.uid	the pod's UID
	metadata.labels['<KEY>']	the value of the pod's label <KEY> (for example, metadata.labels['mylabel'])
	metadata.annotations['<KEY>']	the value of the pod's annotation <KEY> (for example, metadata.annotations['myannotation'])
resourceFieldRef:	A Container's CPU limit	
	A Container's CPU request	
	A Container's memory limit	
	A Container's memory request	
	A Container's hugepages limit	providing that the DownwardAPIHugePages feature gate is enabled
	A Container's hugepages request	providing that the DownwardAPIHugePages feature gate is enabled
	A Container's ephemeral-storage limit	
	A Container's ephemeral-storage request	
downwardAPI volume fieldRef:	metadata.labels	all of the pod's labels, formatted as label-key="escaped-label-value" with one label per line
	metadata.annotations	all of the pod's annotations, formatted as annotation-key="escaped-annotation-value" with one annotation per line
environment variables:	status.podIP	the pod's IP address
	spec.serviceAccountName	the pod's service account name, available since v1.4.0-alpha.3
	spec.nodeName	the node's name, available since v1.4.0-alpha.3
	status.hostIP	the node's IP, available since v1.7.0-alpha.1

※ 참고 : <https://kubernetes.io/docs/tasks/inject-data-application/downward-api-volume-expose-pod-information/#capabilities-of-the-downward-api>

# downwardAPI volume example - 1/2

dapi-volume.yaml

```
apiVersion: v1
kind: Pod

metadata:
  name: kubernetes-downwardapi-volume-example
  labels:
    zone: us-est-coast
    cluster: test-cluster1
    rack: rack-22
  annotations:
    build: two
    builder: john-doe

spec:
  containers:
    - name: client-container
      image: k8s.gcr.io/busybox
      command: ["sh", "-c"]
      args:
        - while true; do
            if [[ -e /etc/podinfo/labels ]]; then
              echo -en '$n$n'; cat /etc/podinfo/labels; fi;
            if [[ -e /etc/podinfo/annotations ]]; then
              echo -en '$n$n'; cat /etc/podinfo/annotations; fi;
            sleep 5;
        done;

  volumes:
    - name: podinfo
      downwardAPI:
        items:
          - path: "labels"
            fieldRef:
              fieldPath: metadata.labels
          - path: "annotations"
            fieldRef:
              fieldPath: metadata.annotations

      volumeMounts:
        - name: podinfo
          mountPath: /etc/podinfo
```

※ 참고 : <https://k8s.io/examples/pods/inject/dapi-volume.yaml>

///

**Tip**

# 쿠버네티스 어린이 그림 가이드



※ 참고 : <https://www.youtube.com/watch?v=4ht22ReBjno>

///

<https://kahoot.it/>

[ Score ]

- 이민준 (7)
- 김남형 (4)
- 박남준 (3)
- 이혜정 (3)
- 김상호 (2)
- 이원준 (2)
- 정현찬 (1)
- 김정은 (1)